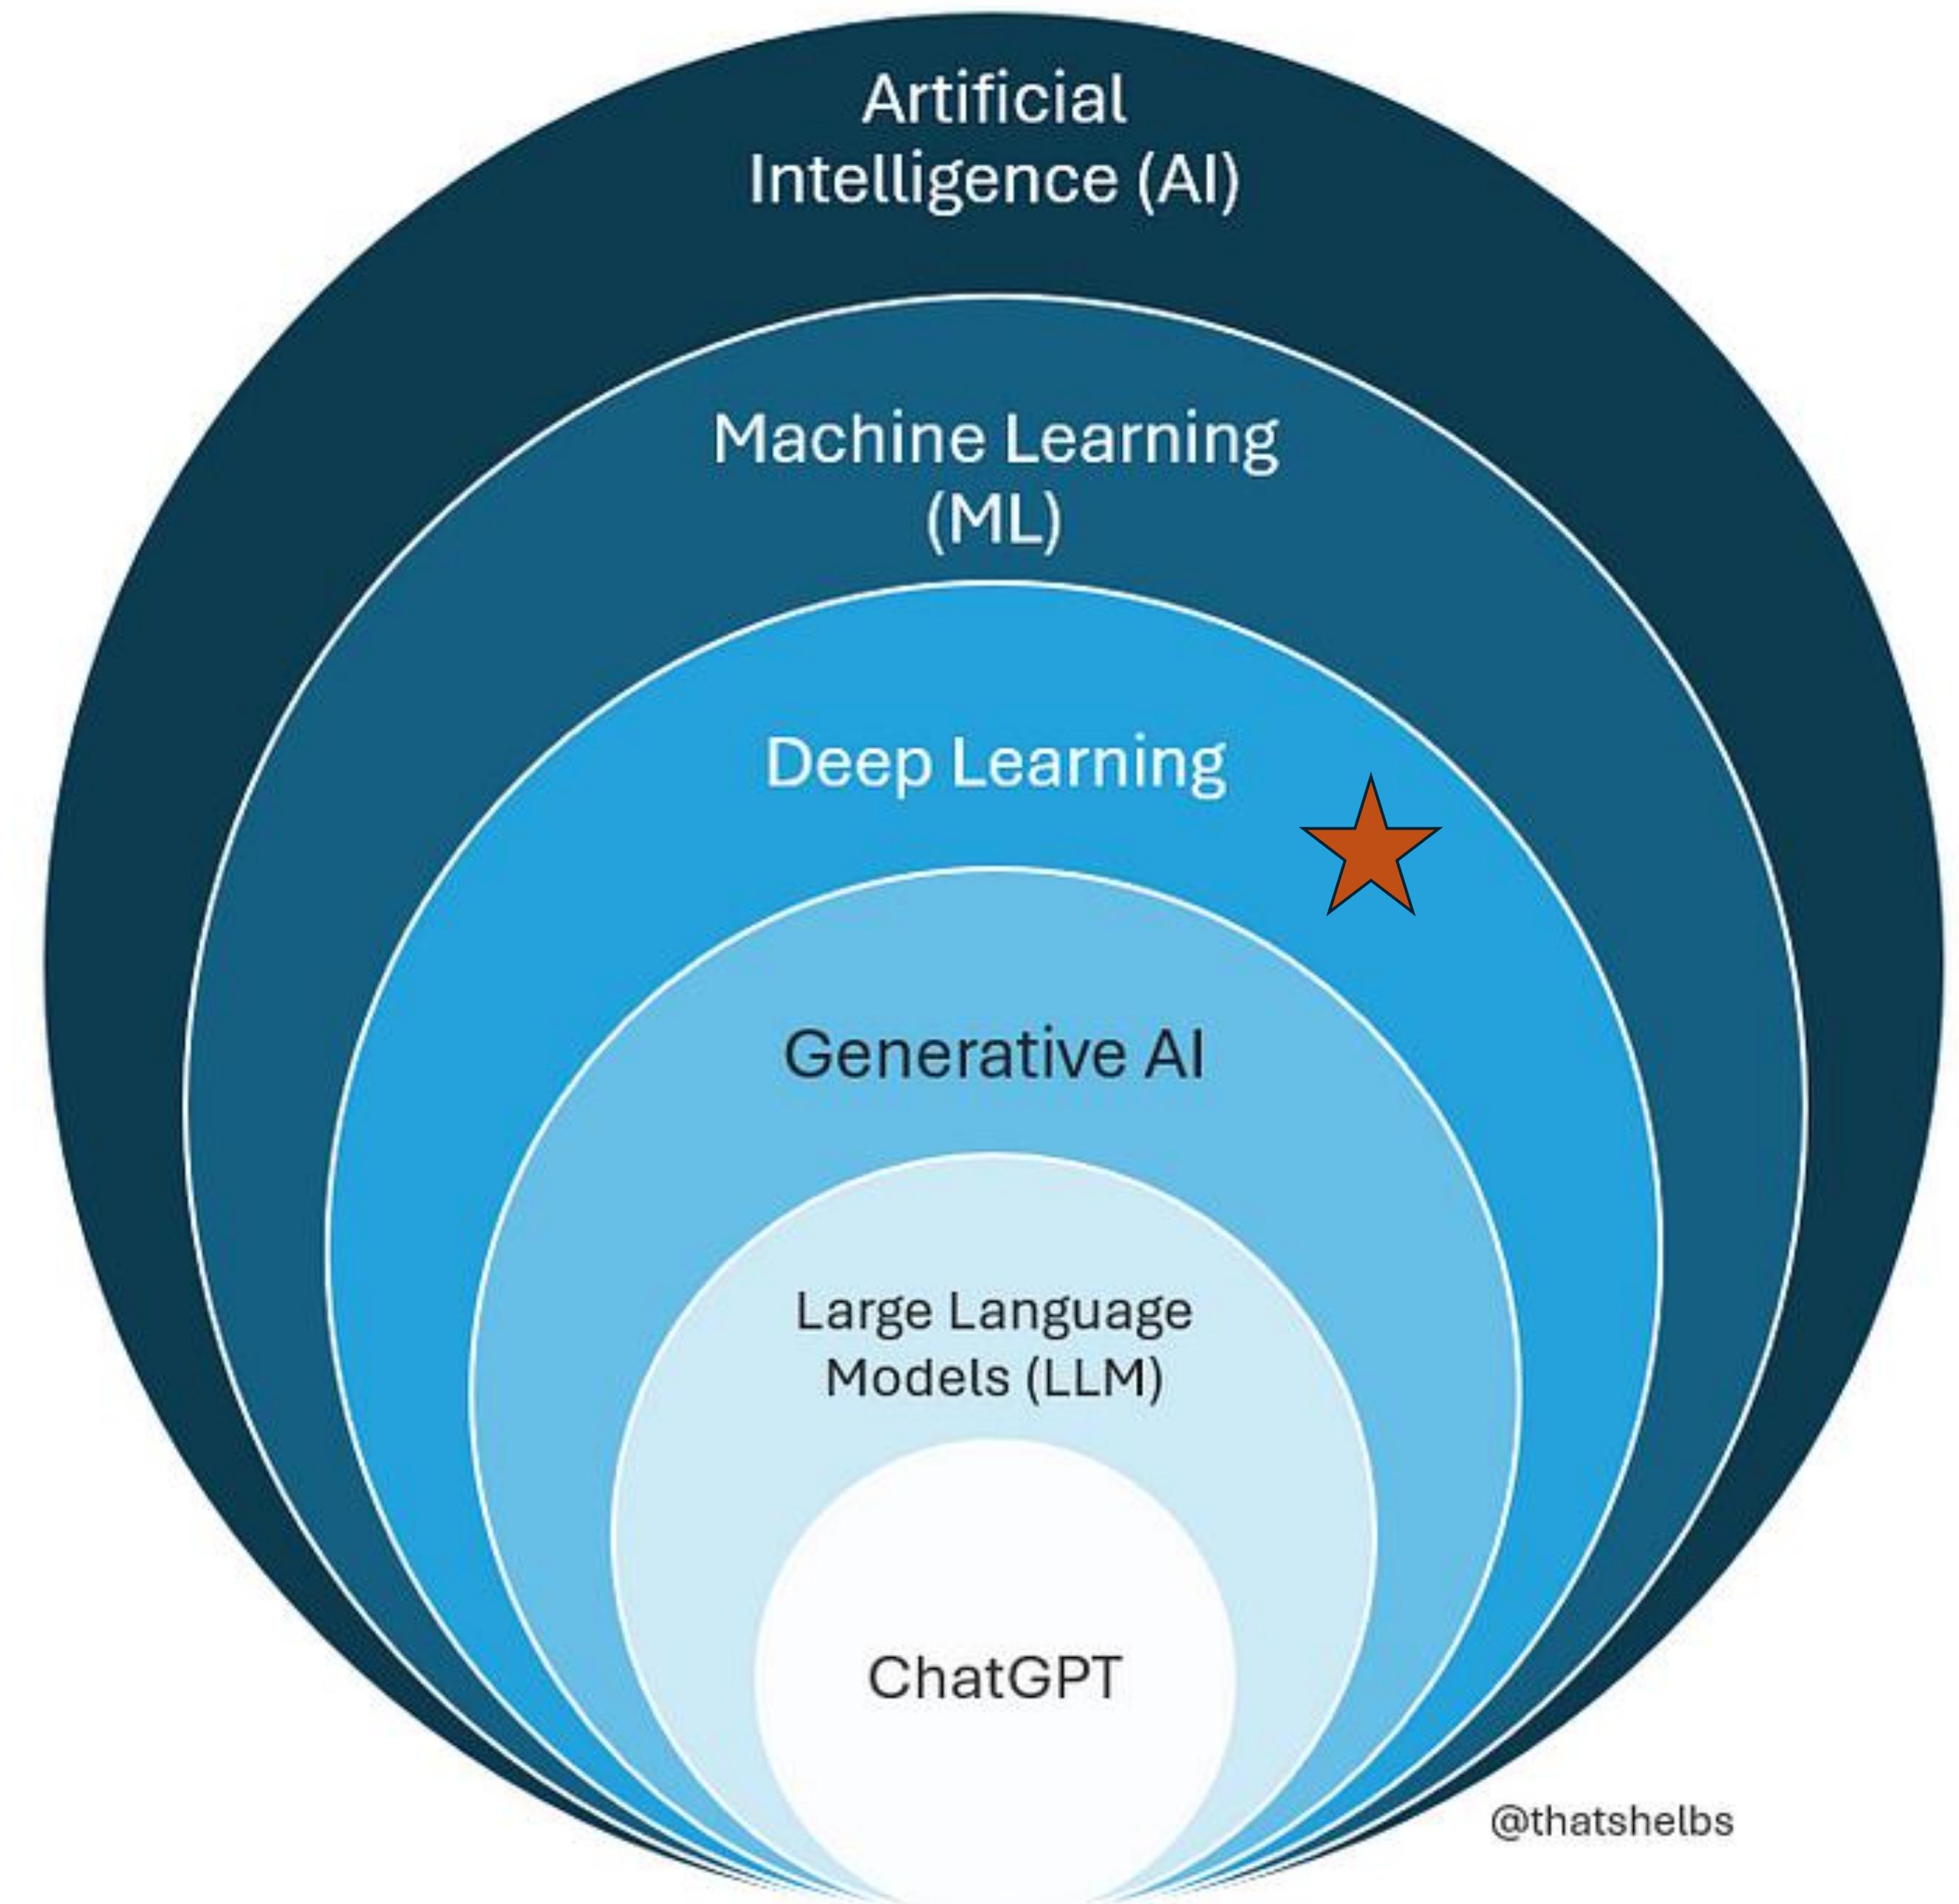


# Introduction to Deep Learning (Computer vision)

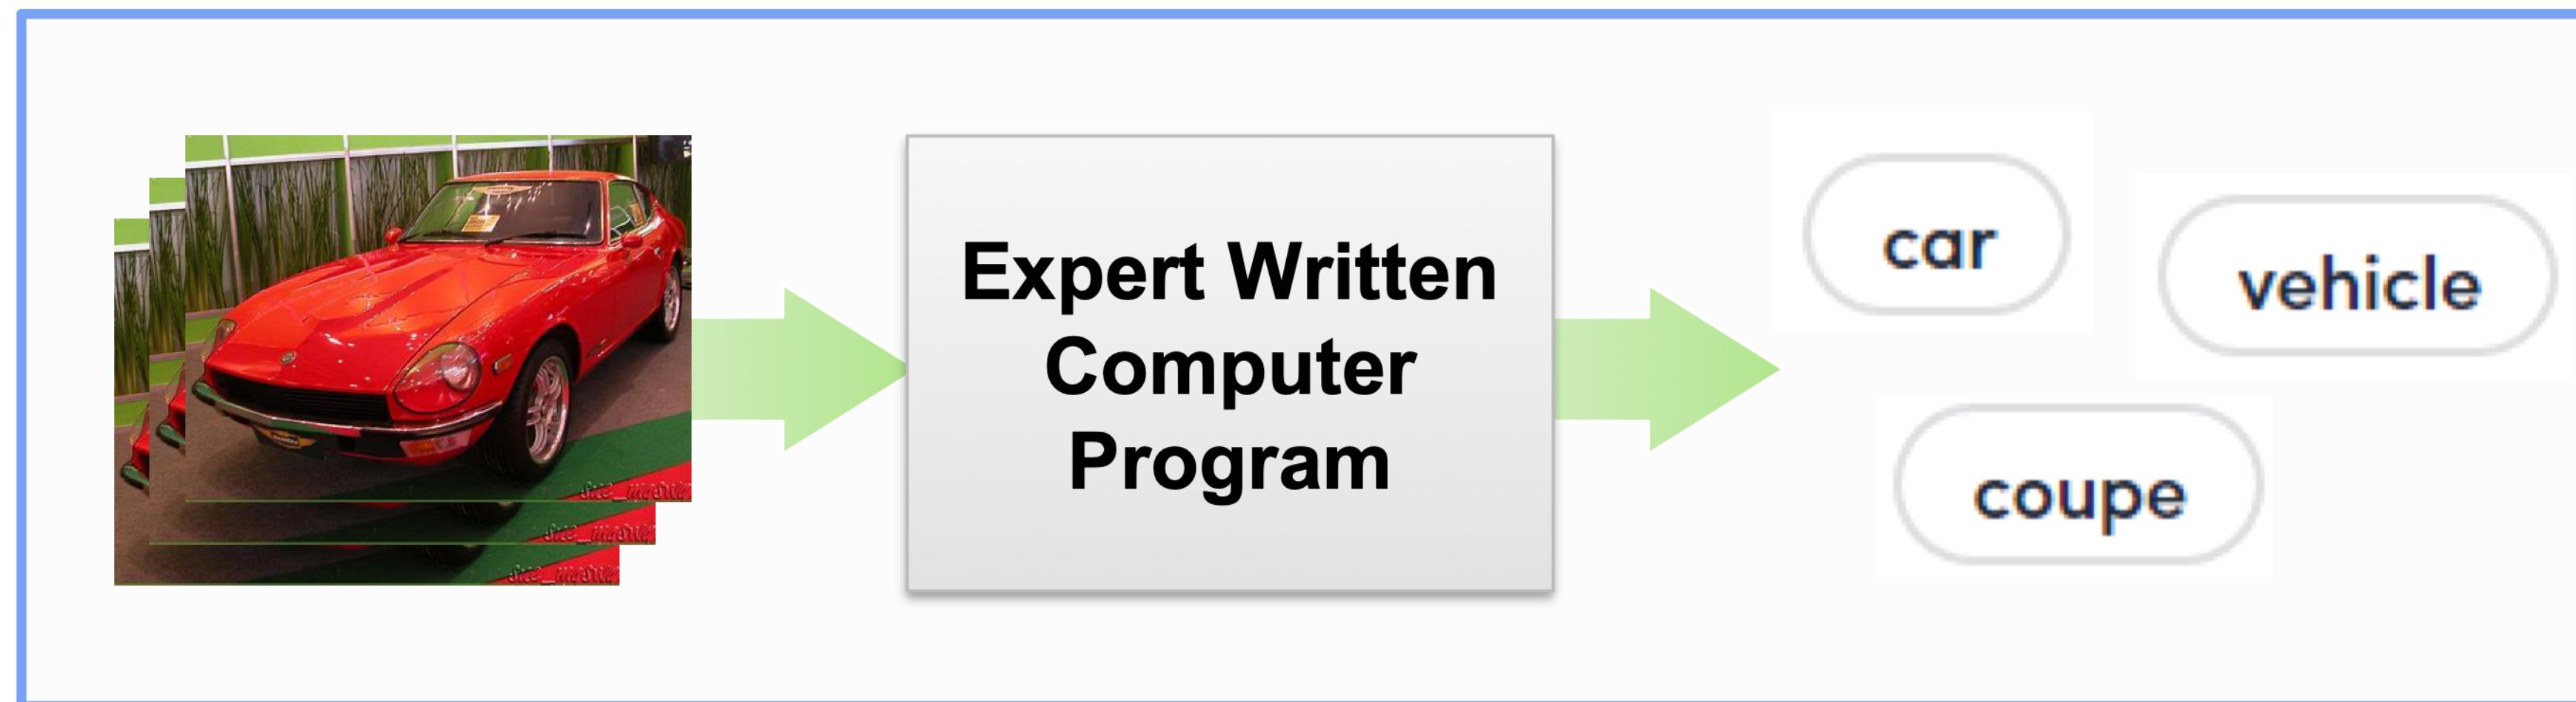
Aaron Y. Lee MD MSCI  
University of Washington



Computer vision:  
Giving machines the ability to “see.”

# A NEW COMPUTING MODEL

## Algorithms that Learn from Examples

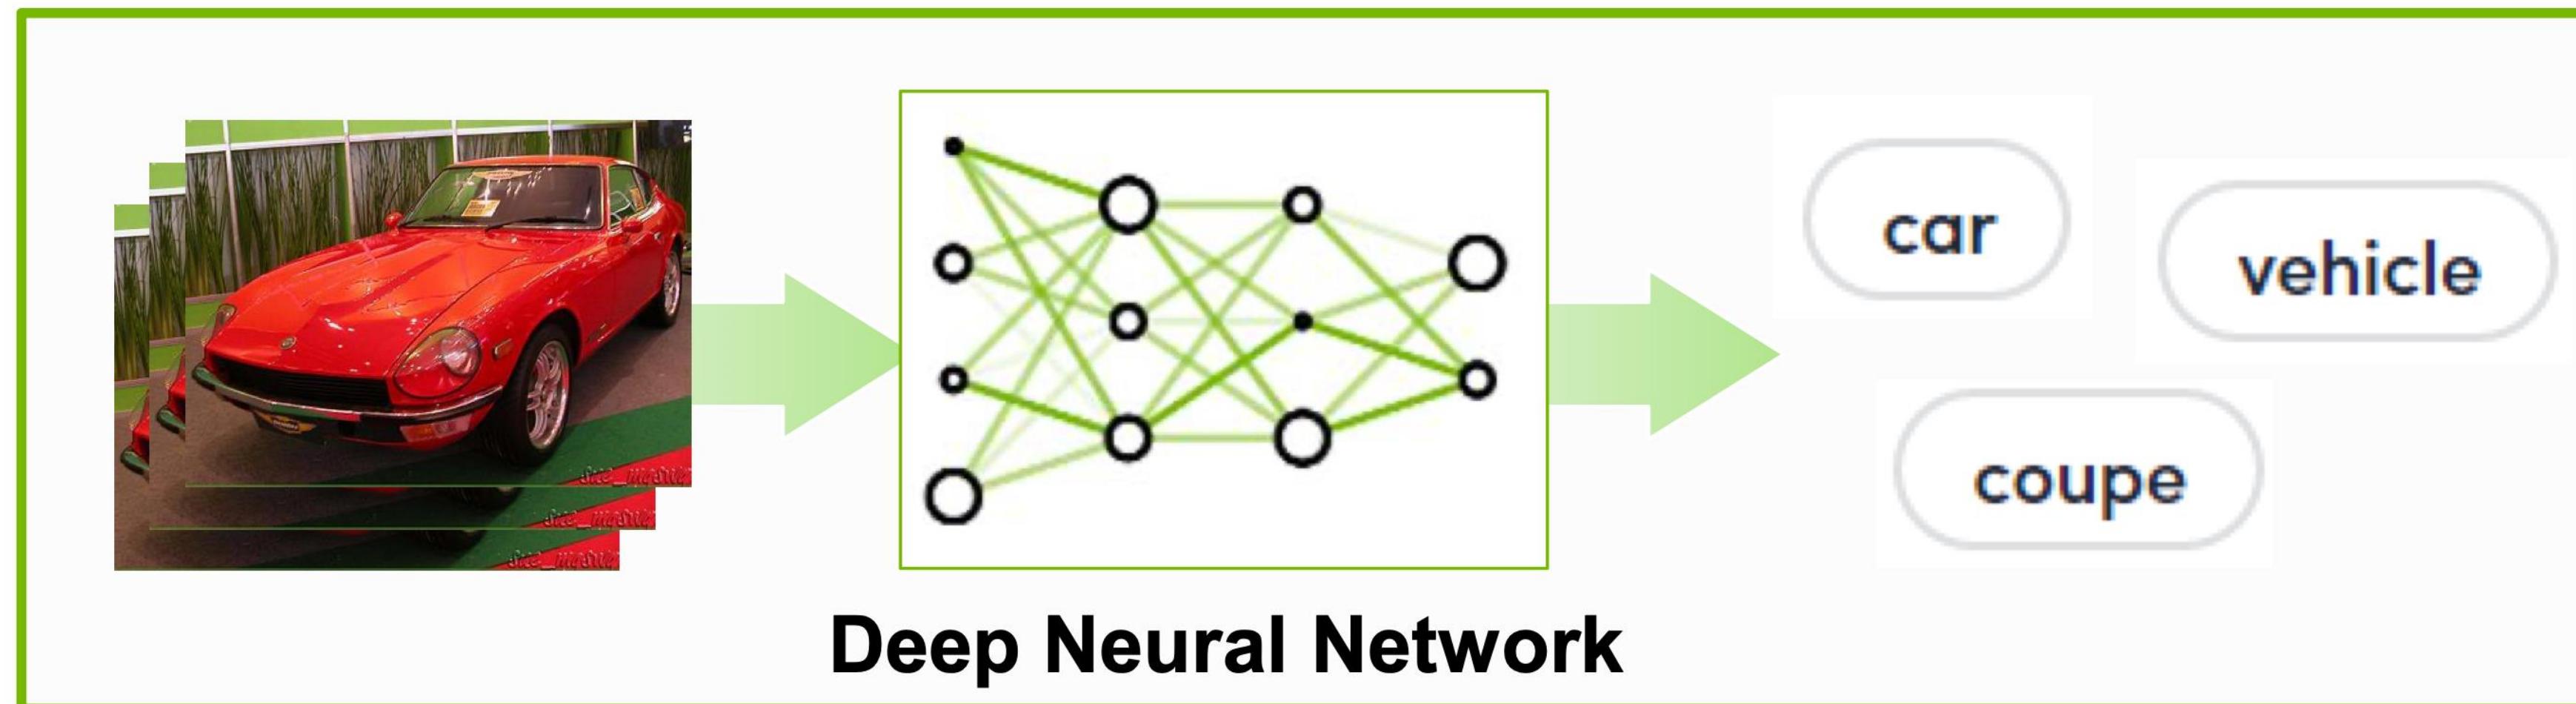
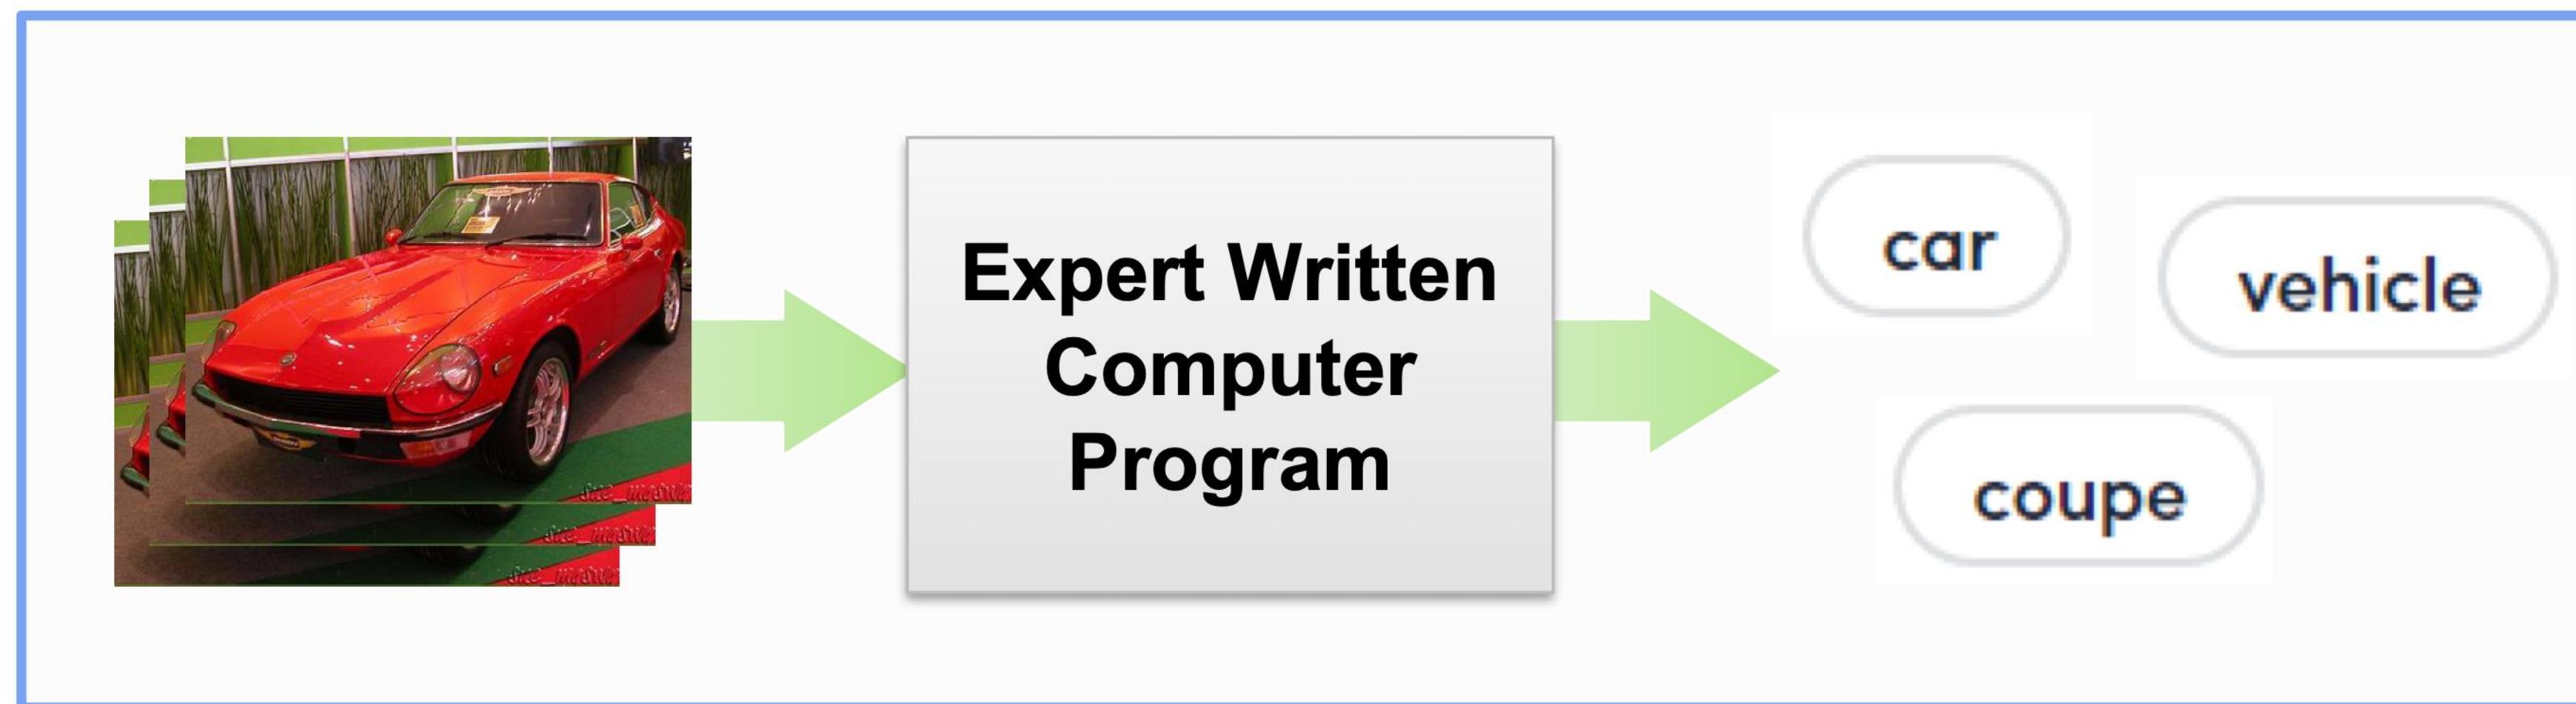


### Traditional Approach

- Requires domain experts
- Time consuming
- Error prone
- Not scalable to new problems

# A NEW COMPUTING MODEL

## Algorithms that Learn from Examples



# Lifecycle of a computer vision DL project

1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

# Lifecycle of a computer vision DL project

1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

**Supervised Learning**  
(We know the right answer)



## **Classification**

CAT

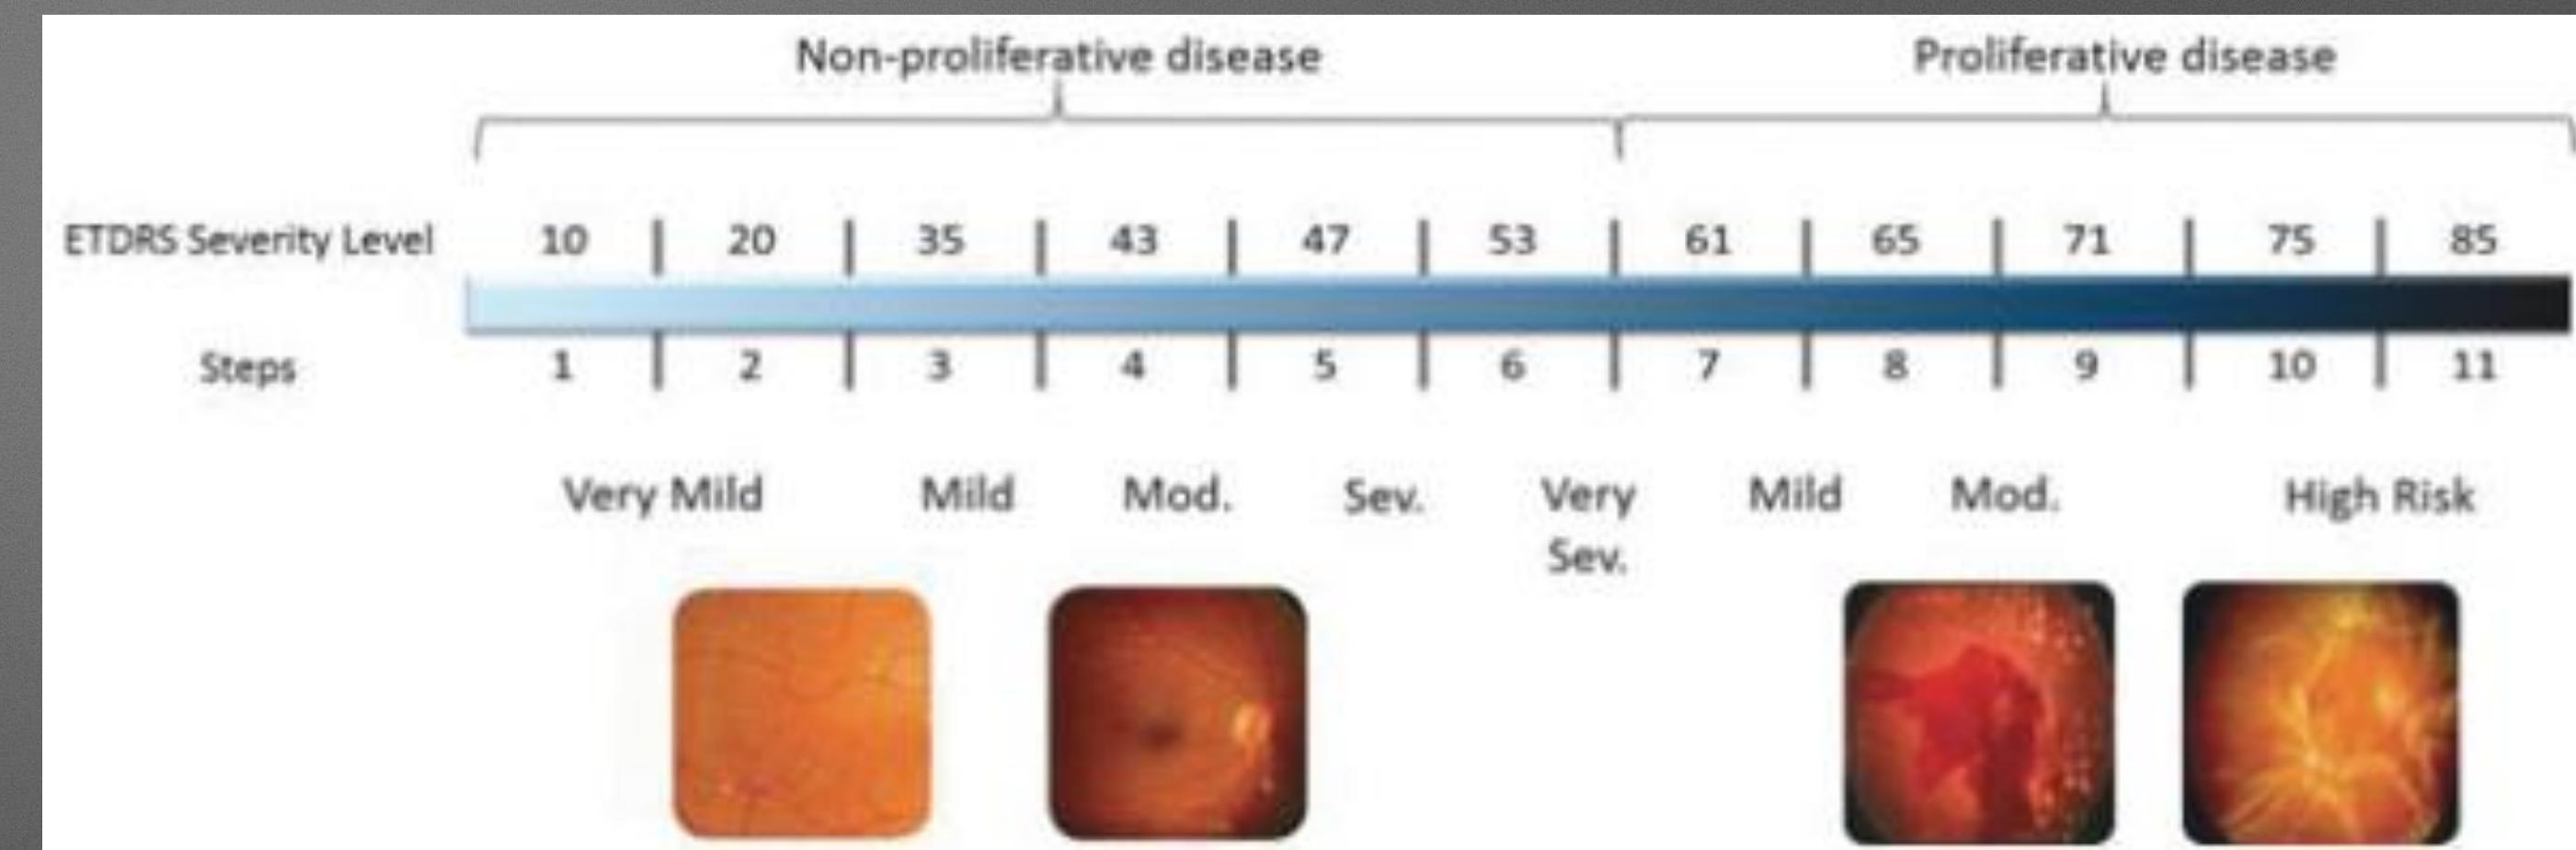
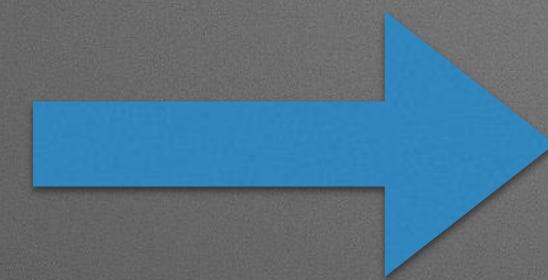
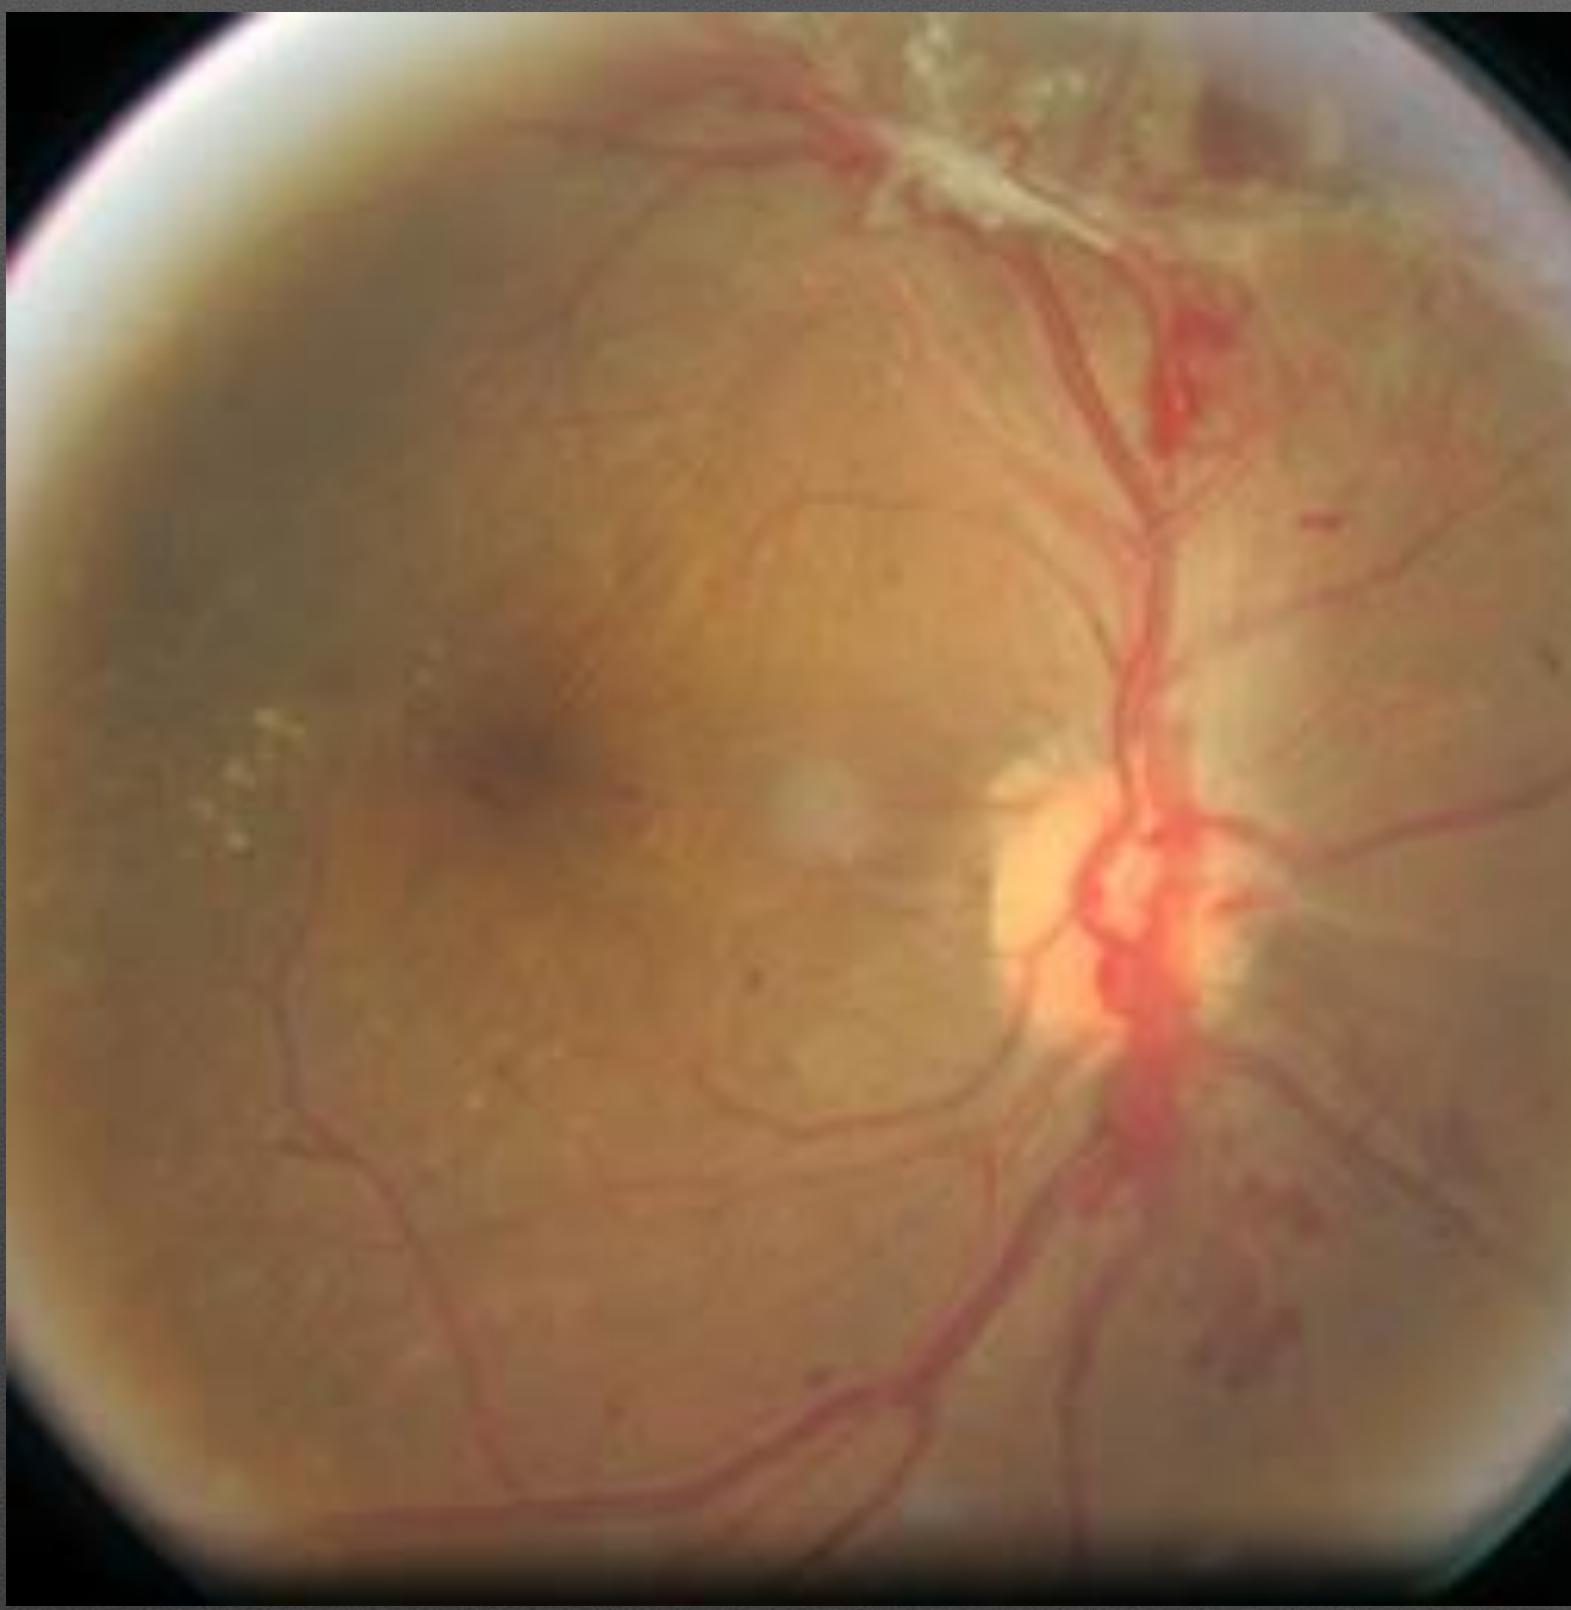


No Diabetic Retinopathy  
Mild NPDR  
Moderate NPDR  
Severe NPDR  
PDR

Given an image, output one of these 5 severities.

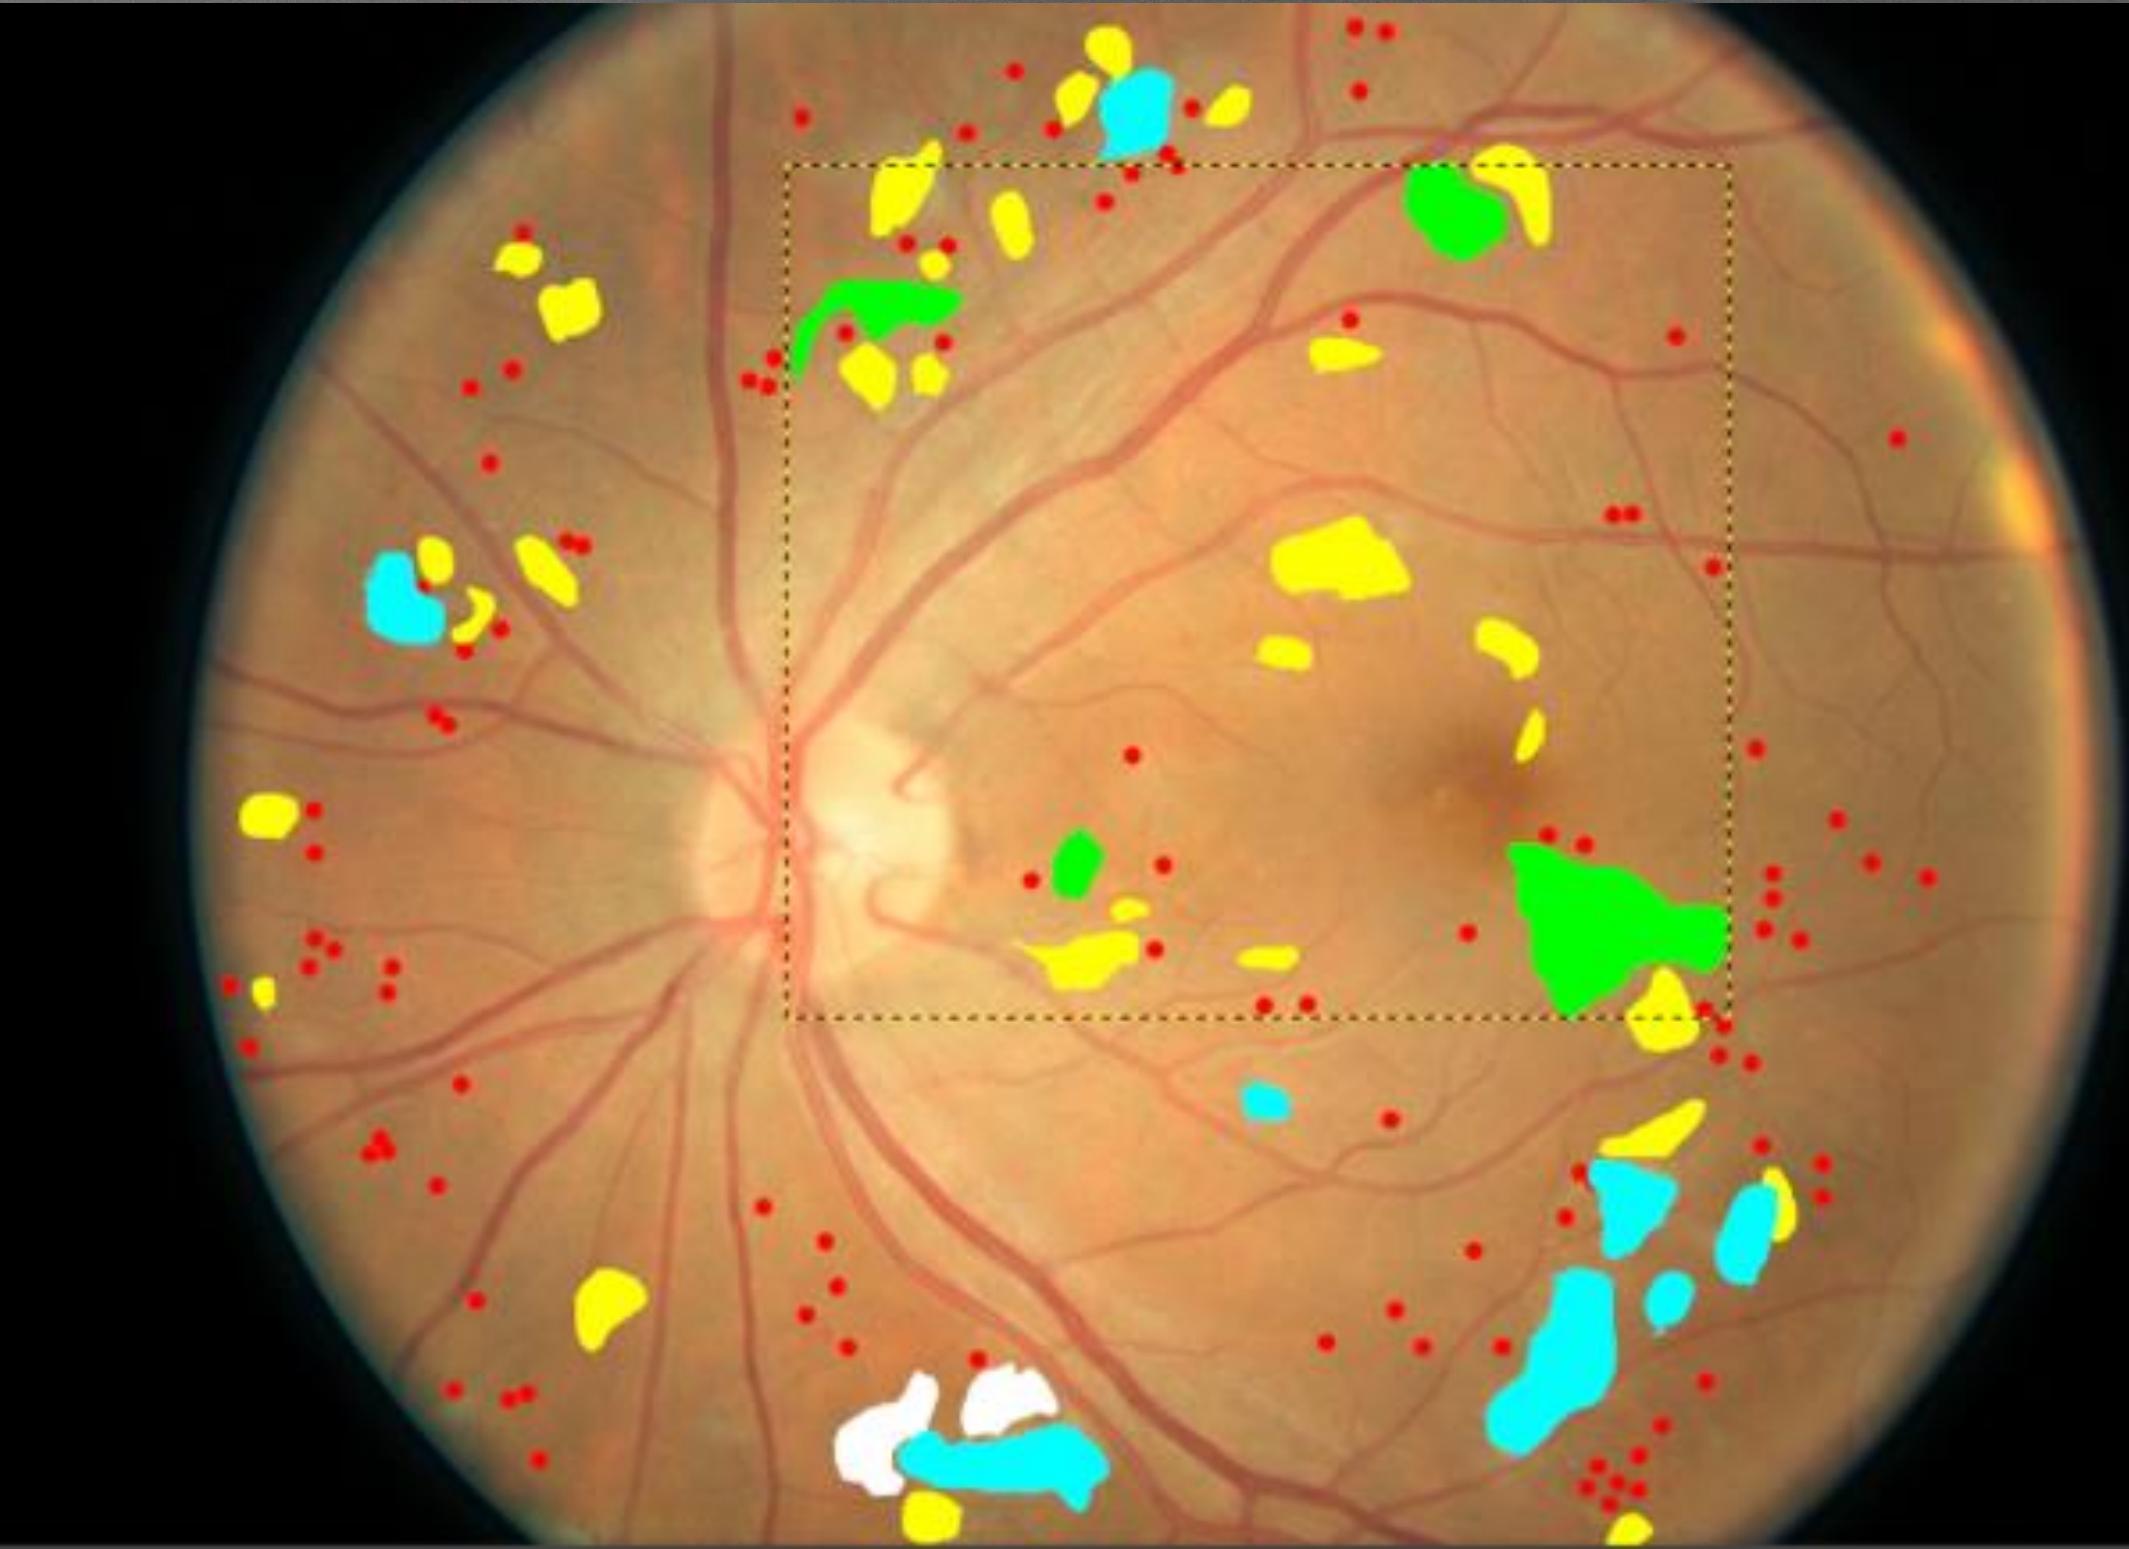
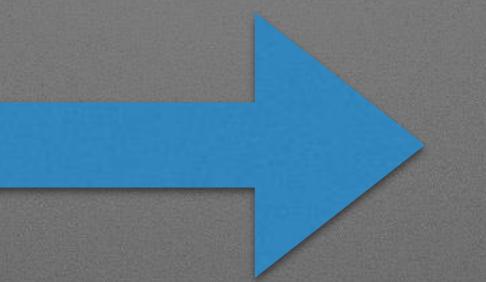
Classification





Given an image, output the ETDRS severity level (0-100).

Regression

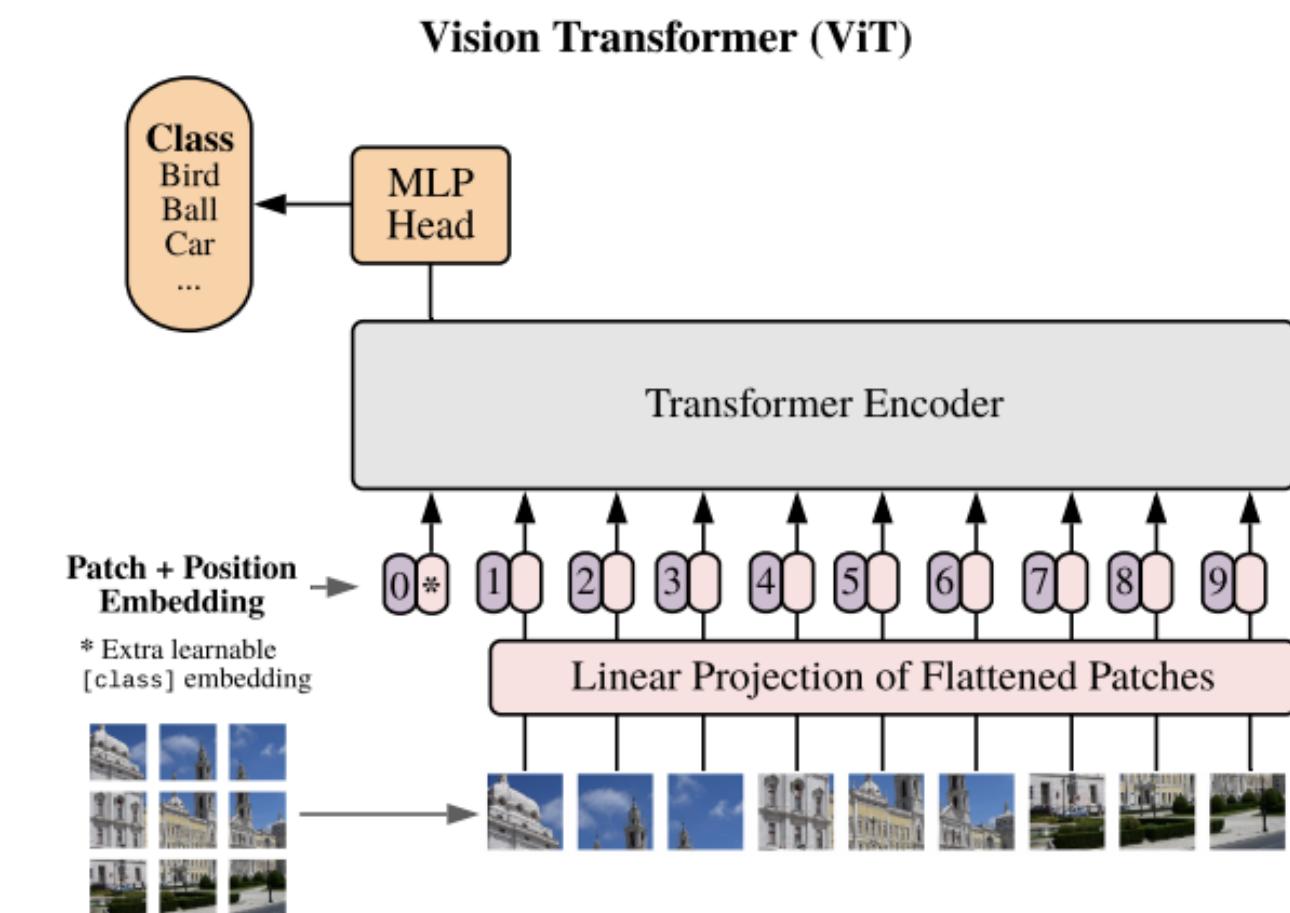
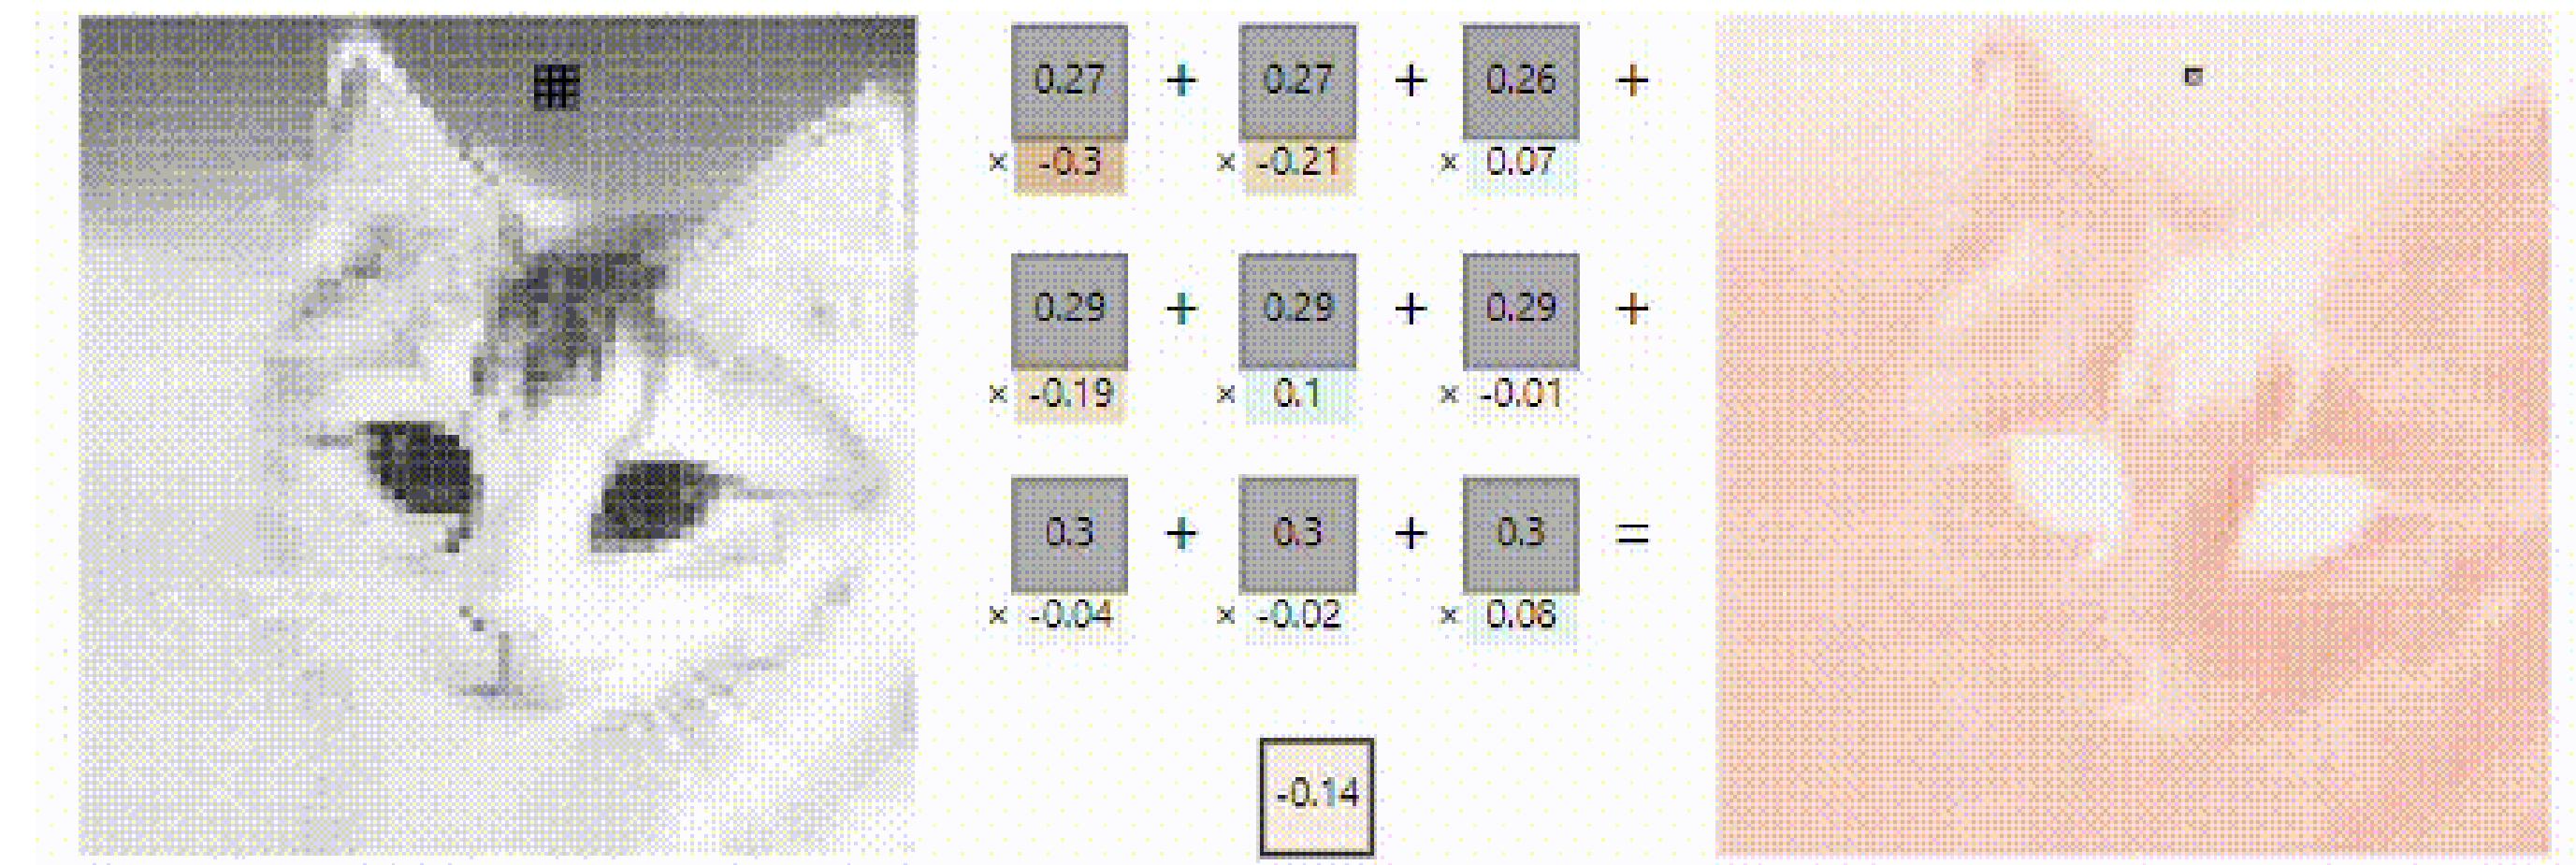


Given an image, output the locations of MA, HEx, CWS, and NV.

Feature Segmentation

# Model choices (Classification)

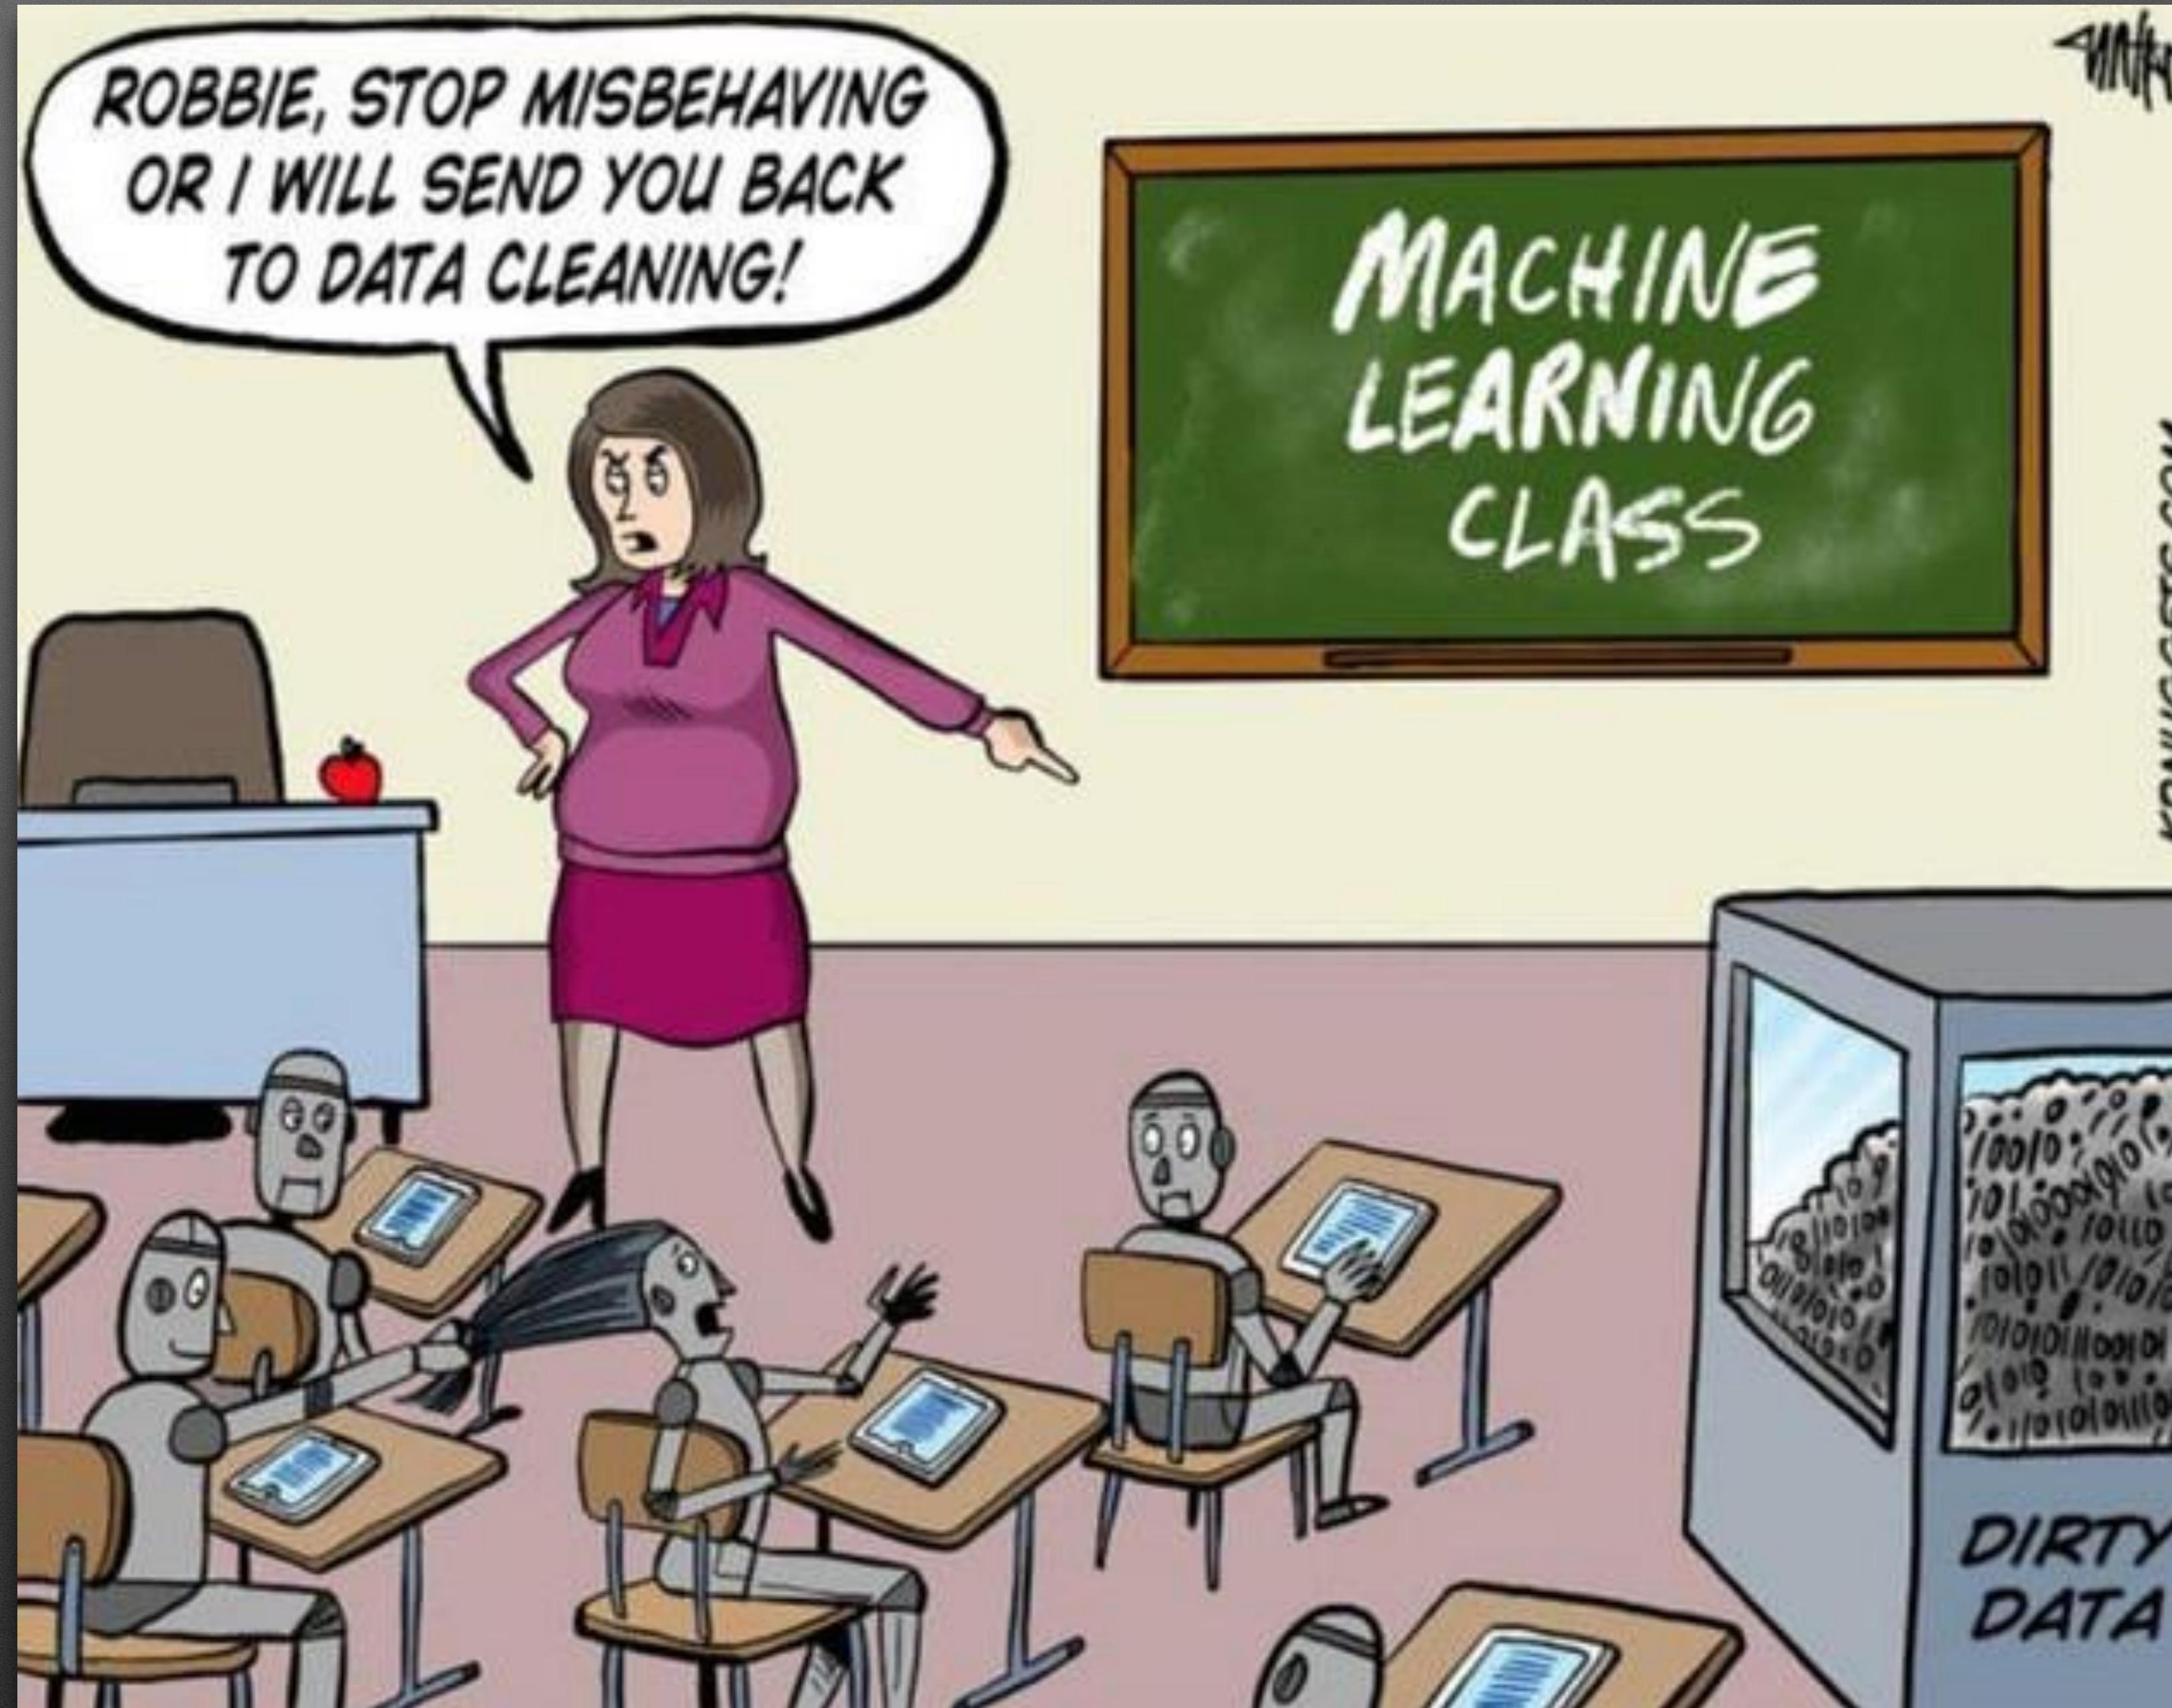
- Convolutional Models
  - VGG series
  - Resnet series
  - EfficientNet series
- Vision Transformer Models
  - Vit series
- Initialization:
  - Random weights
  - Imagenet pretrained weights
  - Foundation model weights



# Lifecycle of a computer vision DL project

1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

Data  
Cleaning =  
99% of the  
work



# Preparing AI-Ready Data

## OMOP

- Demographics
- Survey data
- Physical assessment data
- Medications
- Blood and urine lab values
- MOCA
- Vision testing

## DICOM

- Fundus photography
- OCT
- OCTA
- FLIO

## Waveform Database (WFDB)

- ECG

## mHealth

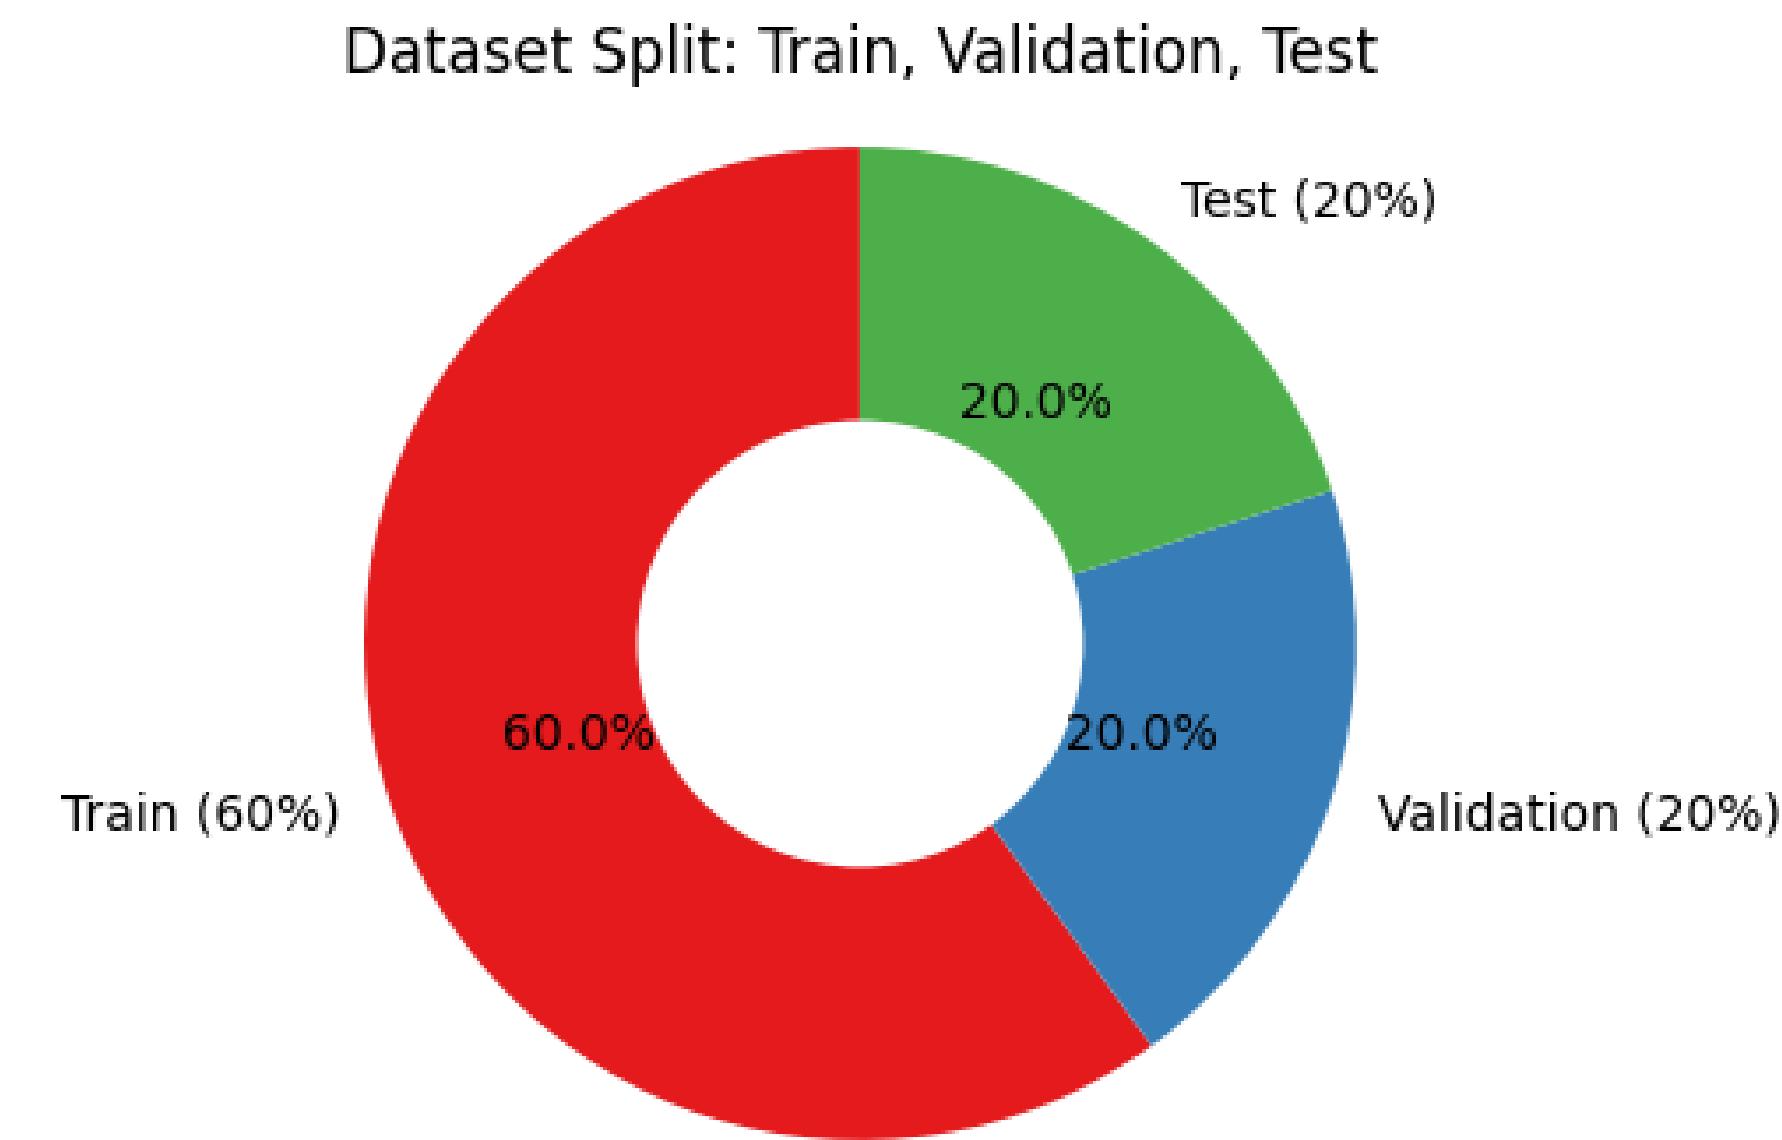
- Physical activity monitor
- Continuous glucose monitor

## Earth Science Data Specification

- Environmental sensor

# Partitioning data

- Training Data (60%)
  - Used to adjust the model to learn from the data
- Validation Data (20%)
  - Used to periodically check how well the model is doing.
  - Model weights are NOT adjusted
  - Used during training session and for hyperparameter tuning.
- Test Data (20%)
  - Used to evaluate the final performance of the final model (ideally once at the end)
  - Model weights are NOT adjusted
- Unit of partition: person



## Suggested split

The suggested split for training, validating, and testing AI/ML models is included in the participants.tsv file that will be included in the dataset. A summary is provided in the table below.

	Train				Val				Test				Total			
	Hispanic	Asian	Black	White	Hispanic	Asian	Black	White	Hispanic	Asian	Black	White	Hispanic	Asian	Black	White
Race/ethnicity (count)	144	167	211	225	40	40	40	40	40	40	40	40	224	247	291	305
Sex (count)	Male	Female			Male	Female			Male	Female			Male	Female		
Diabetes status (count)	No DM	Lifestyle	Oral	Insulin	No DM	Lifestyle	Oral	Insulin	No DM	Lifestyle	Oral	Insulin	No DM	Lifestyle	Oral	Insulin
Mean age (years ± sd)	60.3 ± 11.13	60.2 ± 10.5			60.4 ± 11.0			60.3 ± 11.1								
Total	747	160			160			1067								

# Lifecycle of a computer vision DL project

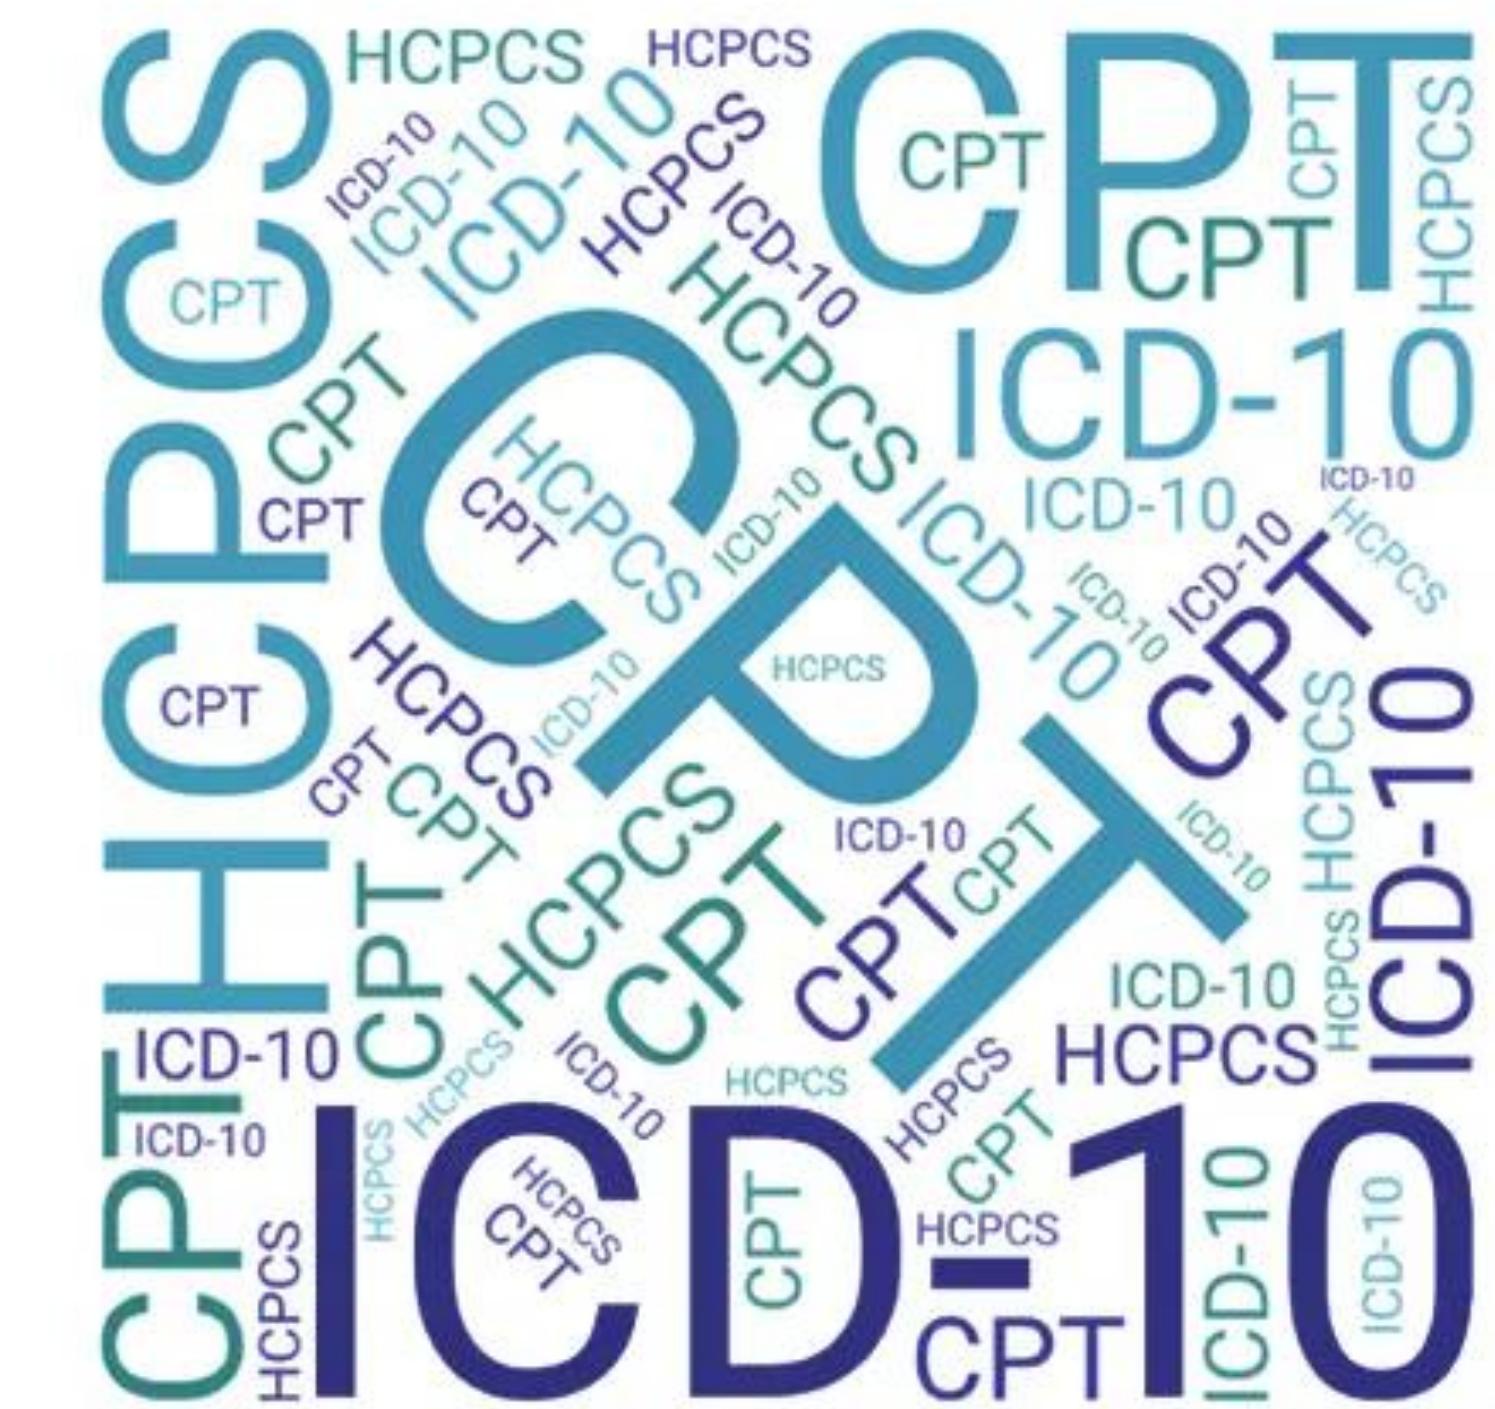
1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
- 3. Torture humans to label the data (output)**
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

# Tradeoffs

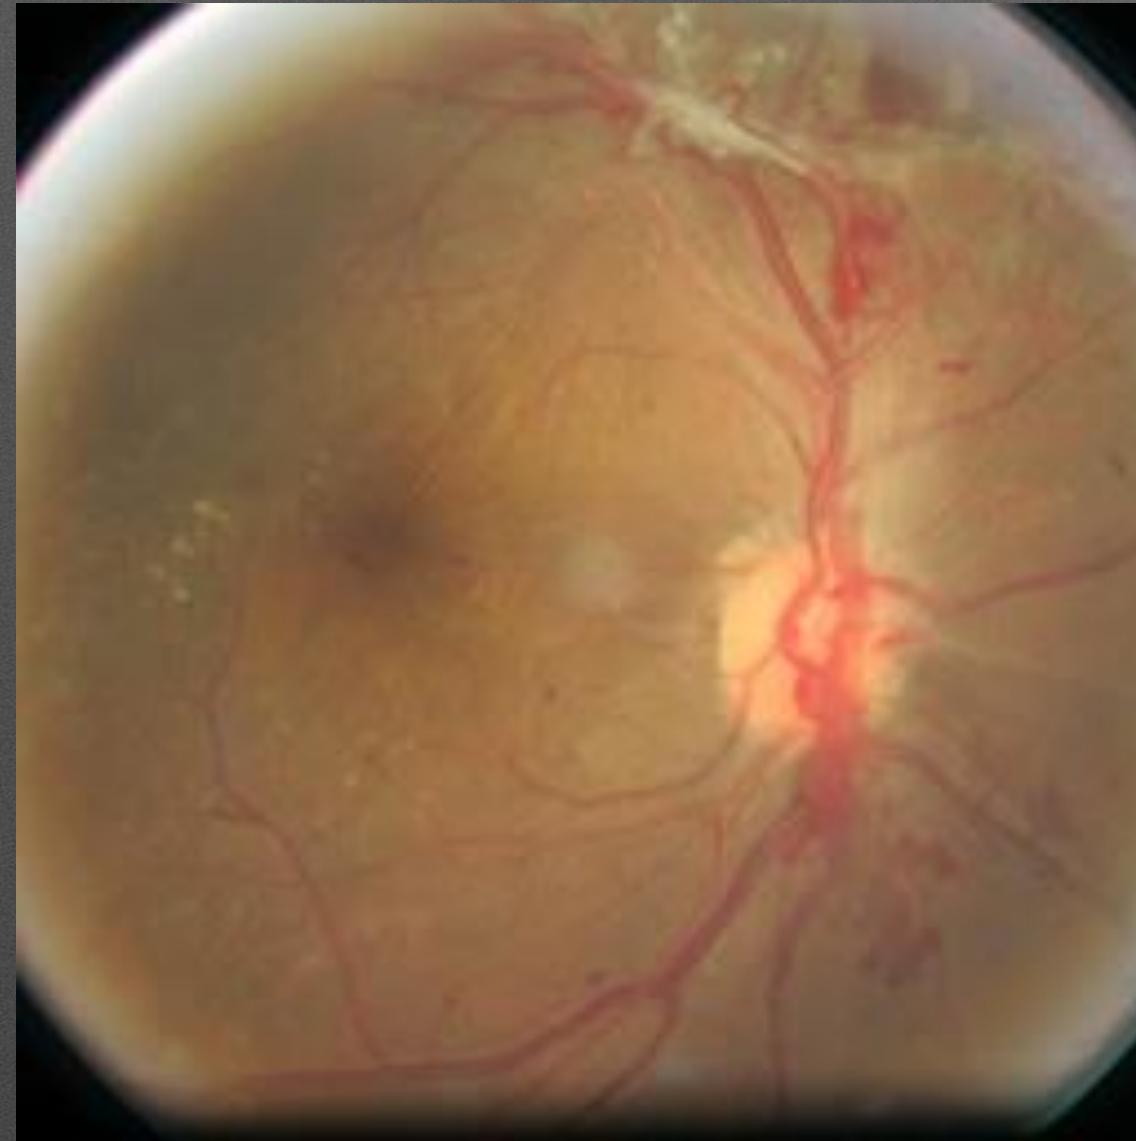
- Ideal setup:
  - 2 expert graders all grade the same images
  - A third arbiter (even more of an expert), masked to the identities of the graders, regrades and provides a final grade
  - Con: **very labor intensive**
- Most common setup:
  - Have  $n$  graders each grade on their own
  - Not everyone is the same level of expertise
  - Not everyone grades the same number of images
  - Con: model performance will likely be **a mixed bag**.

# Leverage existing labels

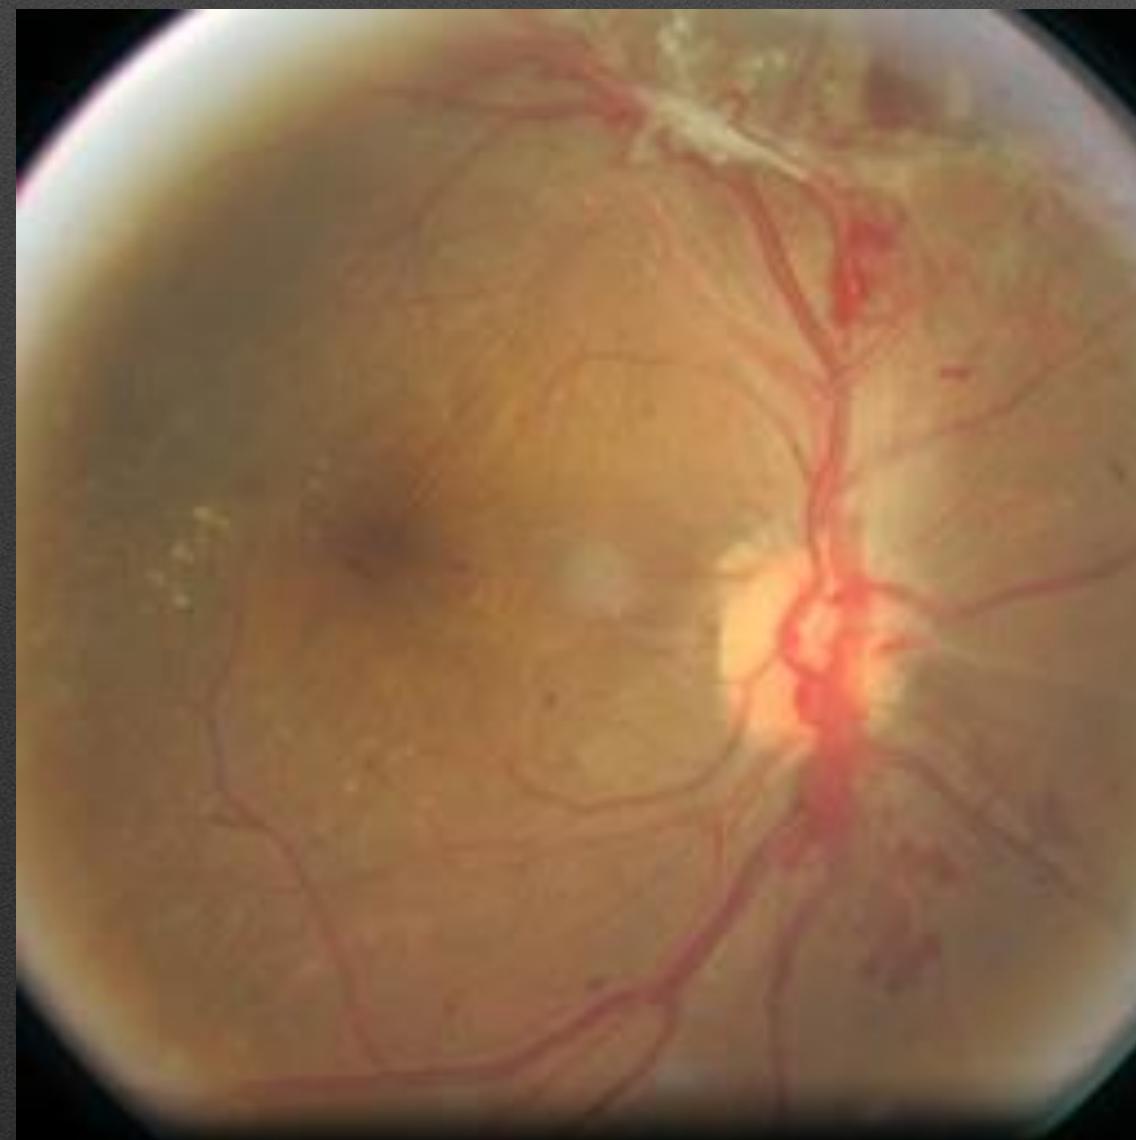
- Derive labels from diagnosis codes, CPT codes, etc. from EHR
  - Noisy and can be **unreliable**
  - CPT procedure codes are better labels as they are true “events”
- Derive labels from lab tests, future events (eg readmission, need for surgery)
  - Can be **acutely altered** in extreme situations



Models will **never** do better than the average  
human error.



Moderate



Severe NPDR

Try to use **objective** ground truth for labels  
(not derived from humans).

# Lifecycle of a computer vision DL project

1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
- 4. Train model (hyperparameters)**
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

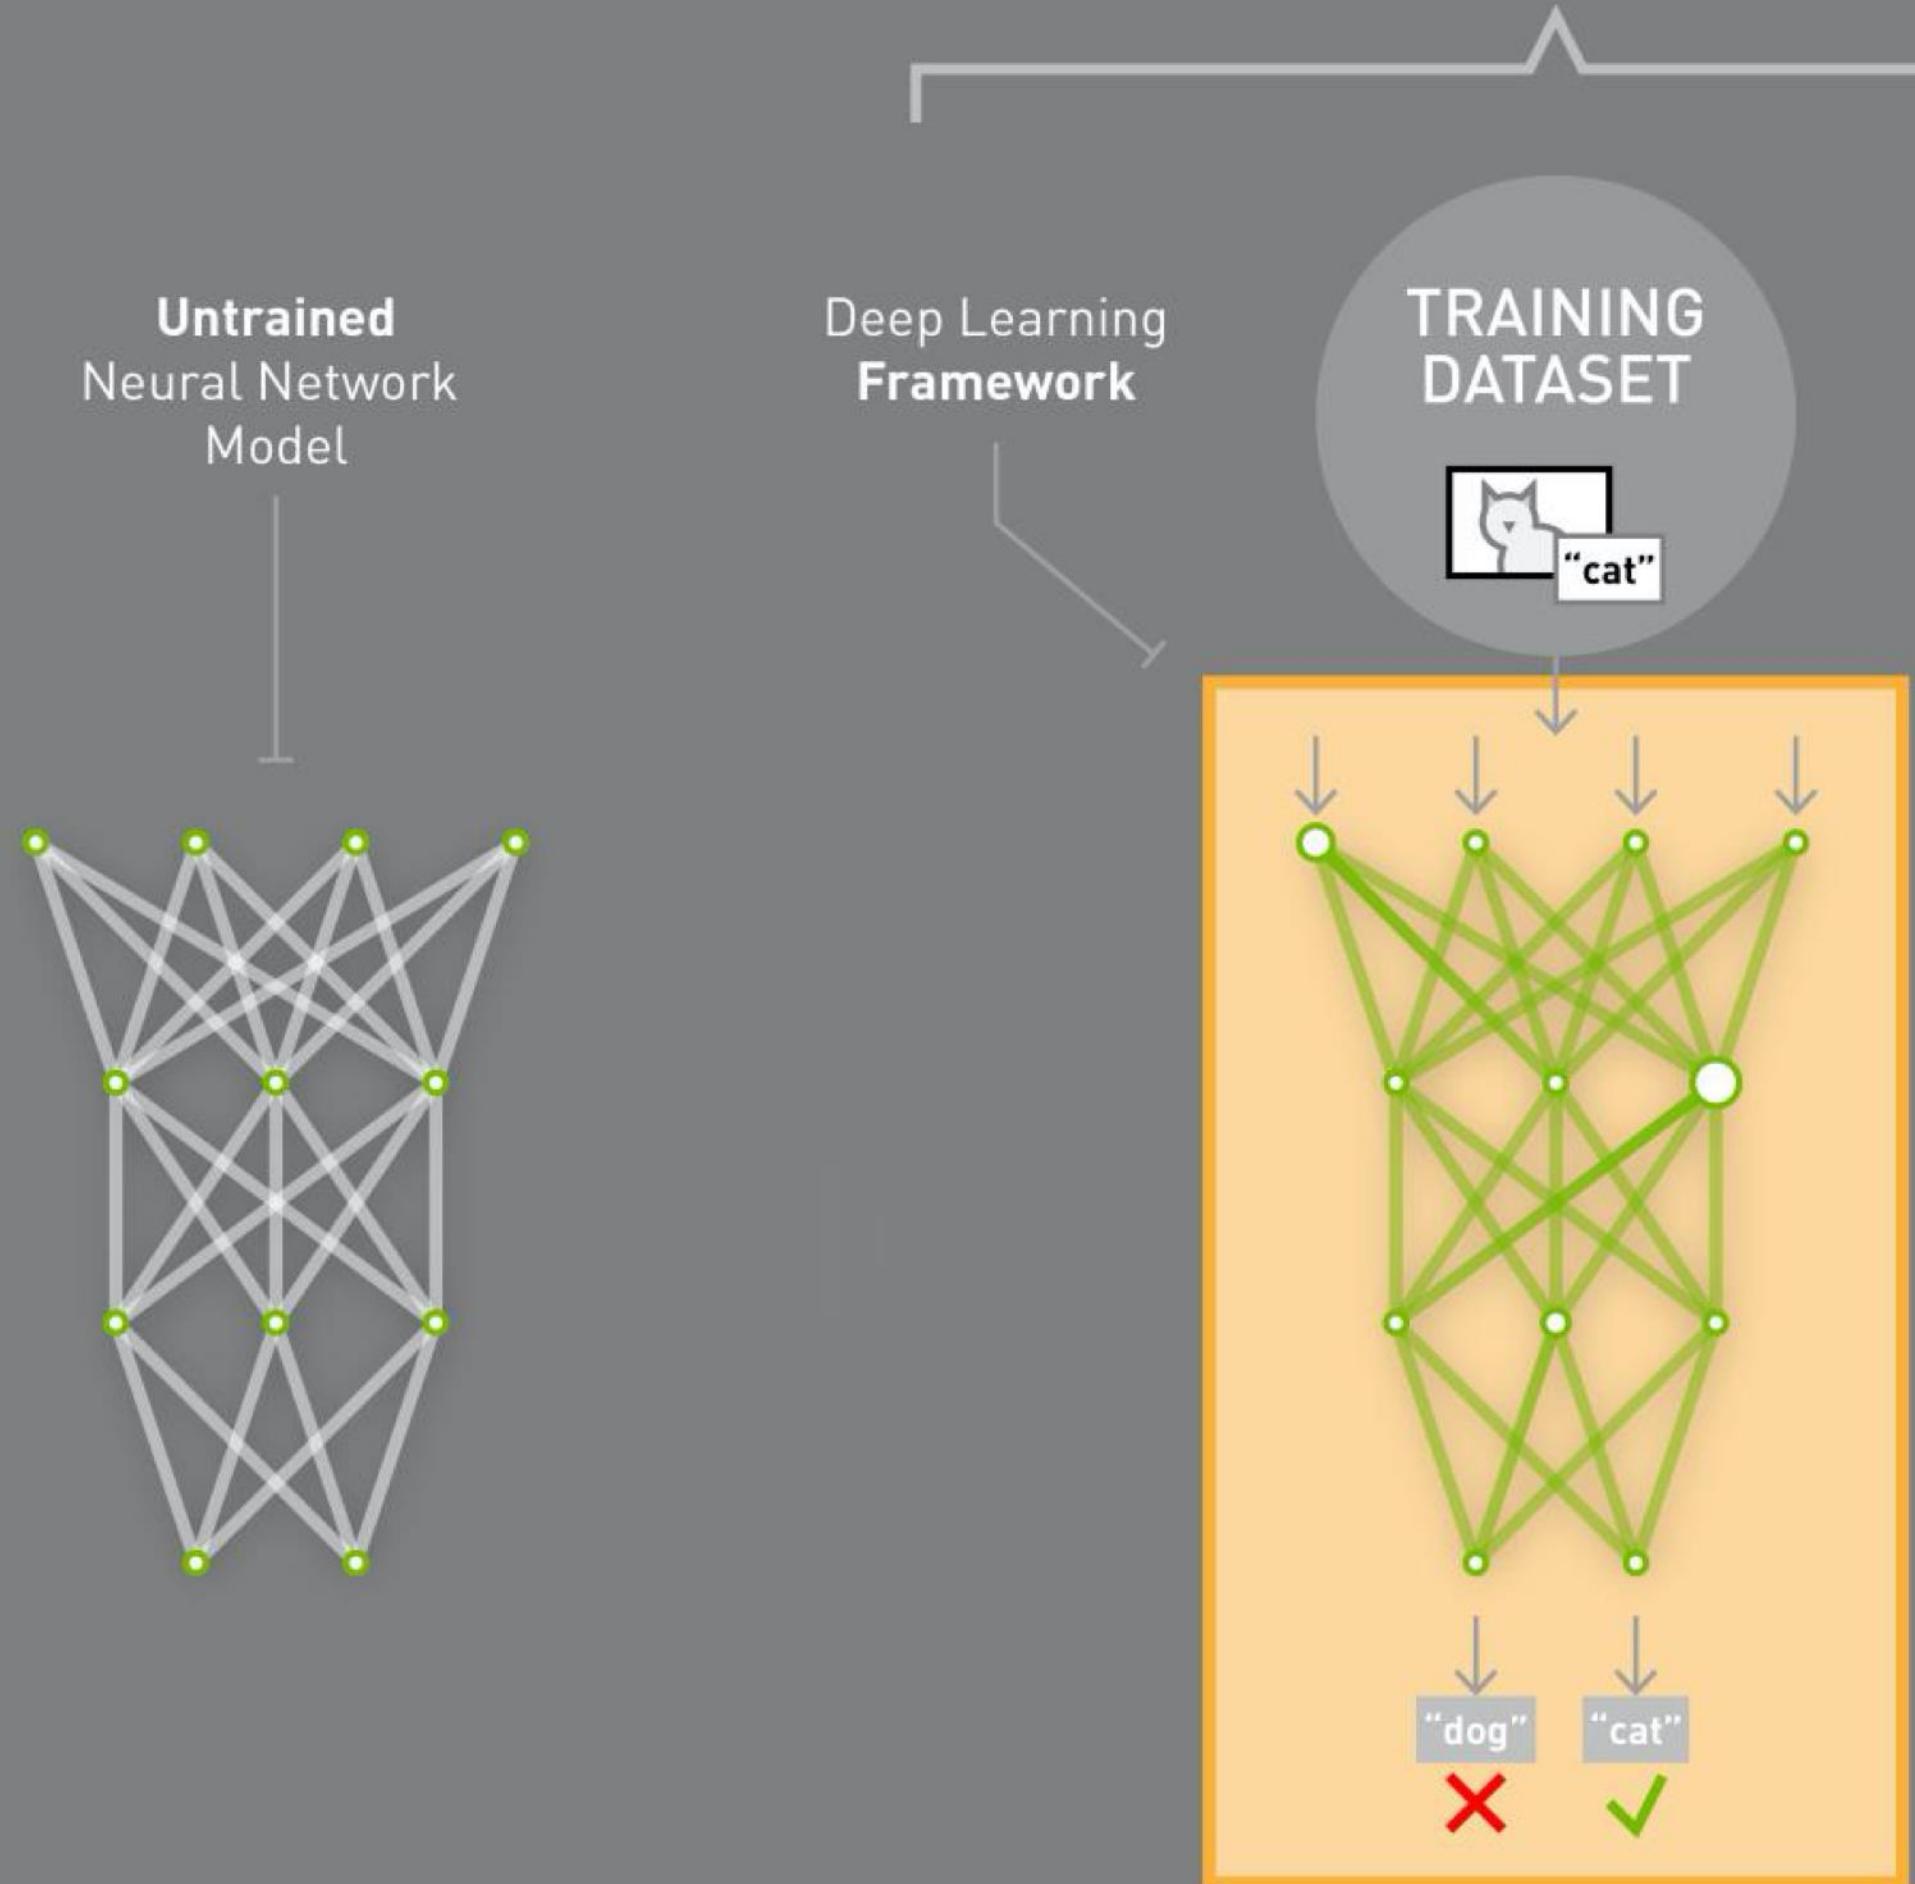
# DEEP LEARNING



# DEEP LEARNING

## TRAINING

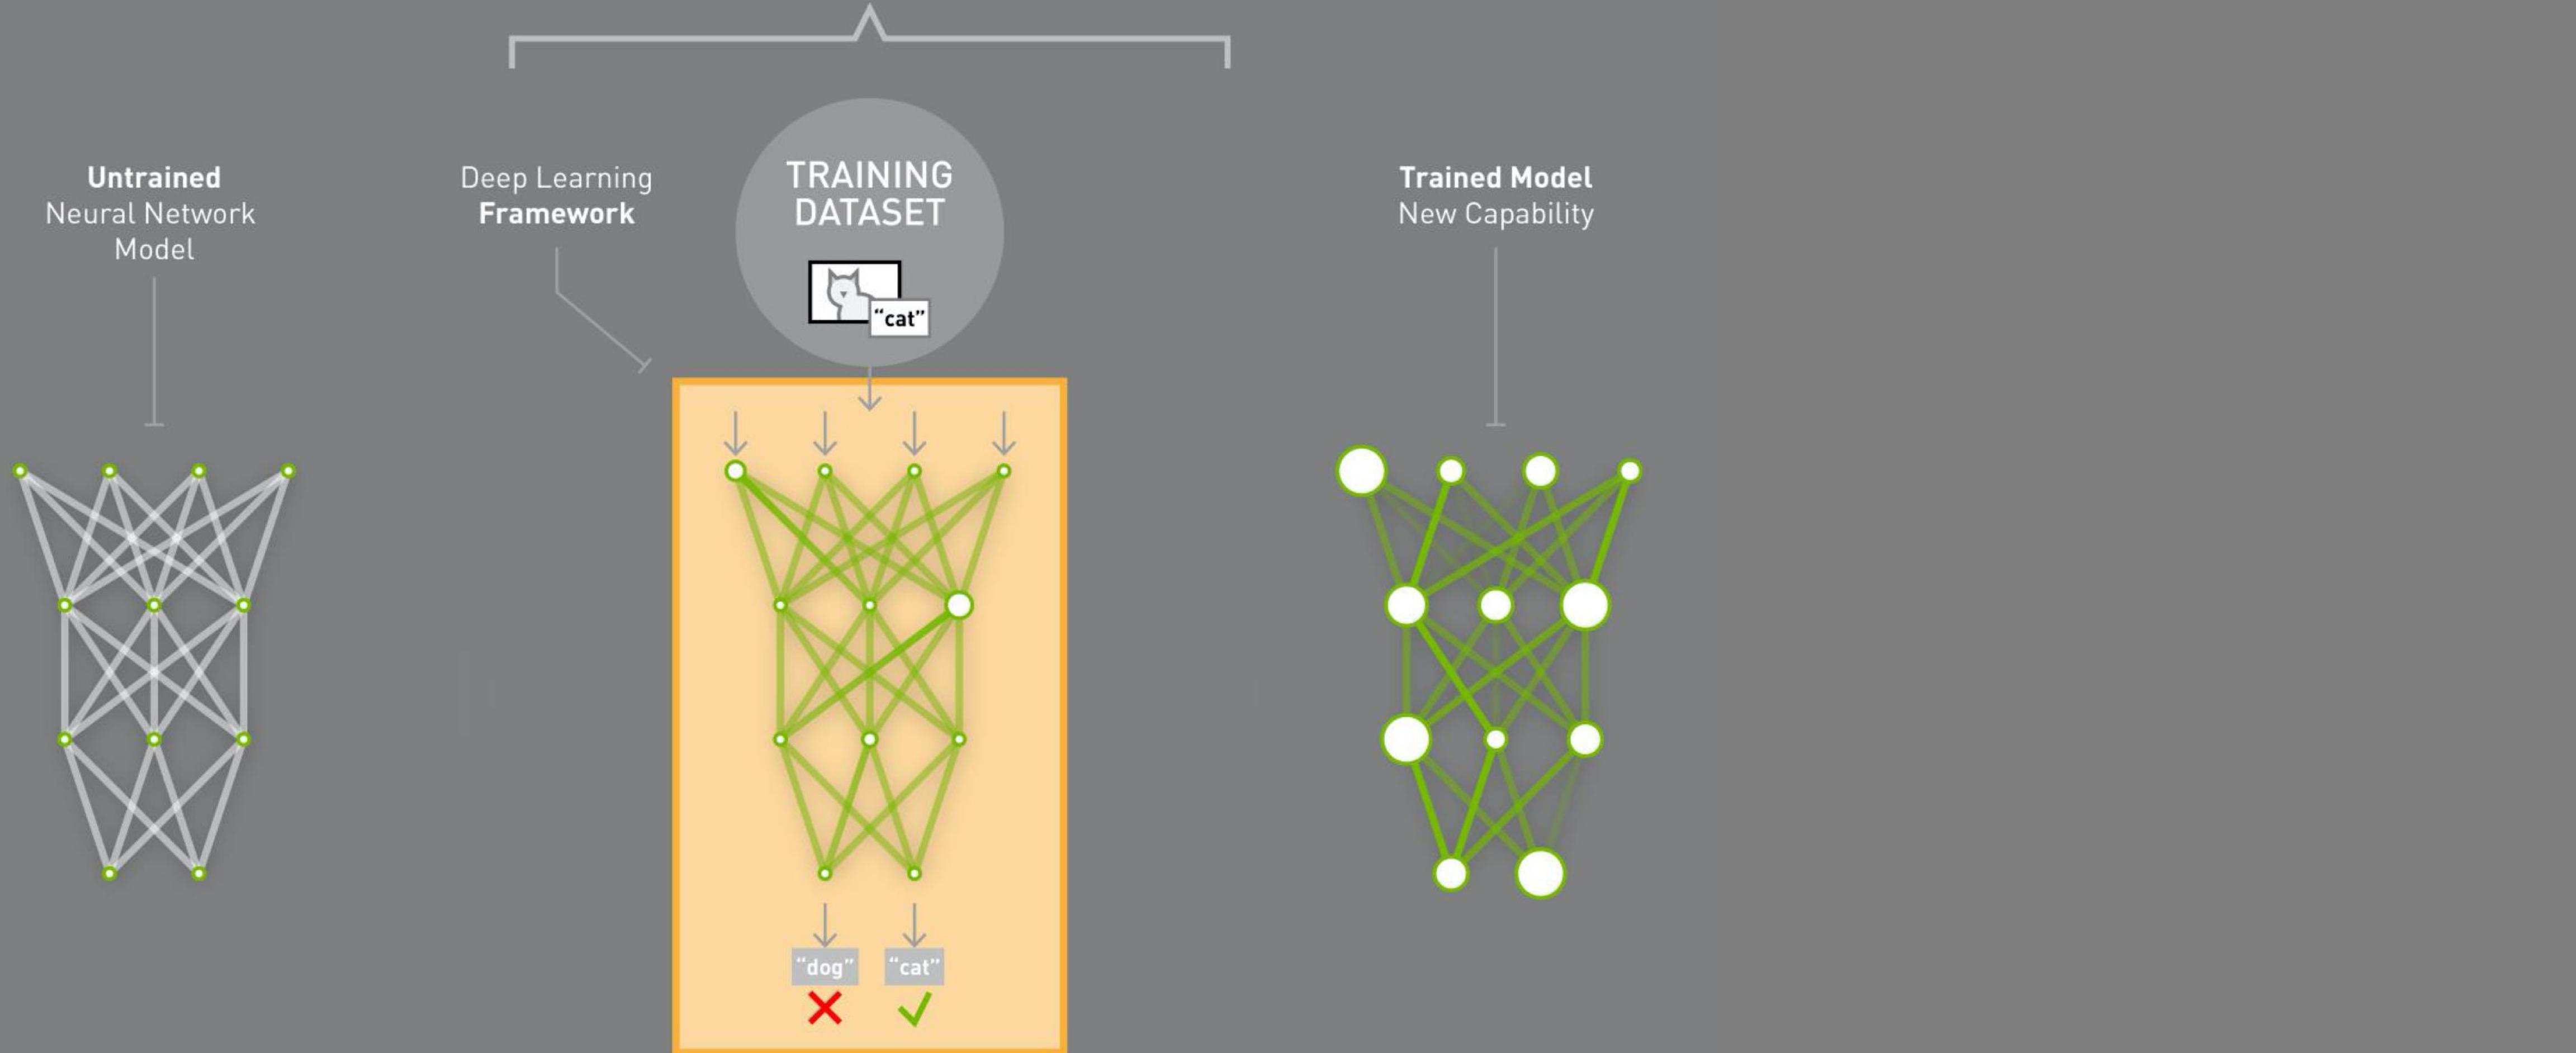
Learning a new capability  
from existing data



# DEEP LEARNING

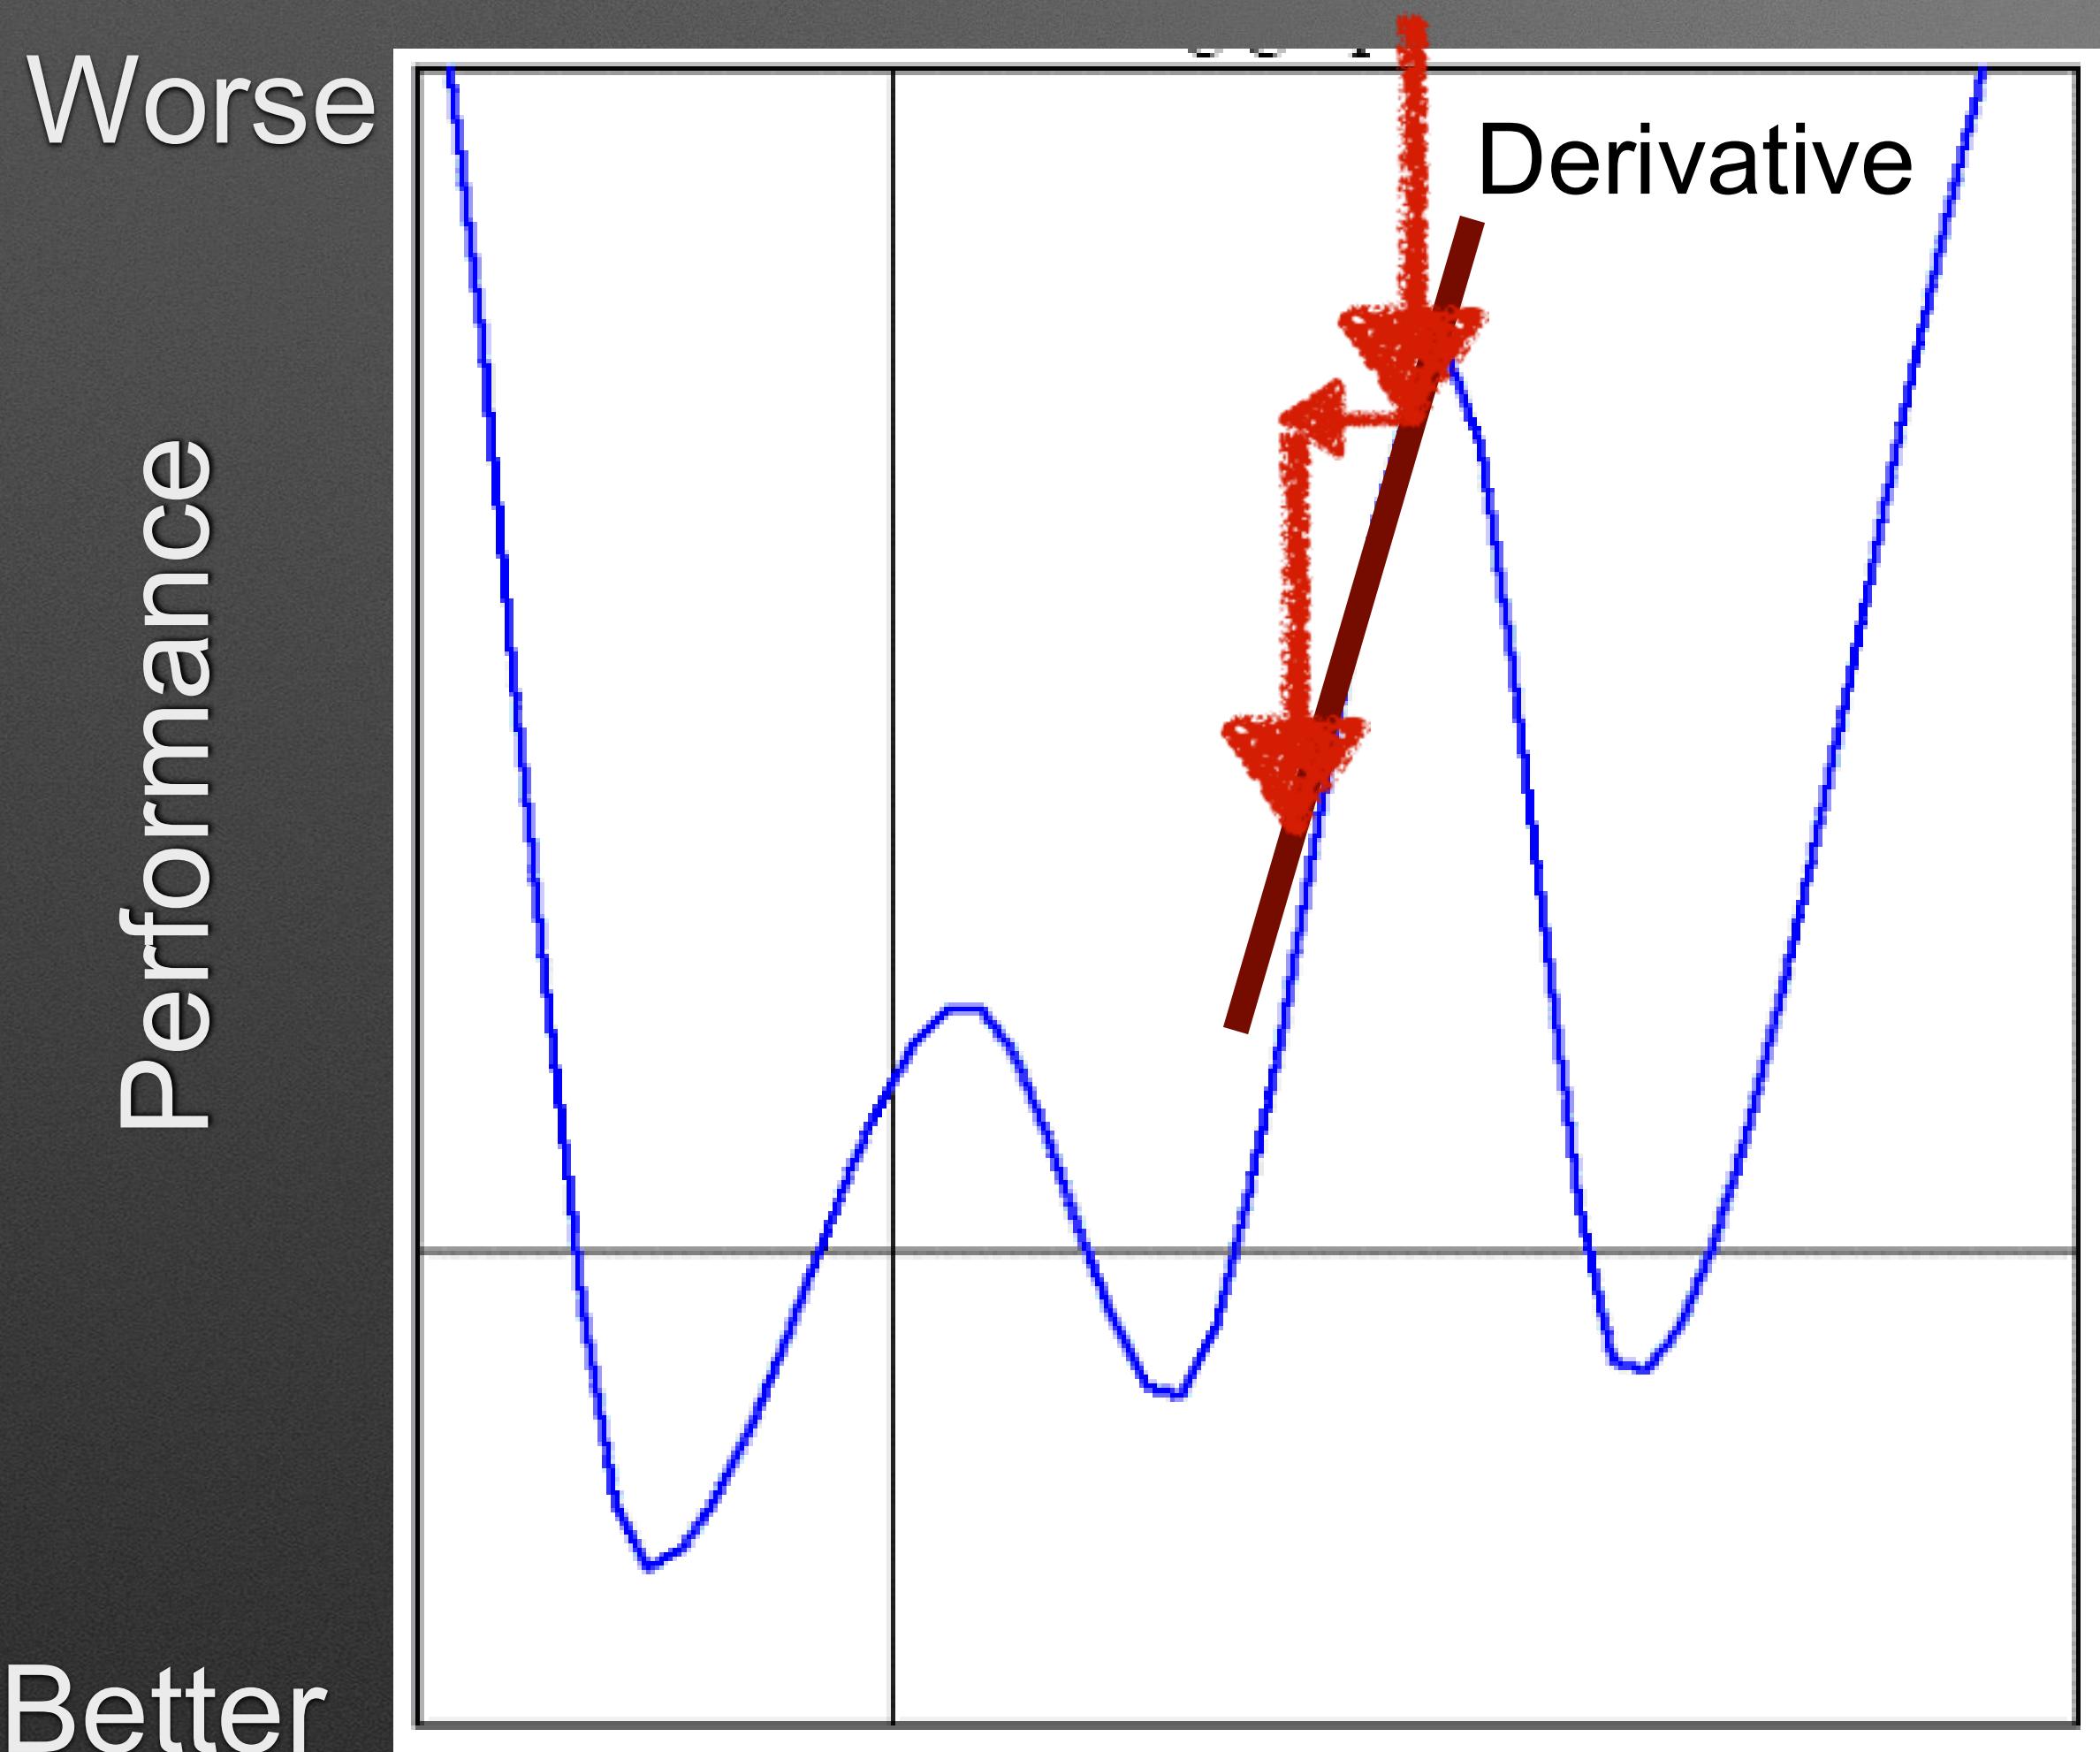
## TRAINING

Learning a new capability  
from existing data



# How does it actually learn?

Random initialization



$$w := w - \eta \sum_{i=1}^n \frac{\nabla Q_i(w)}{n}$$

Random  
initialization

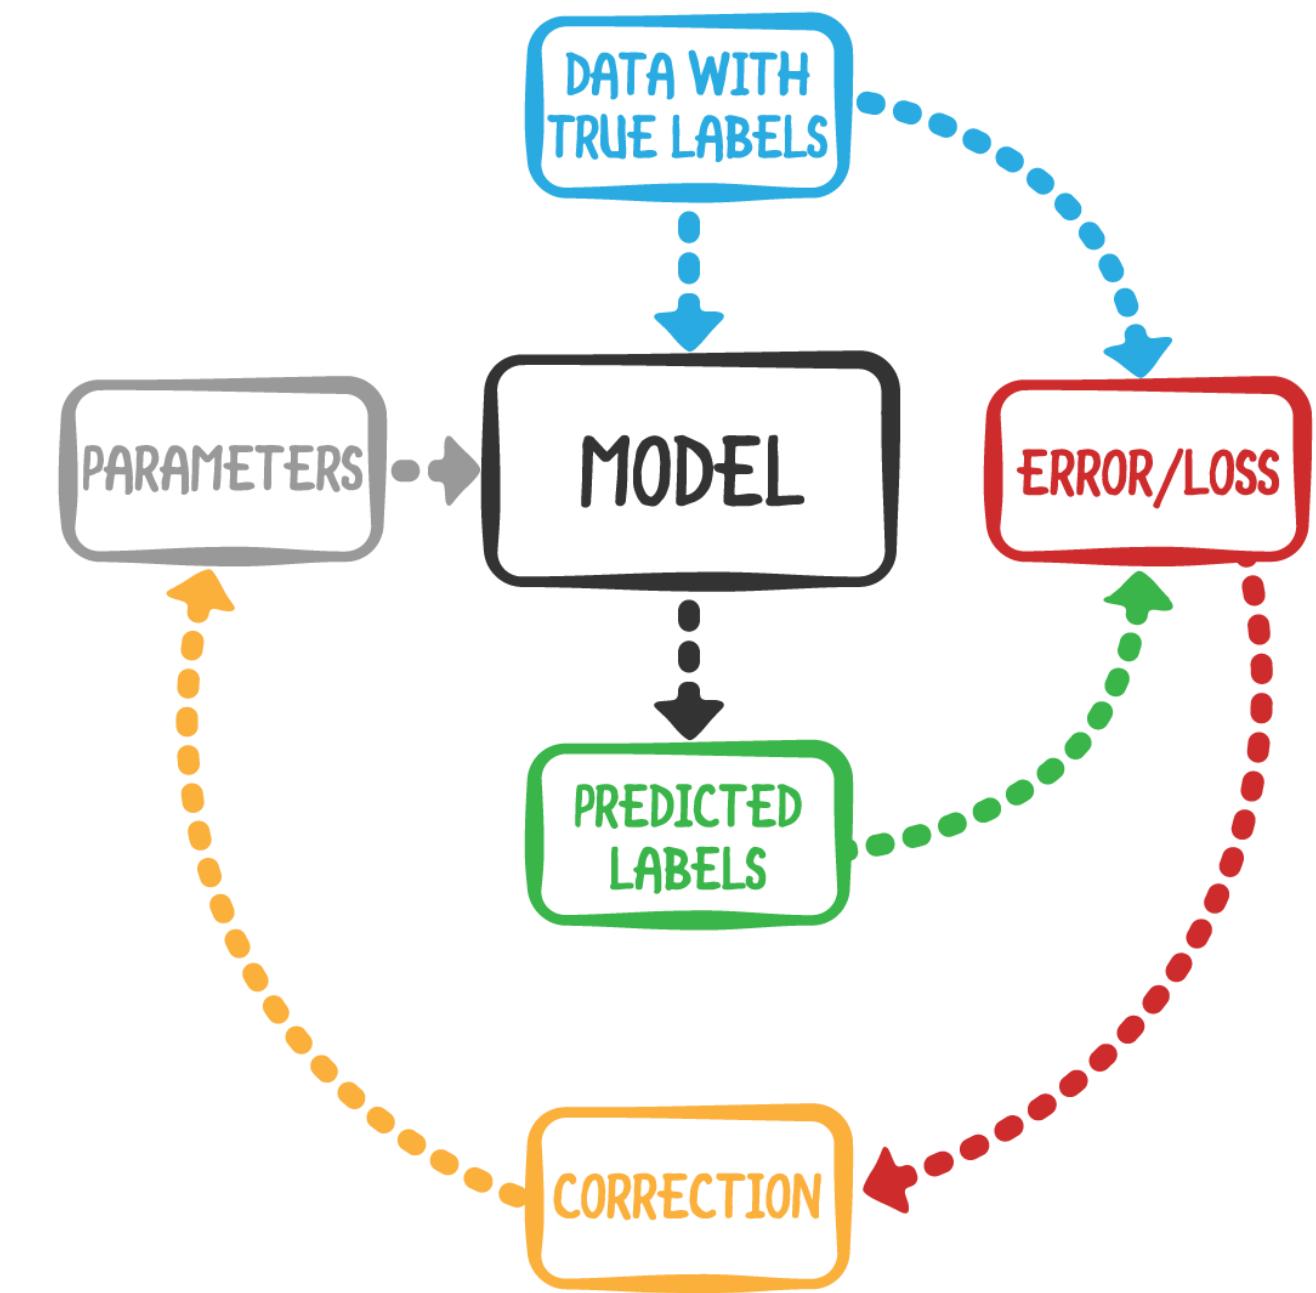
Learning  
rate

Gradient

Stochastic Gradient Descent

# Training loop

- **Epoch 1:**
  - Shuffle training data
  - Divide into small batches
  - For each batch: adjust the model slightly
  - After going through all the training data, freeze the model and test on validation dataset
- **Epoch 2:**
  - Repeat
- **Epoch 3:**
  - Repeat

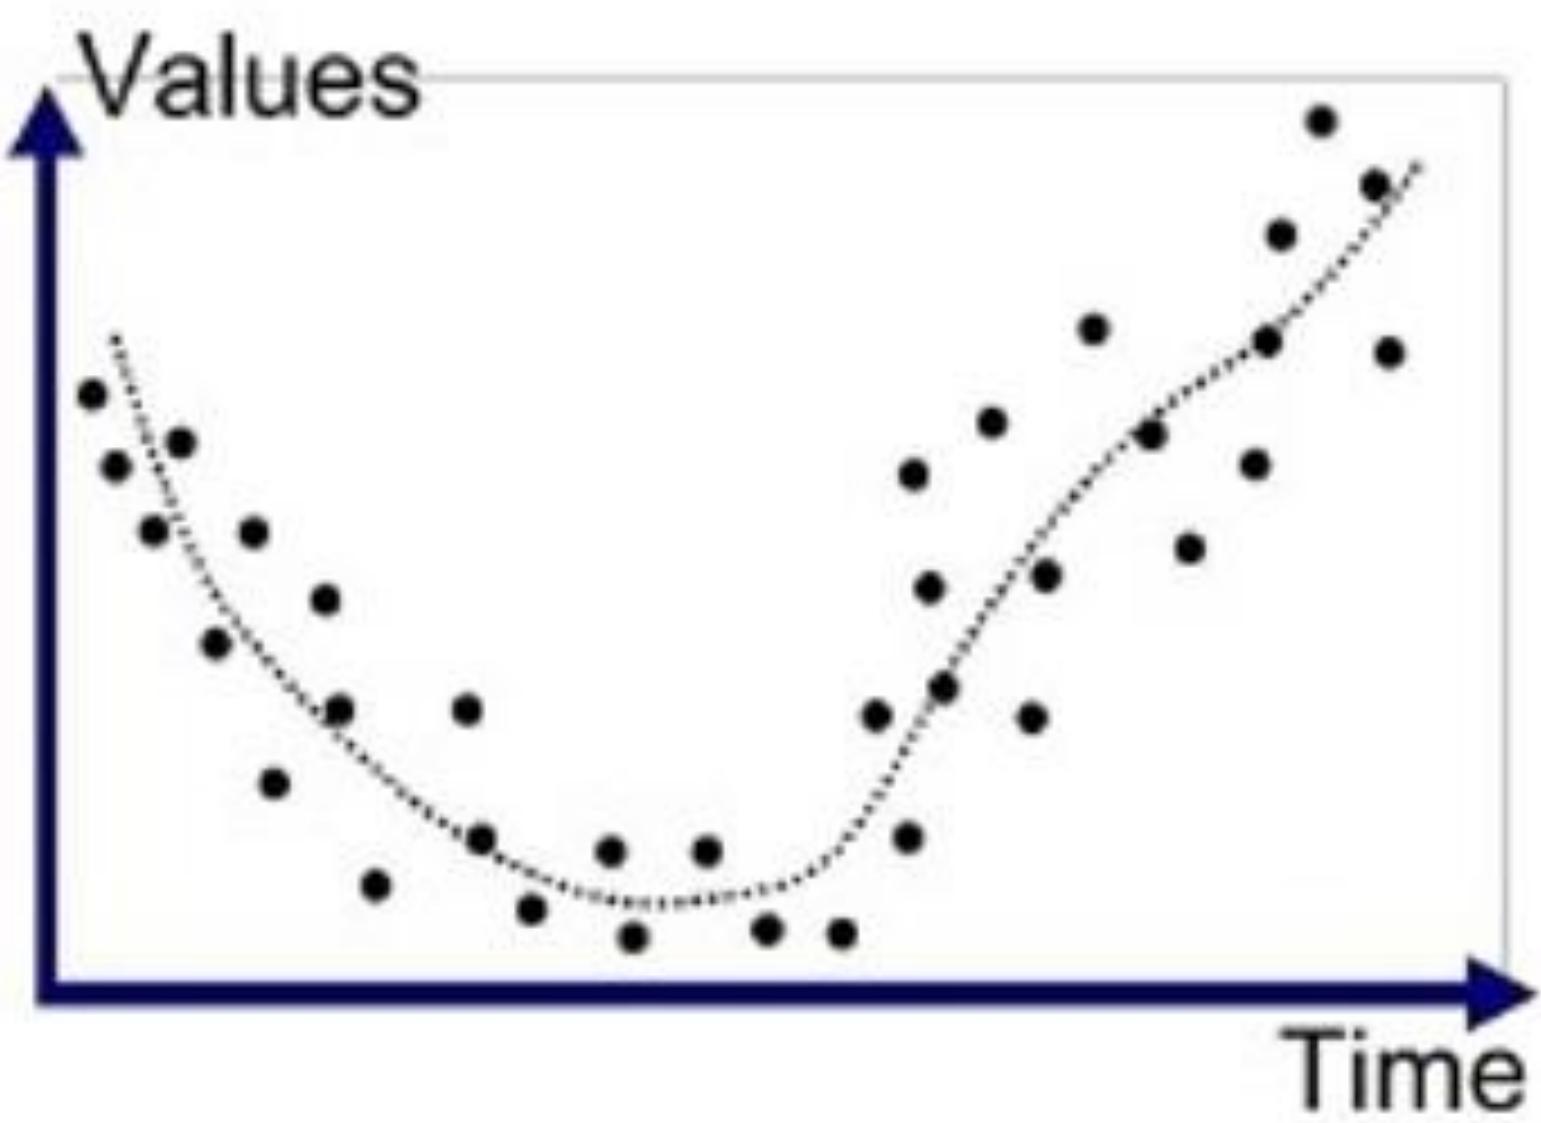


# Lifecycle of a computer vision DL project

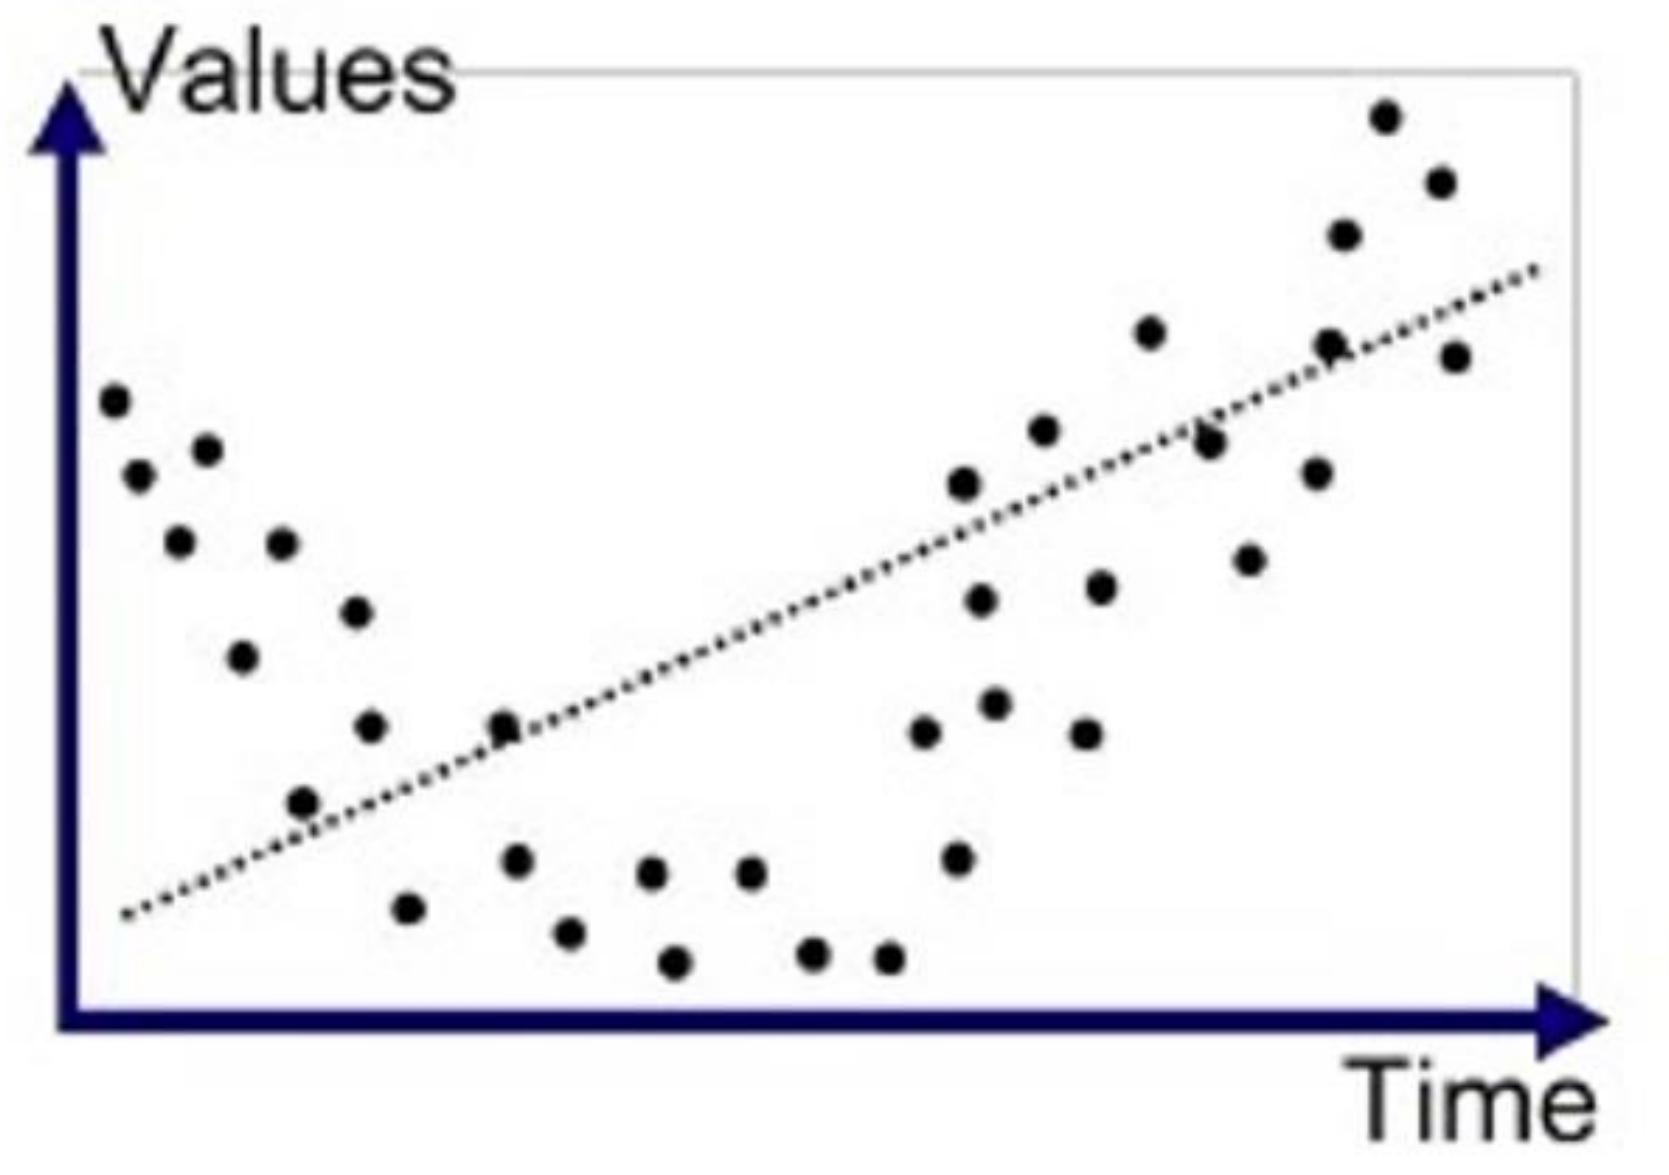
1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

# What are hyperparameters?

- Parameters = how the model will behave for a given input.
- Hyperparameters = things we can adjust that will ultimately lead to better parameters
- Examples: batch size, learning rate, optimizer, number of epochs
- They affect how a training session behaves but not how a deployed model behaves.

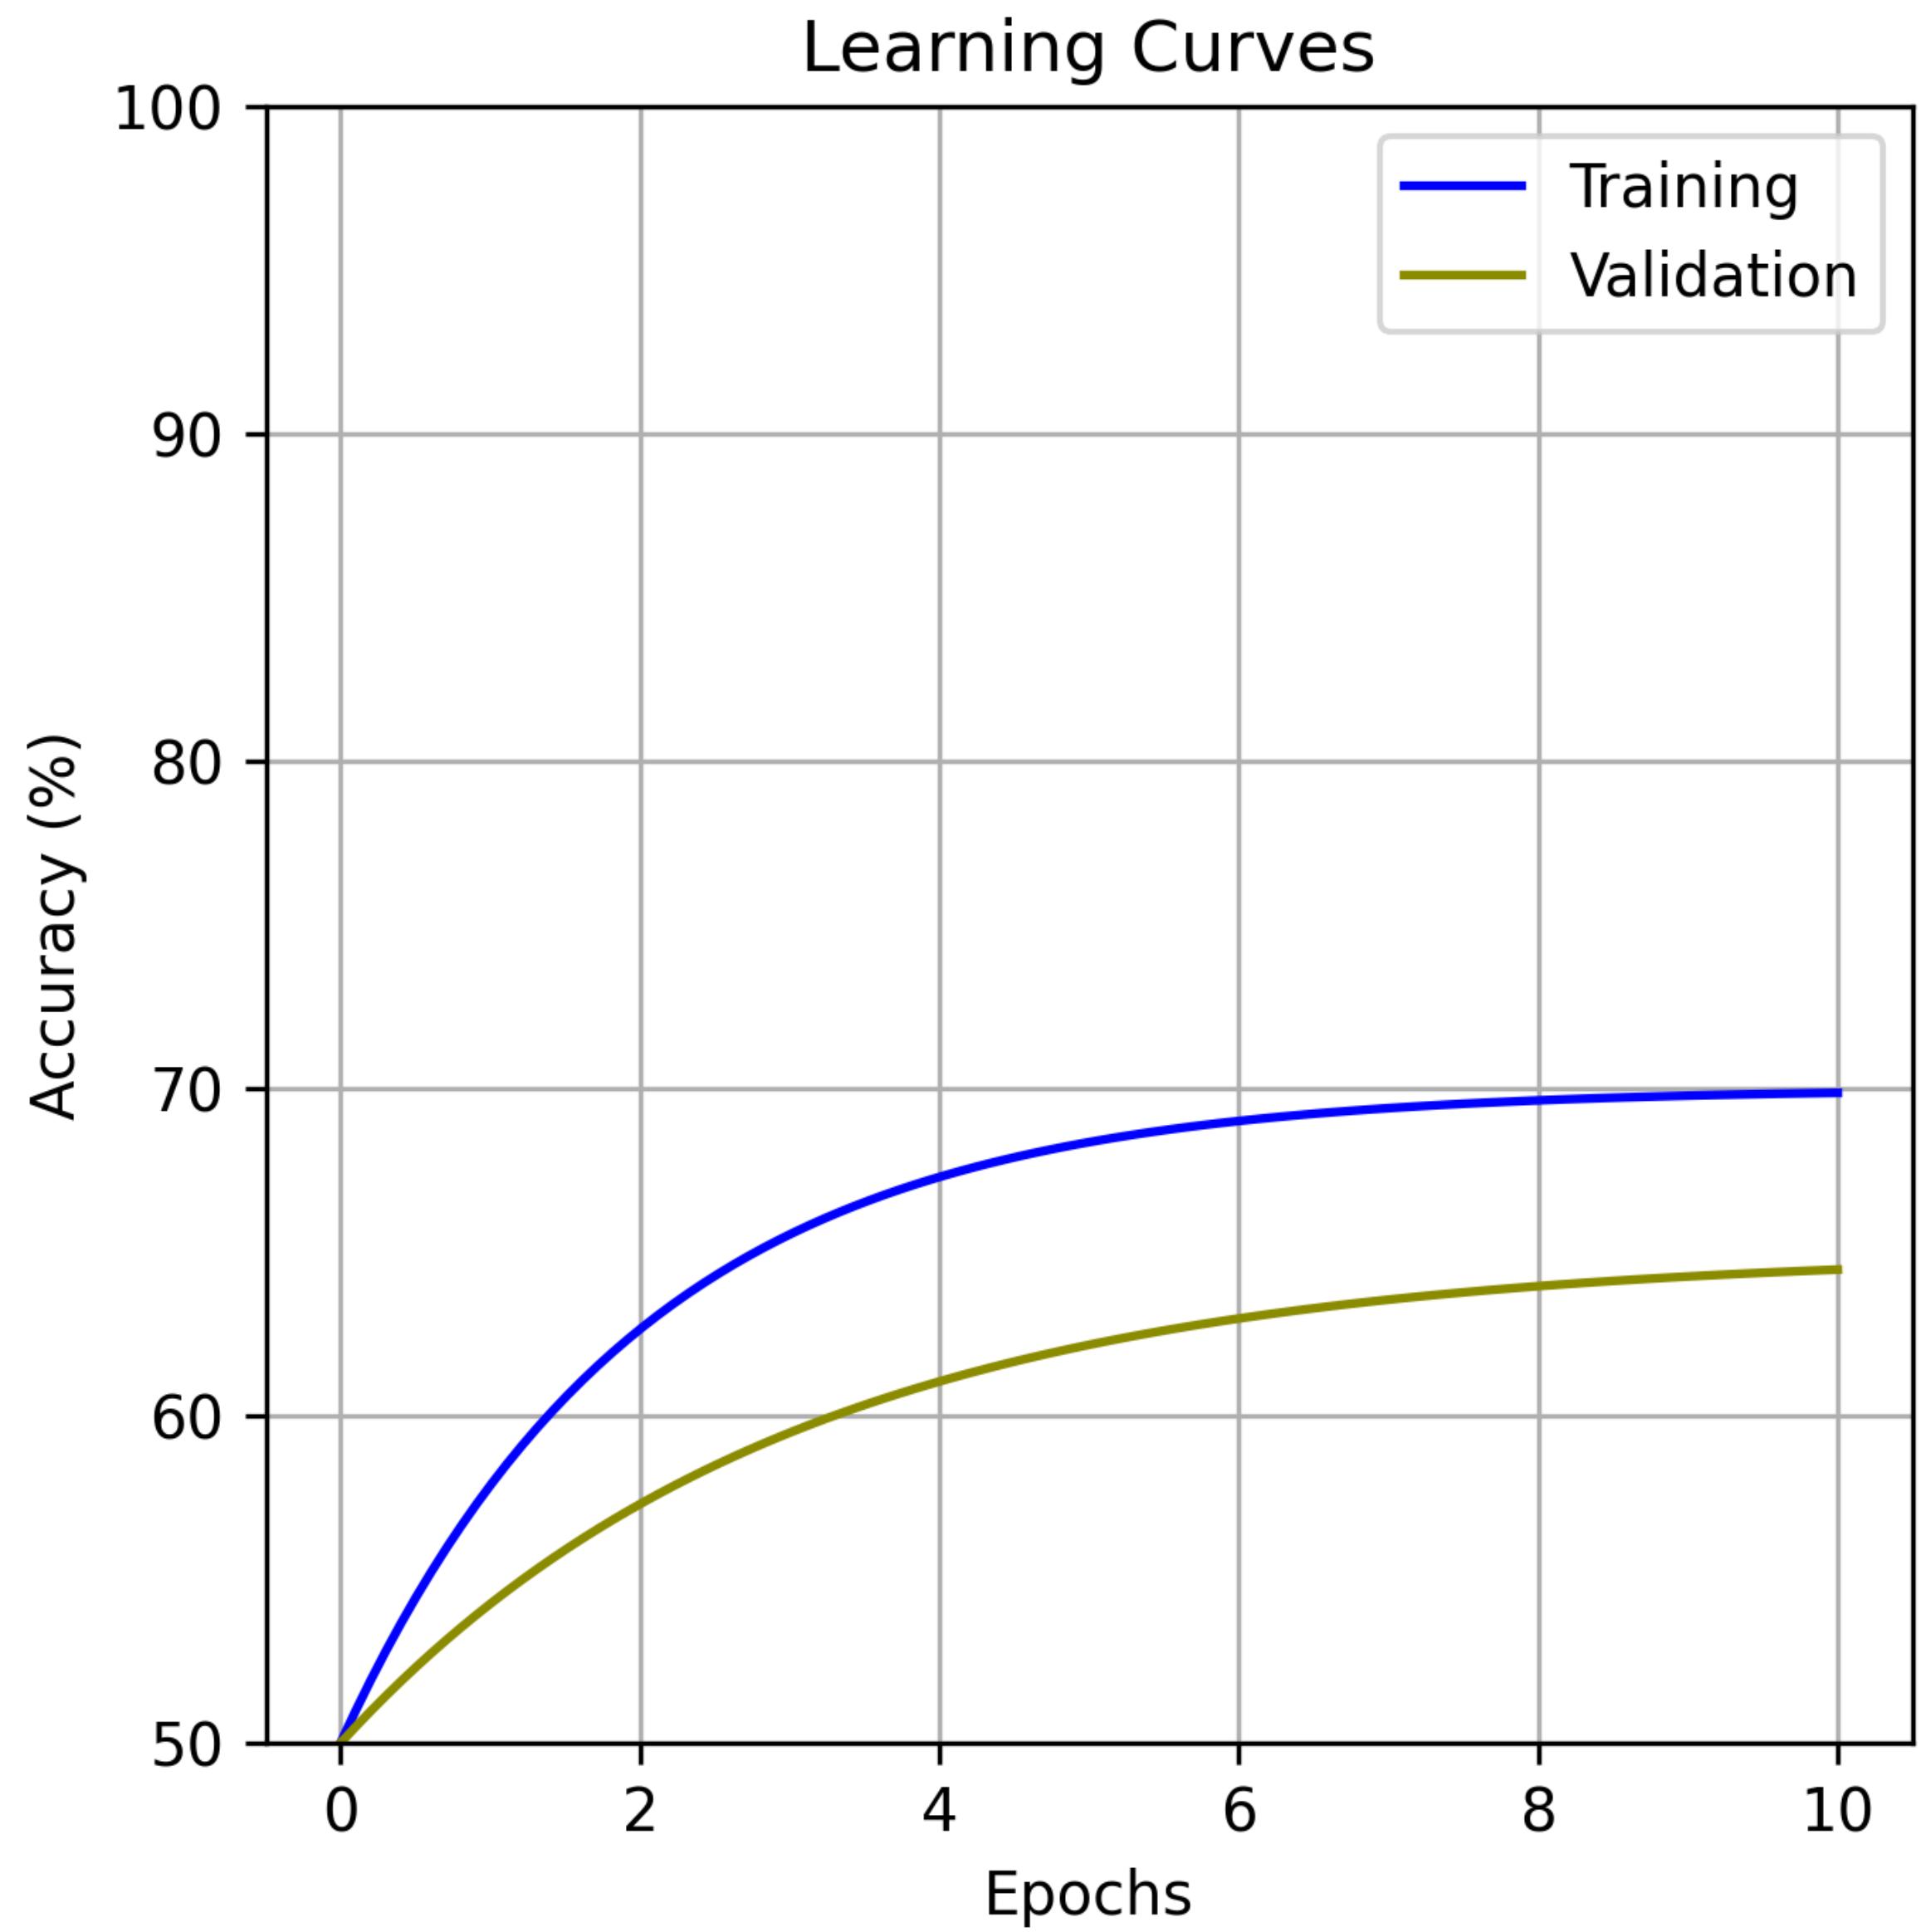


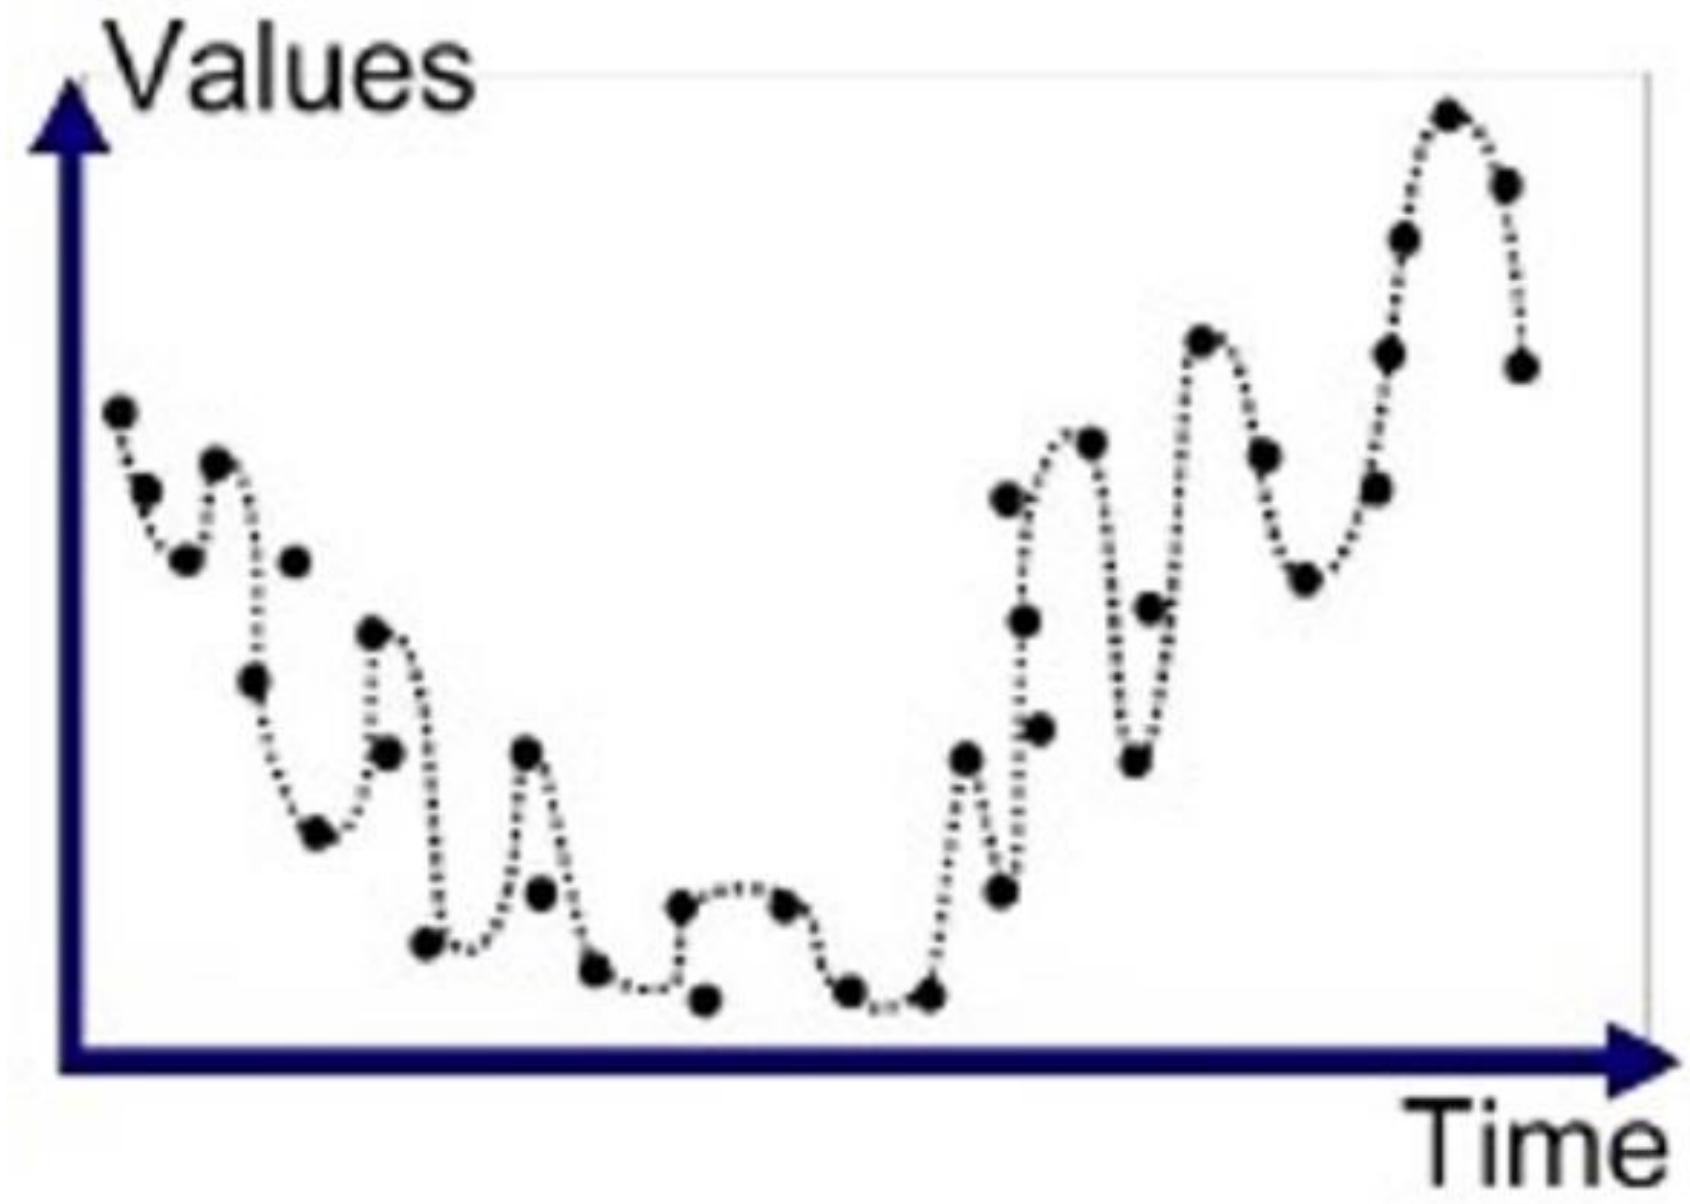
Good Fit/Robust



Underfitted

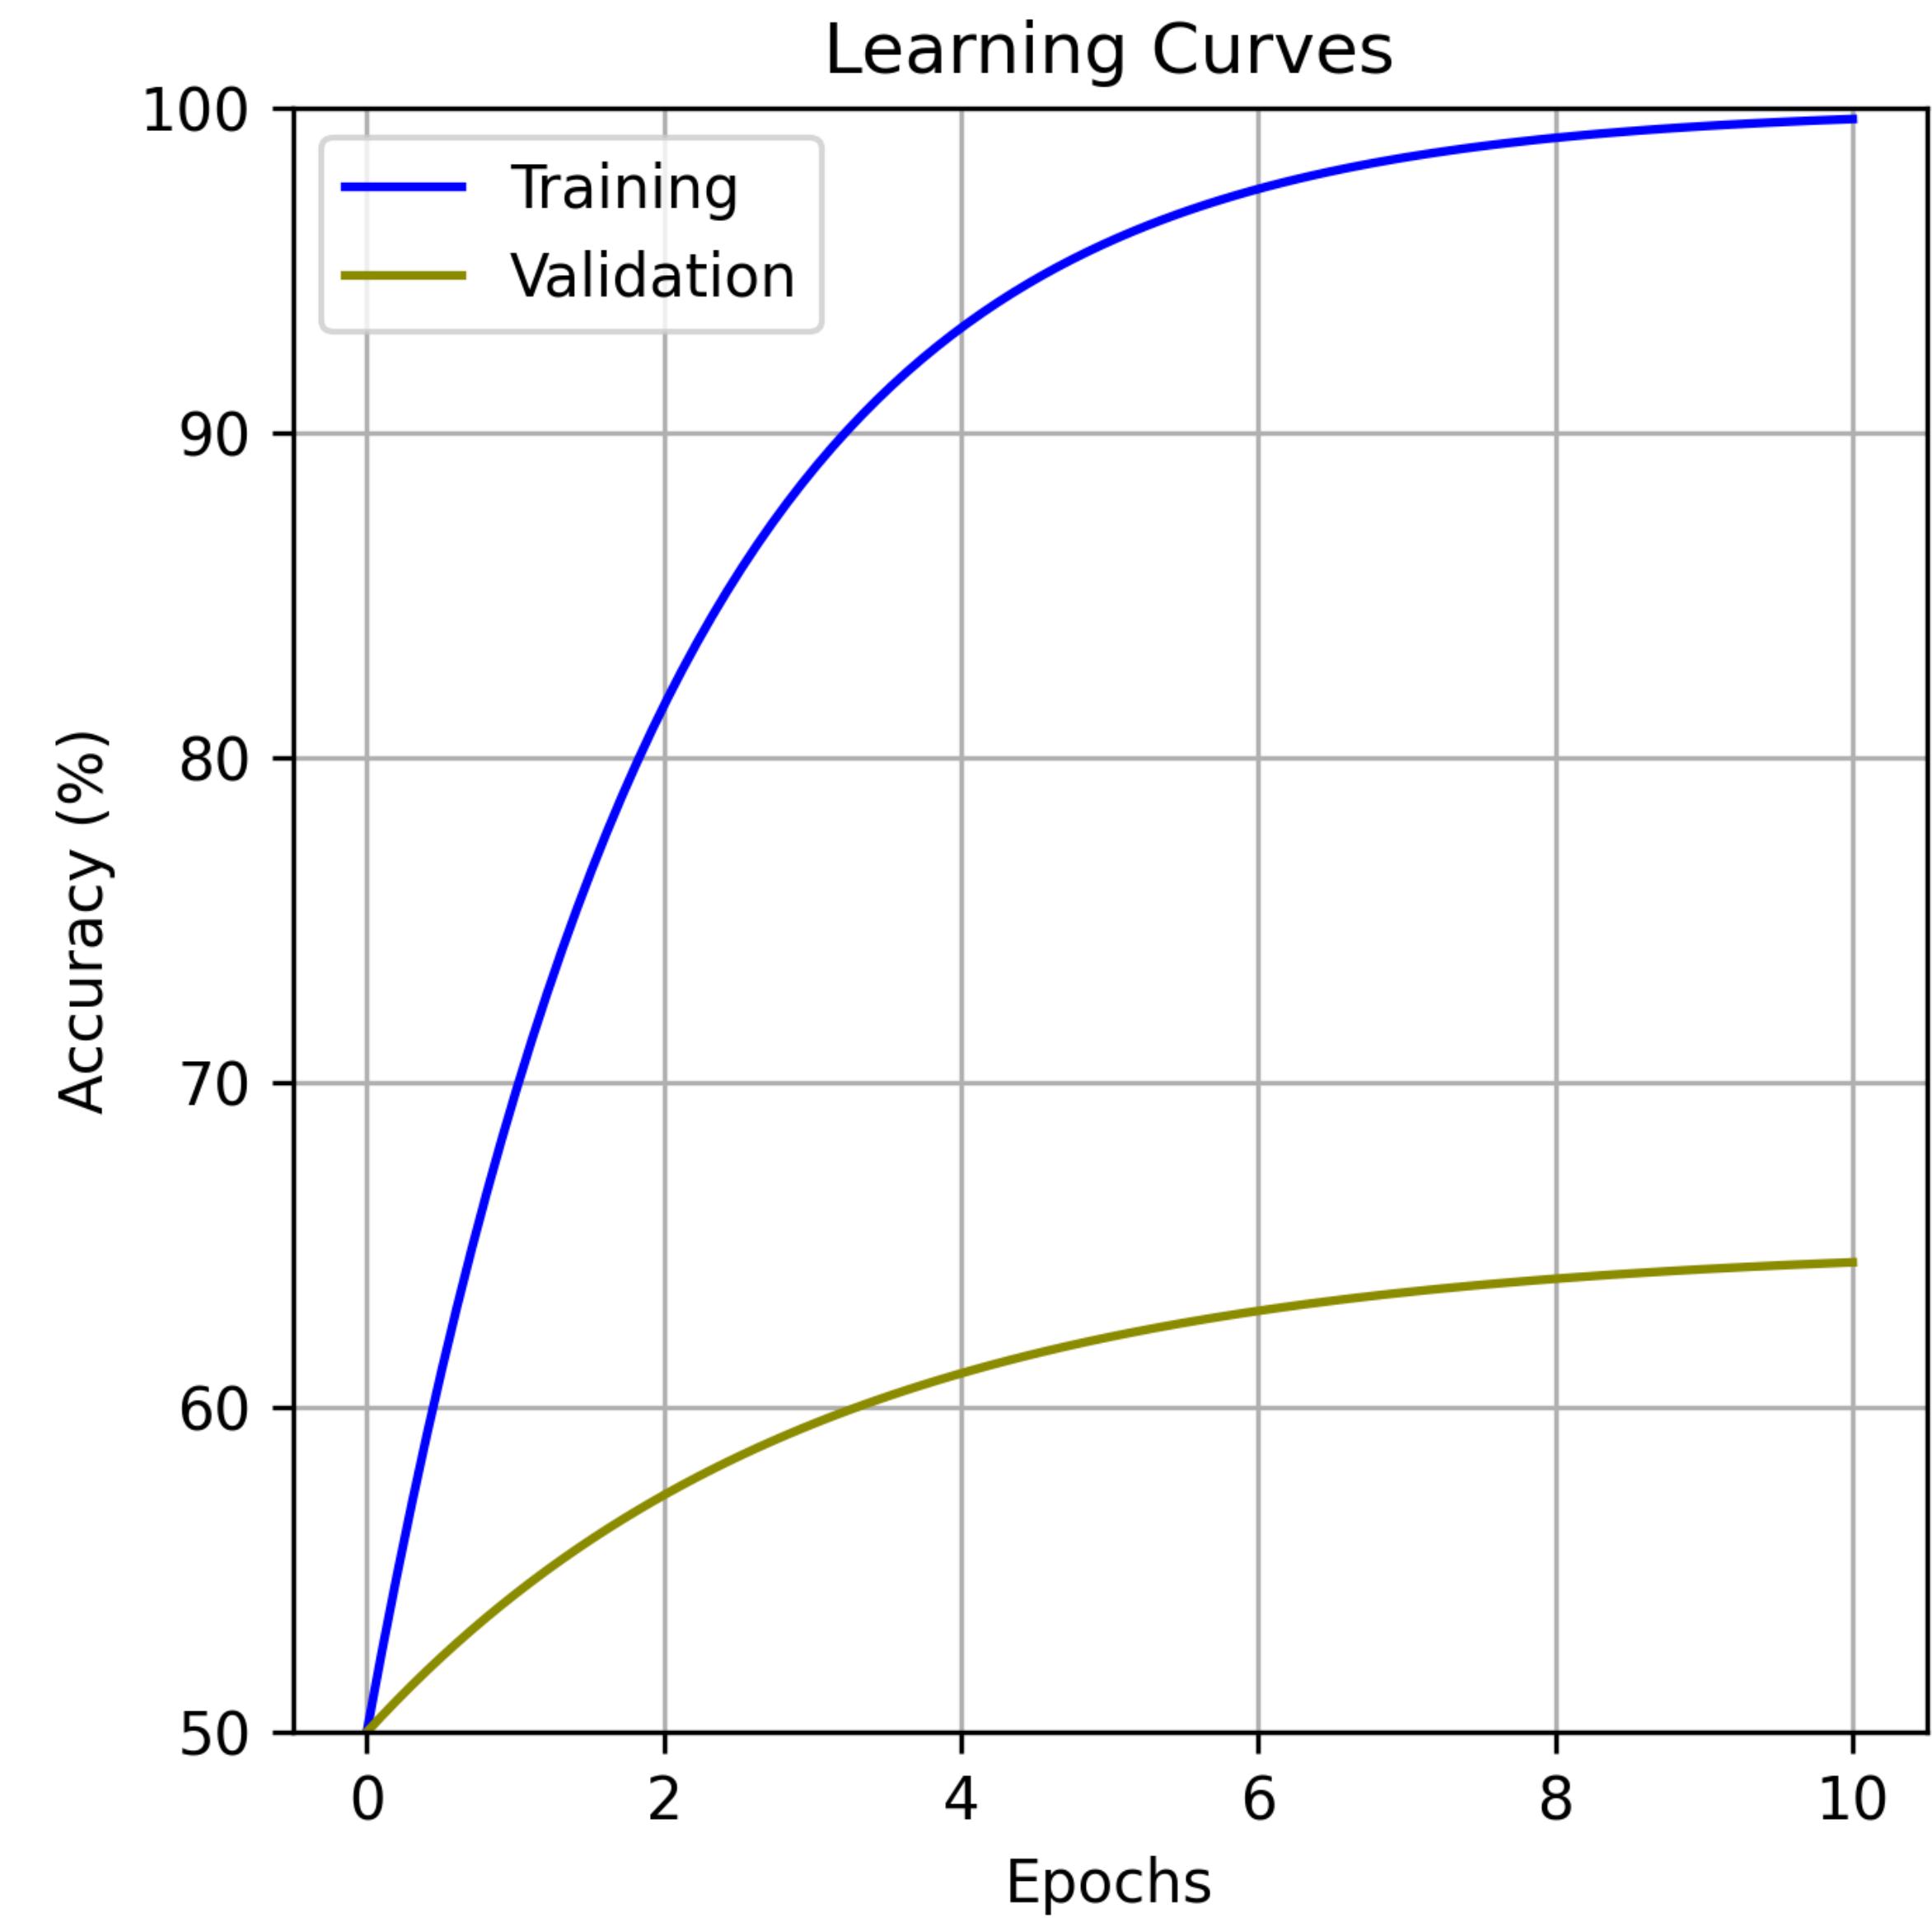
Increase model complexity

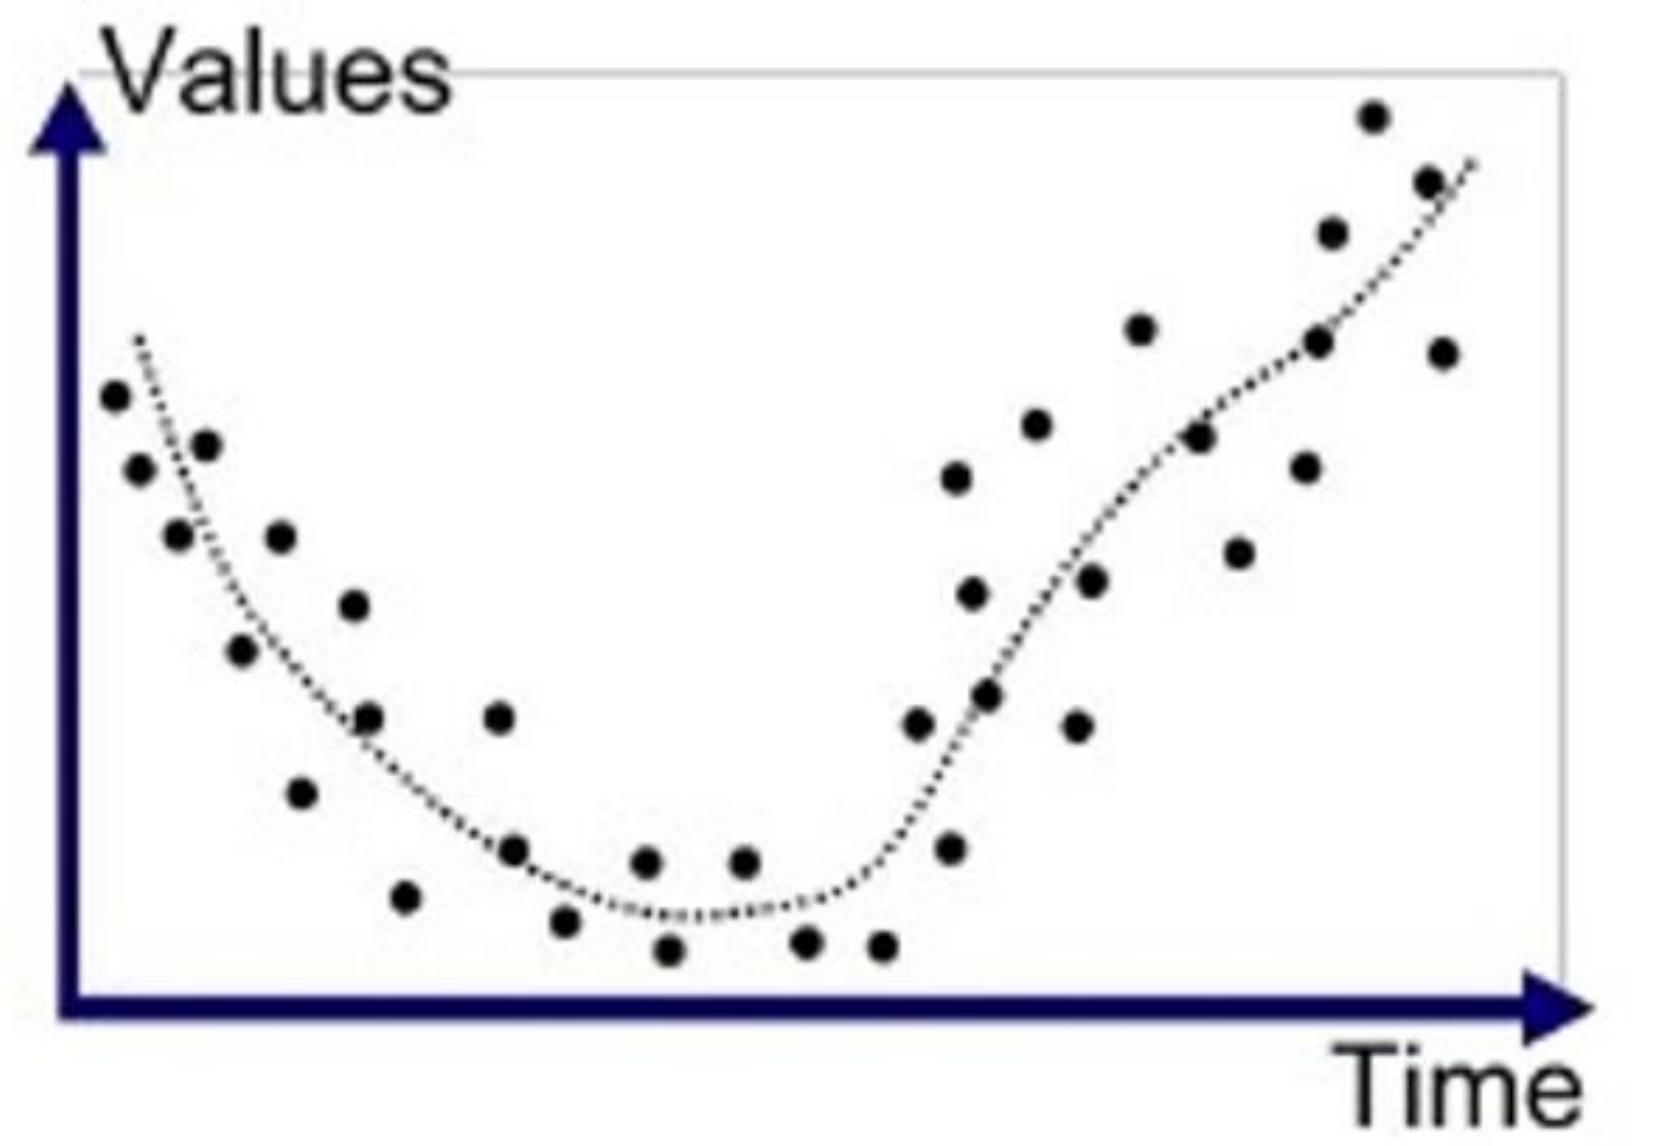




Overfitted

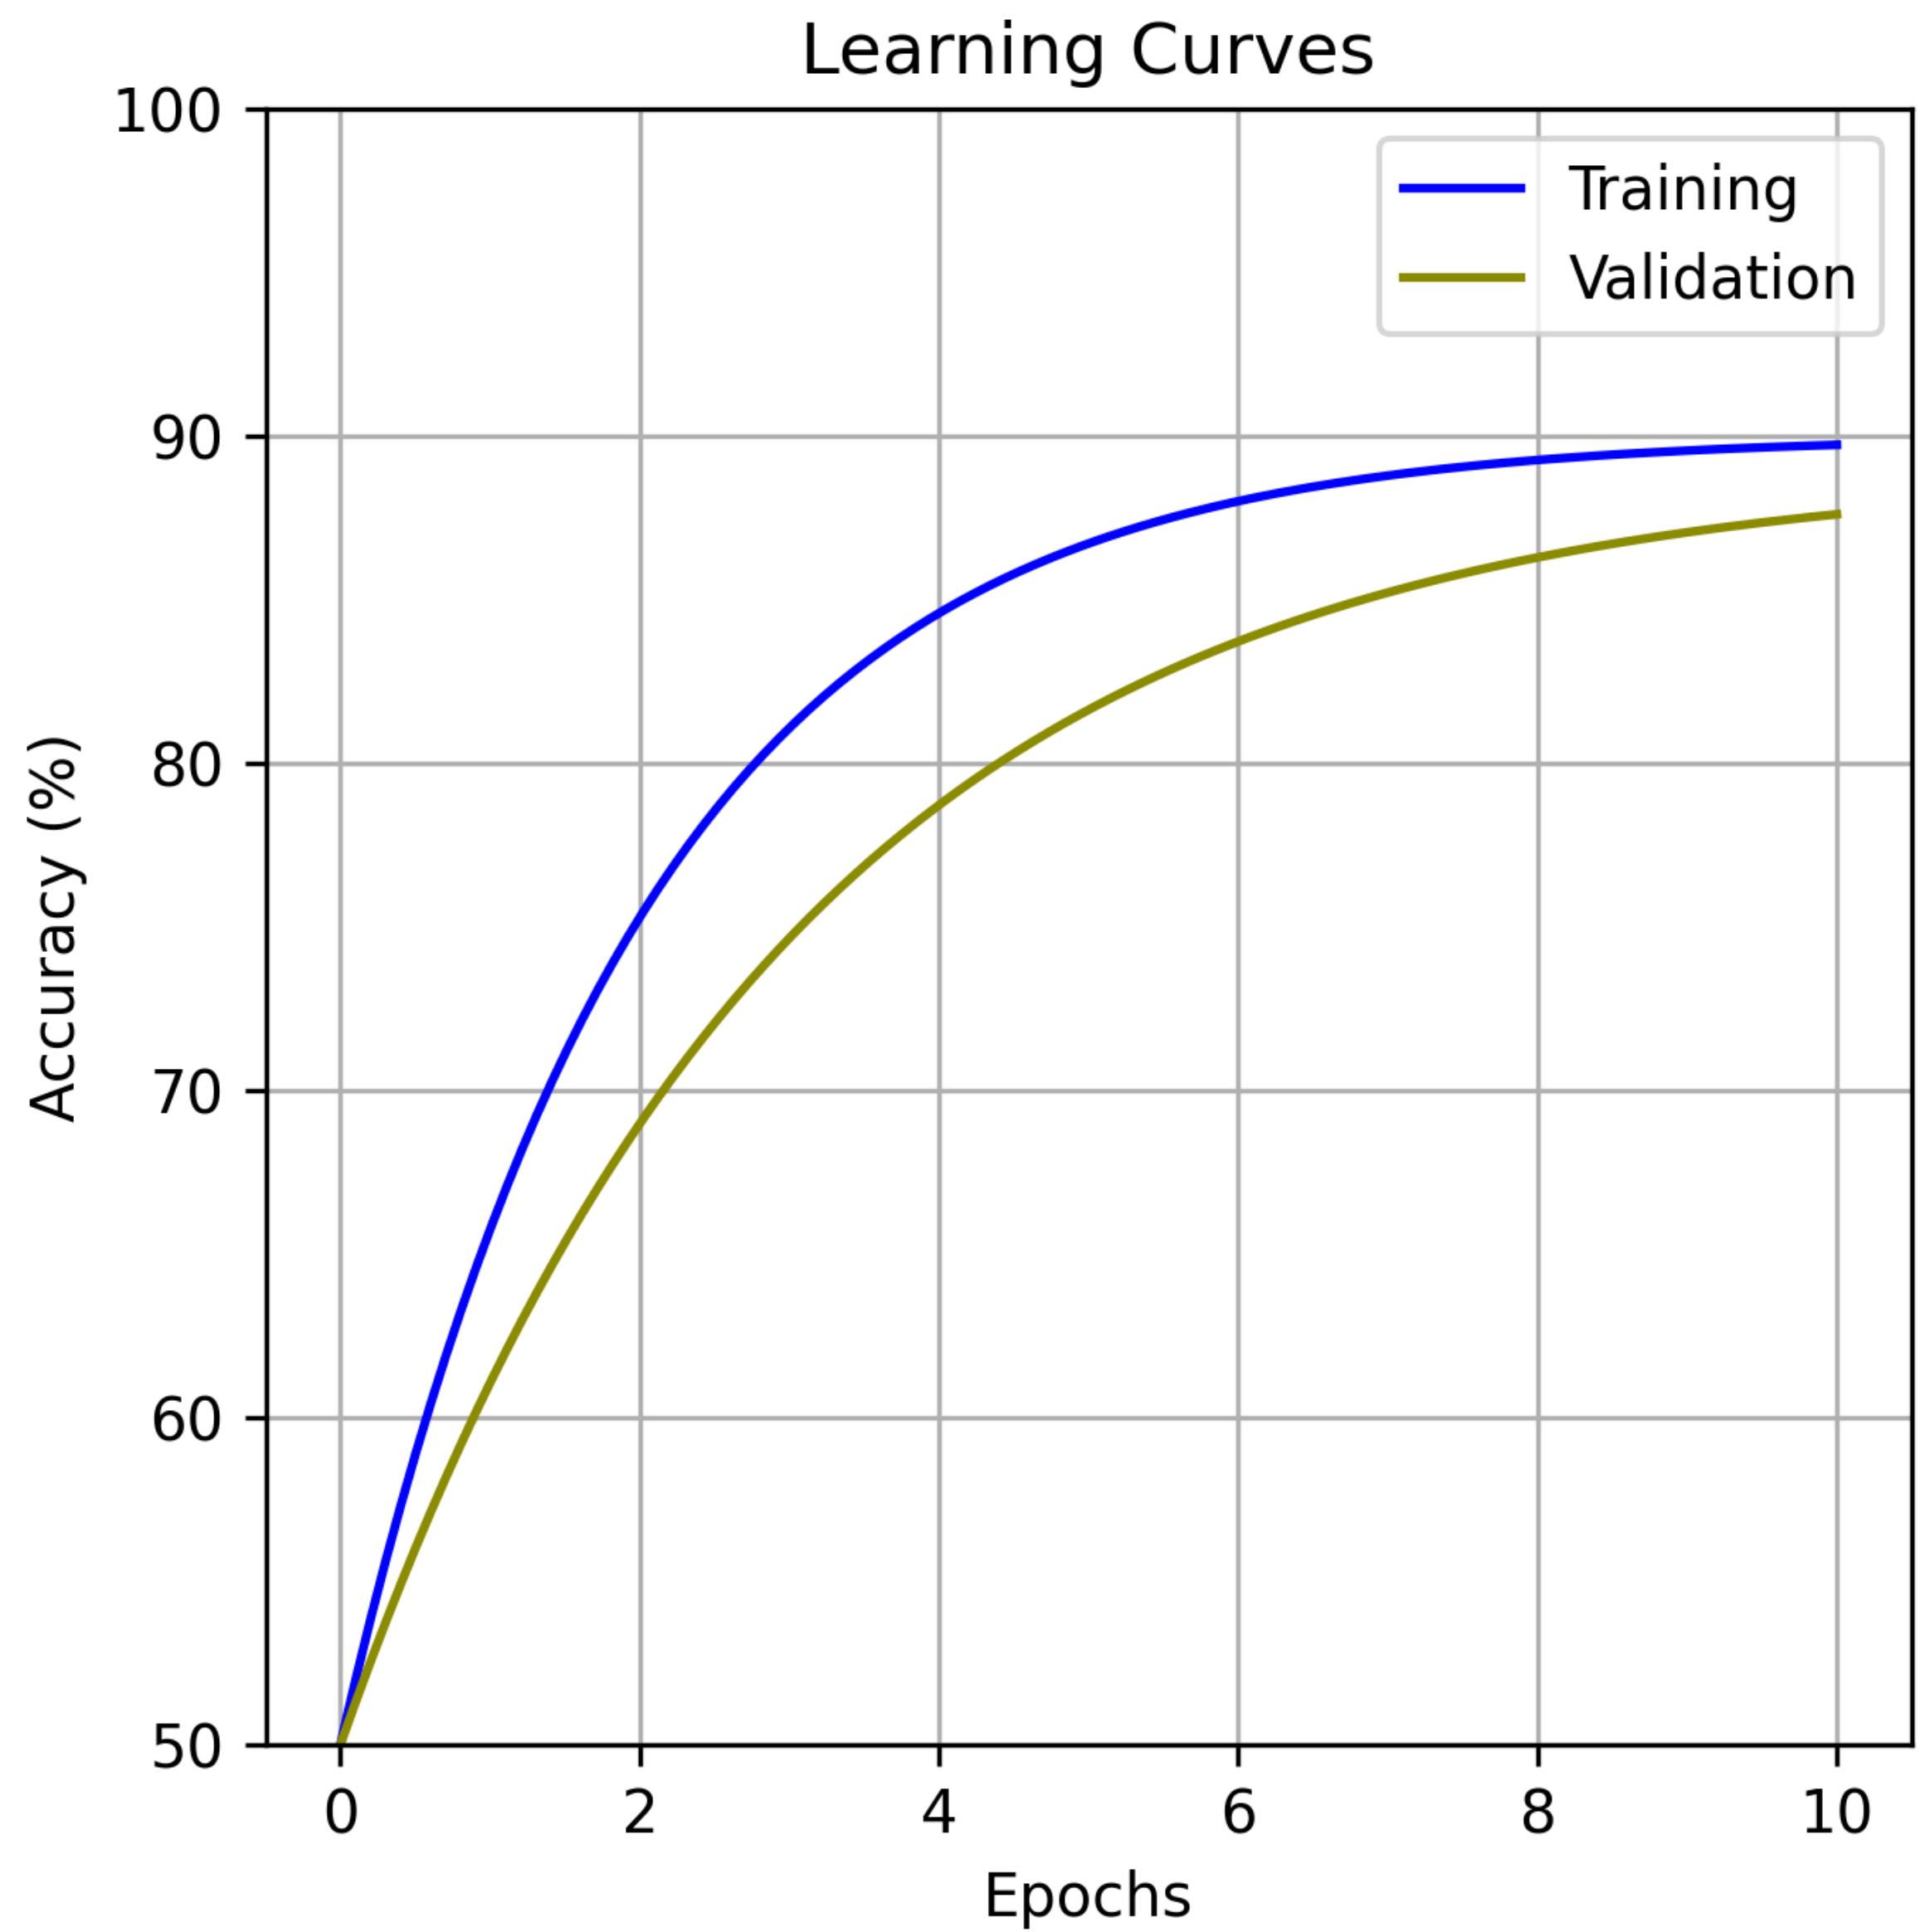
Get more data,  
add regularization

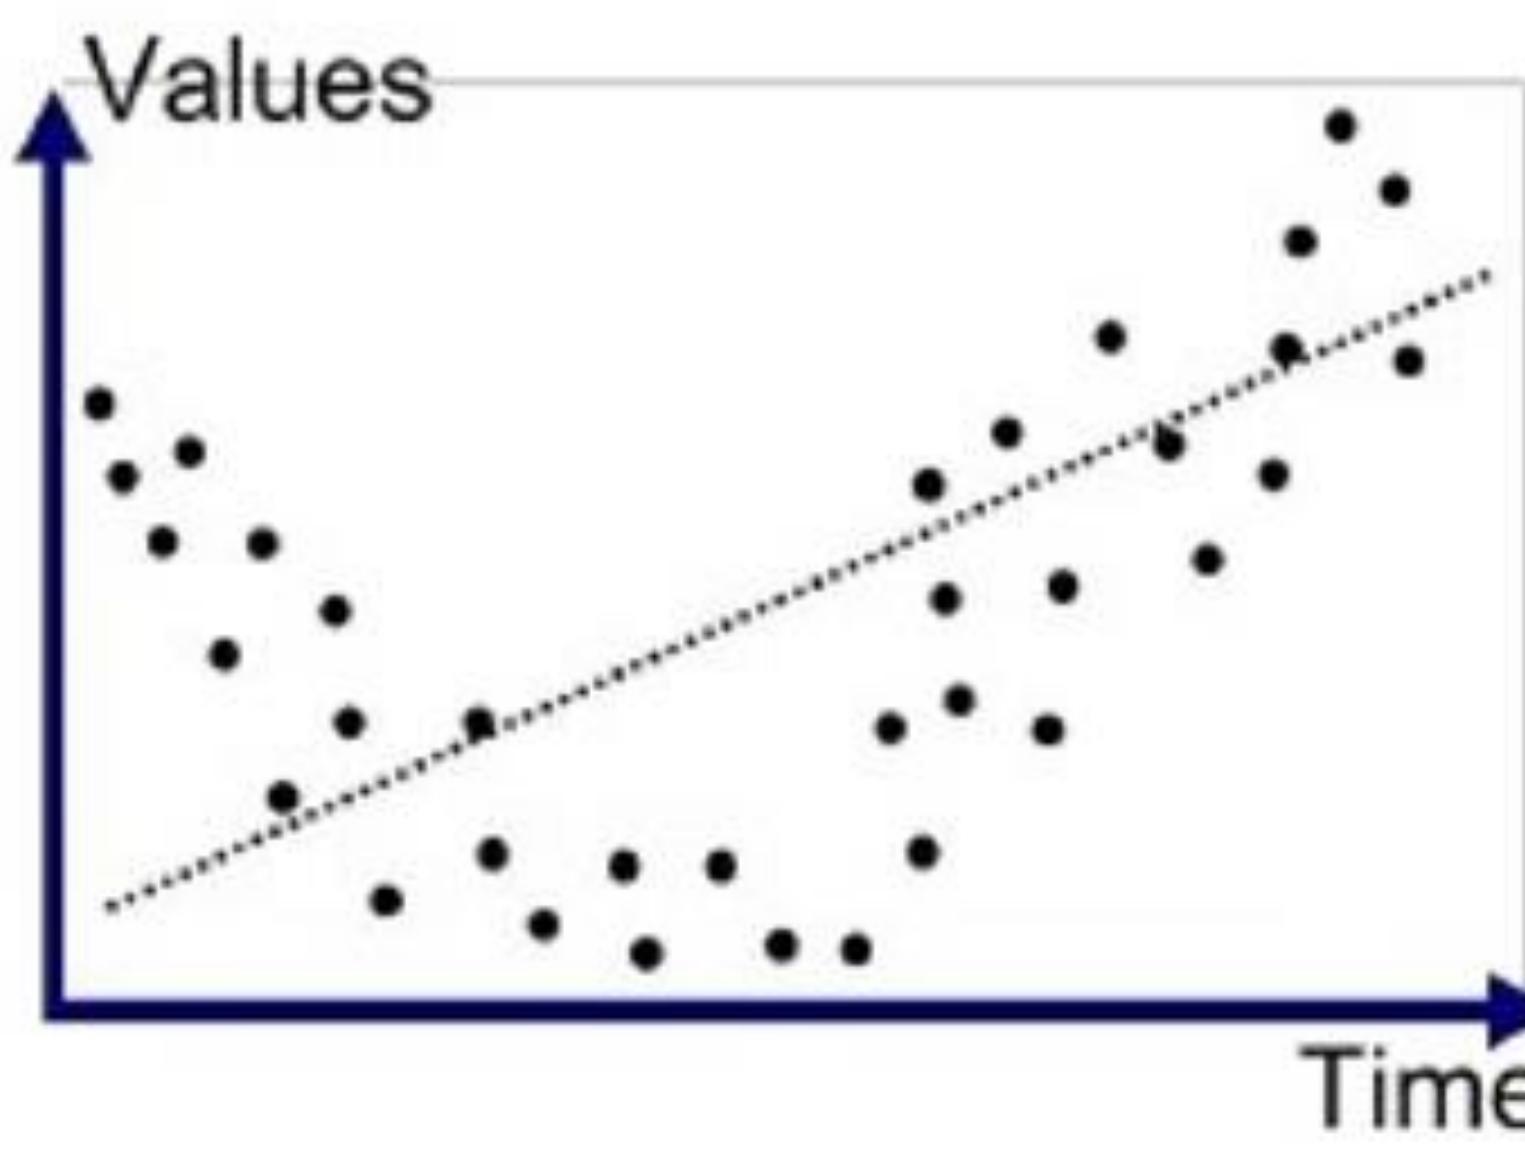




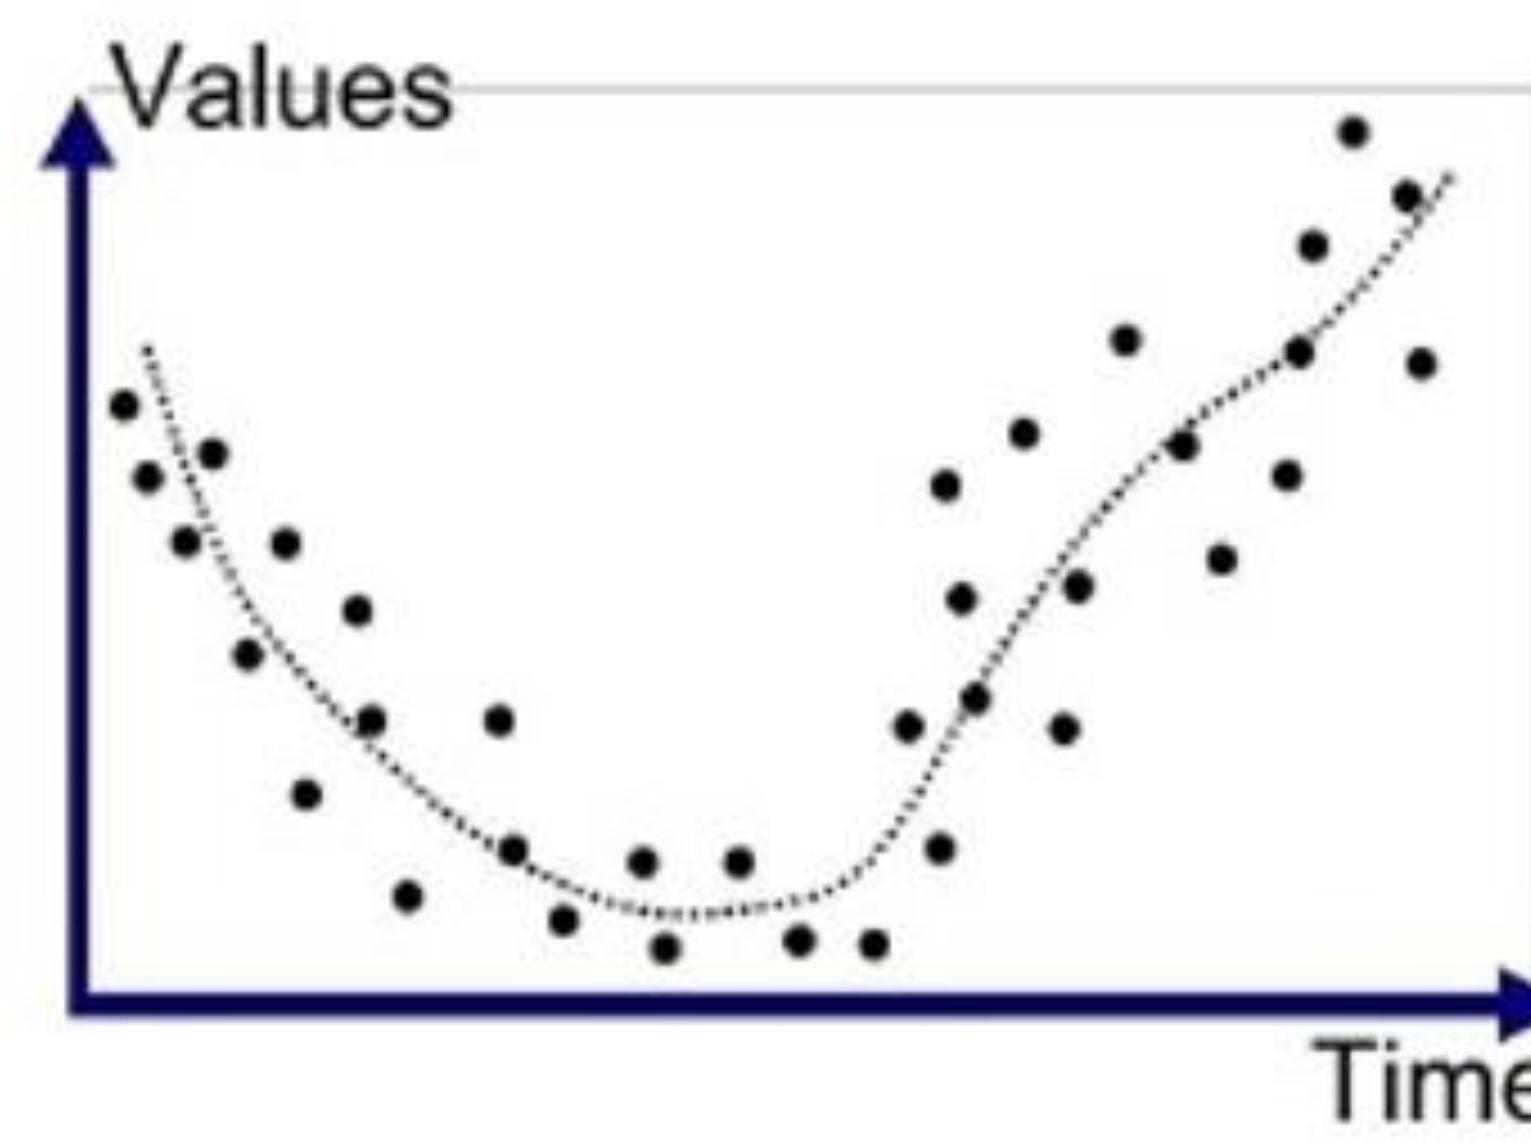
Good Fit/Robust

All done!

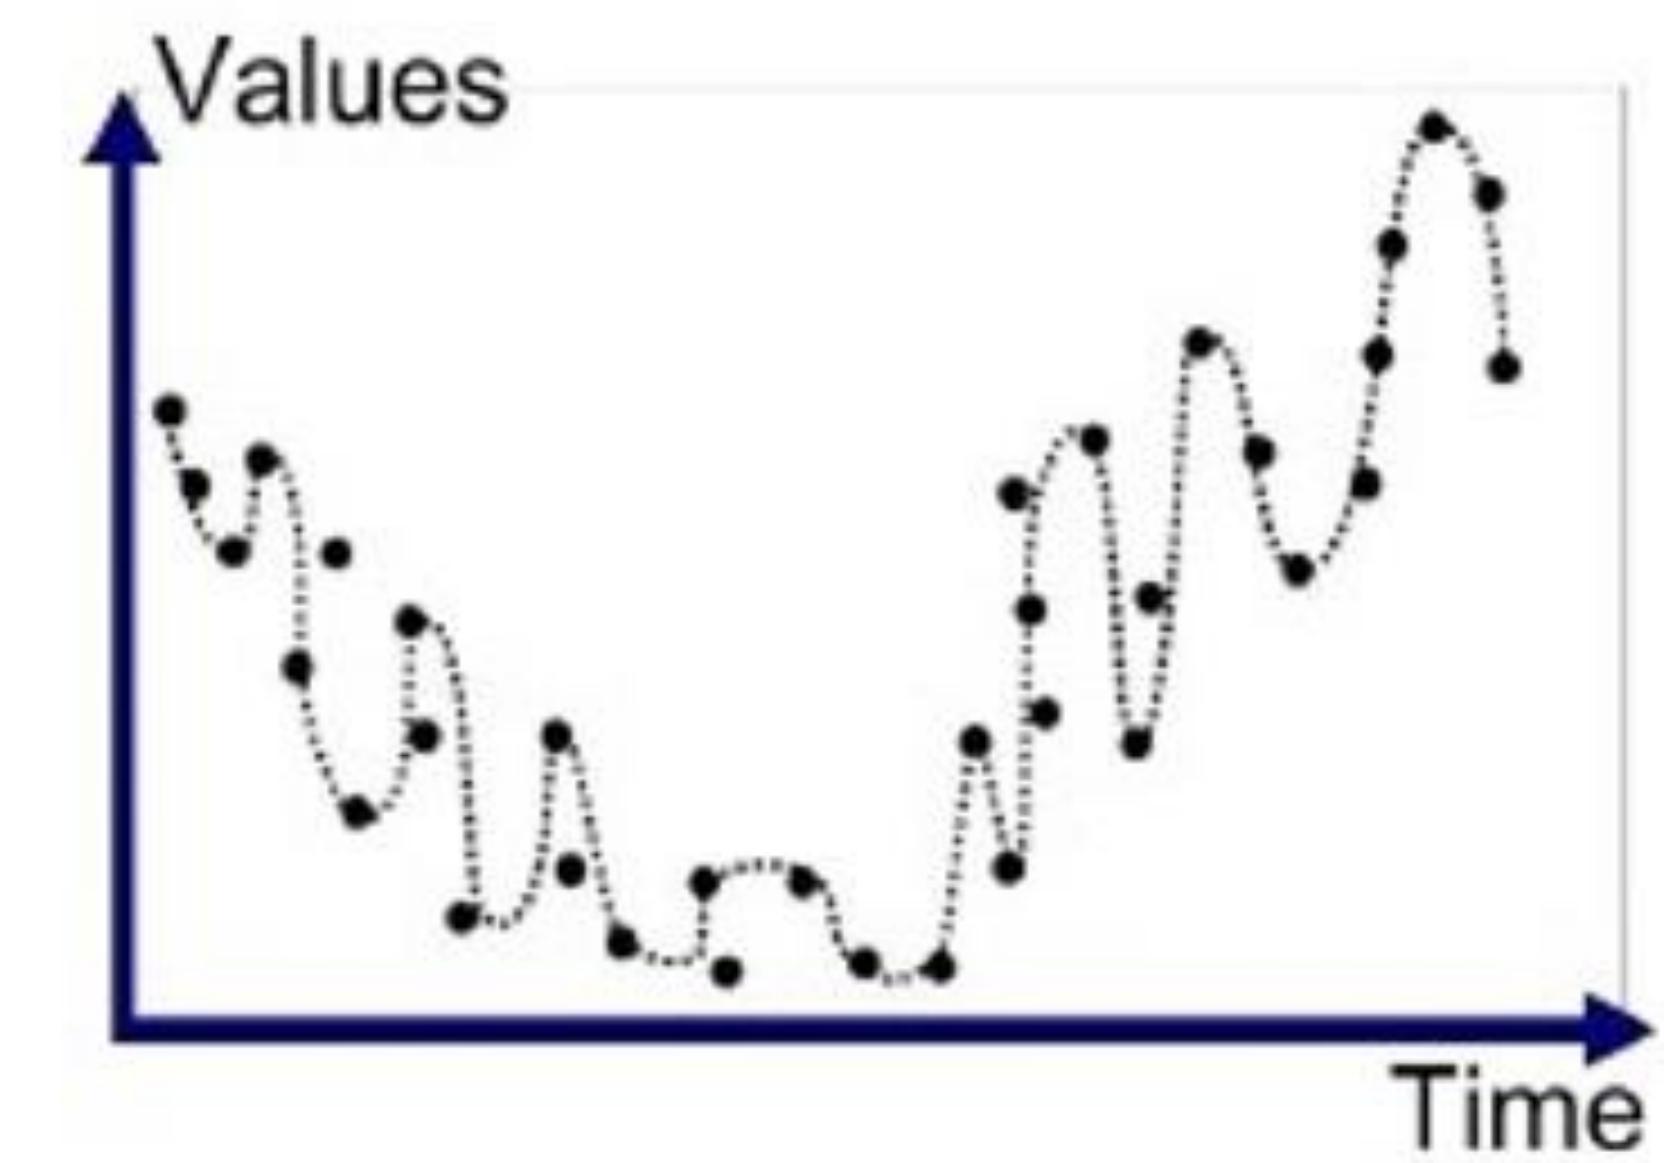




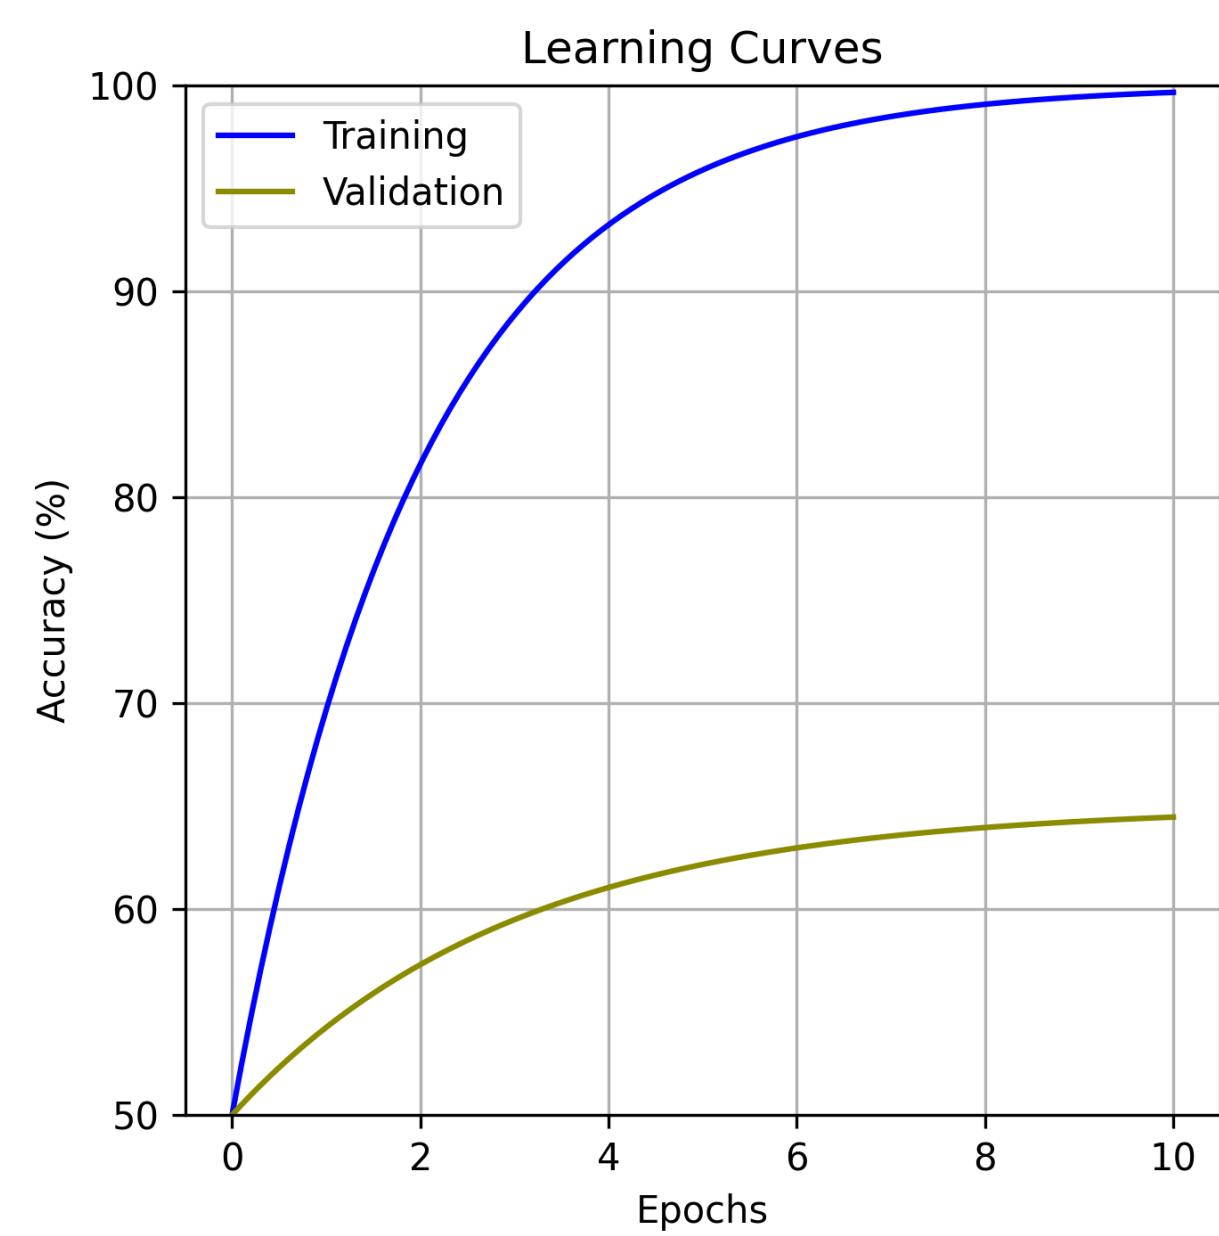
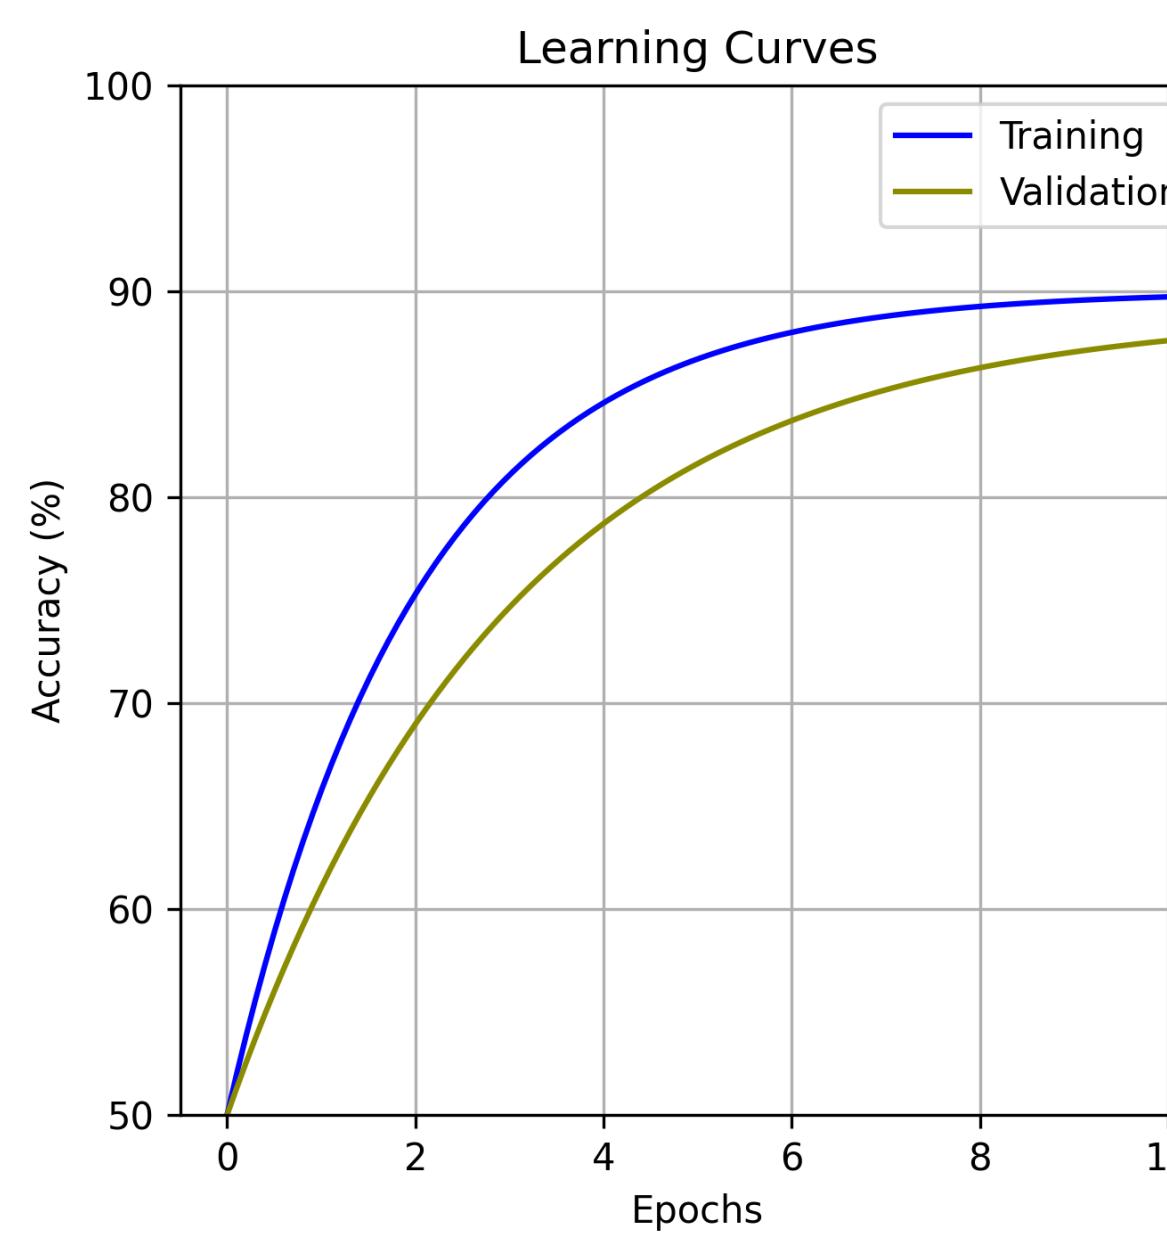
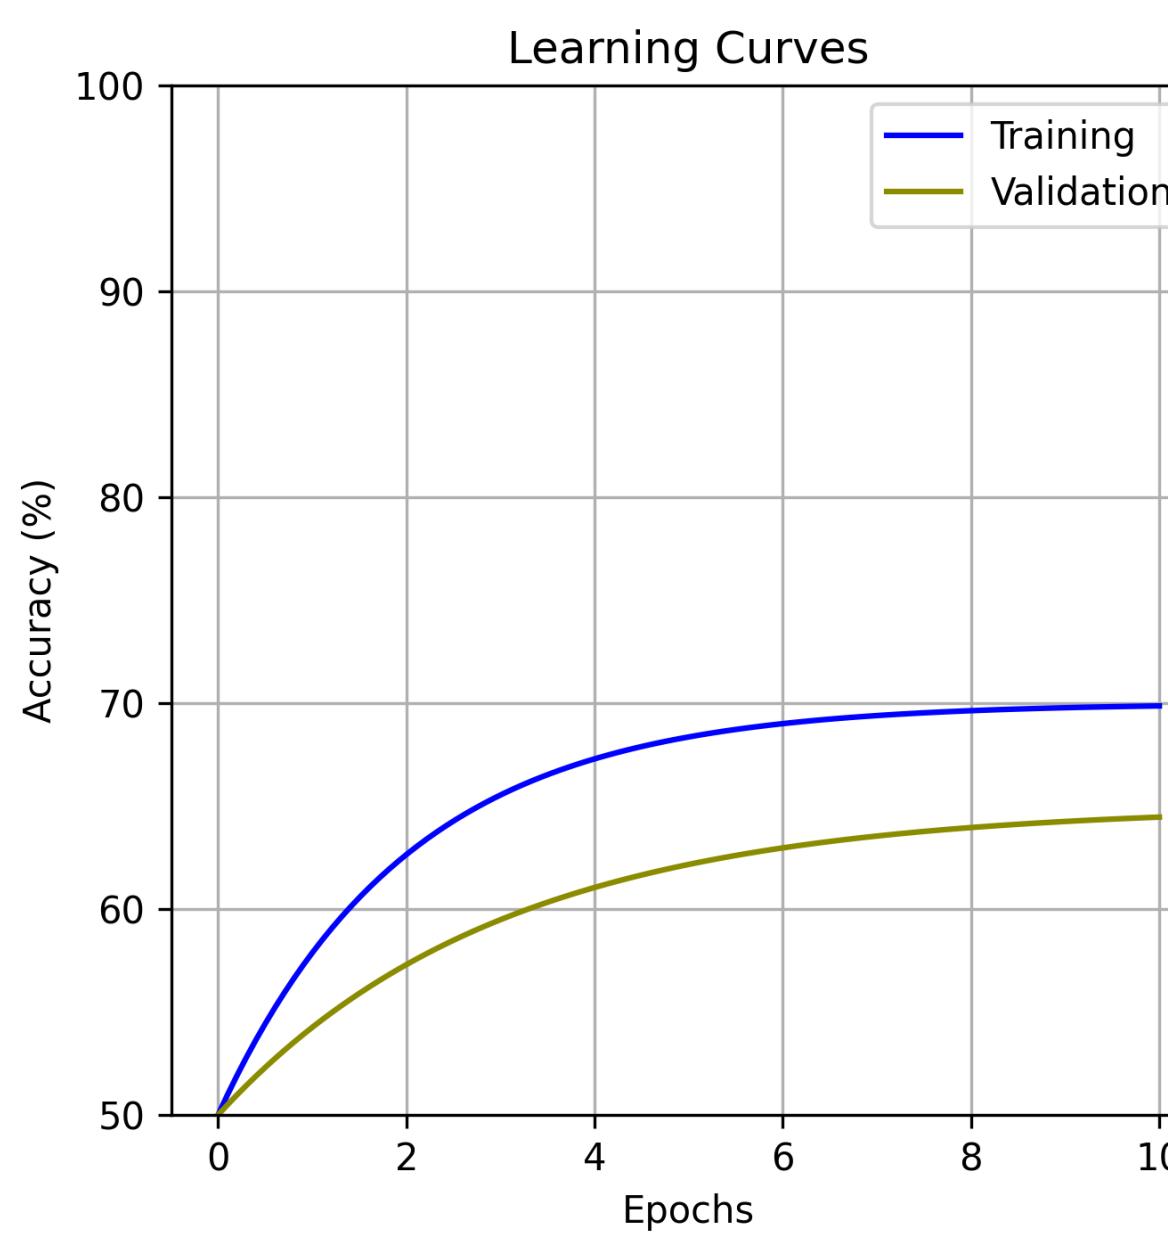
Underfitted



Good Fit/R robust



Overfitted

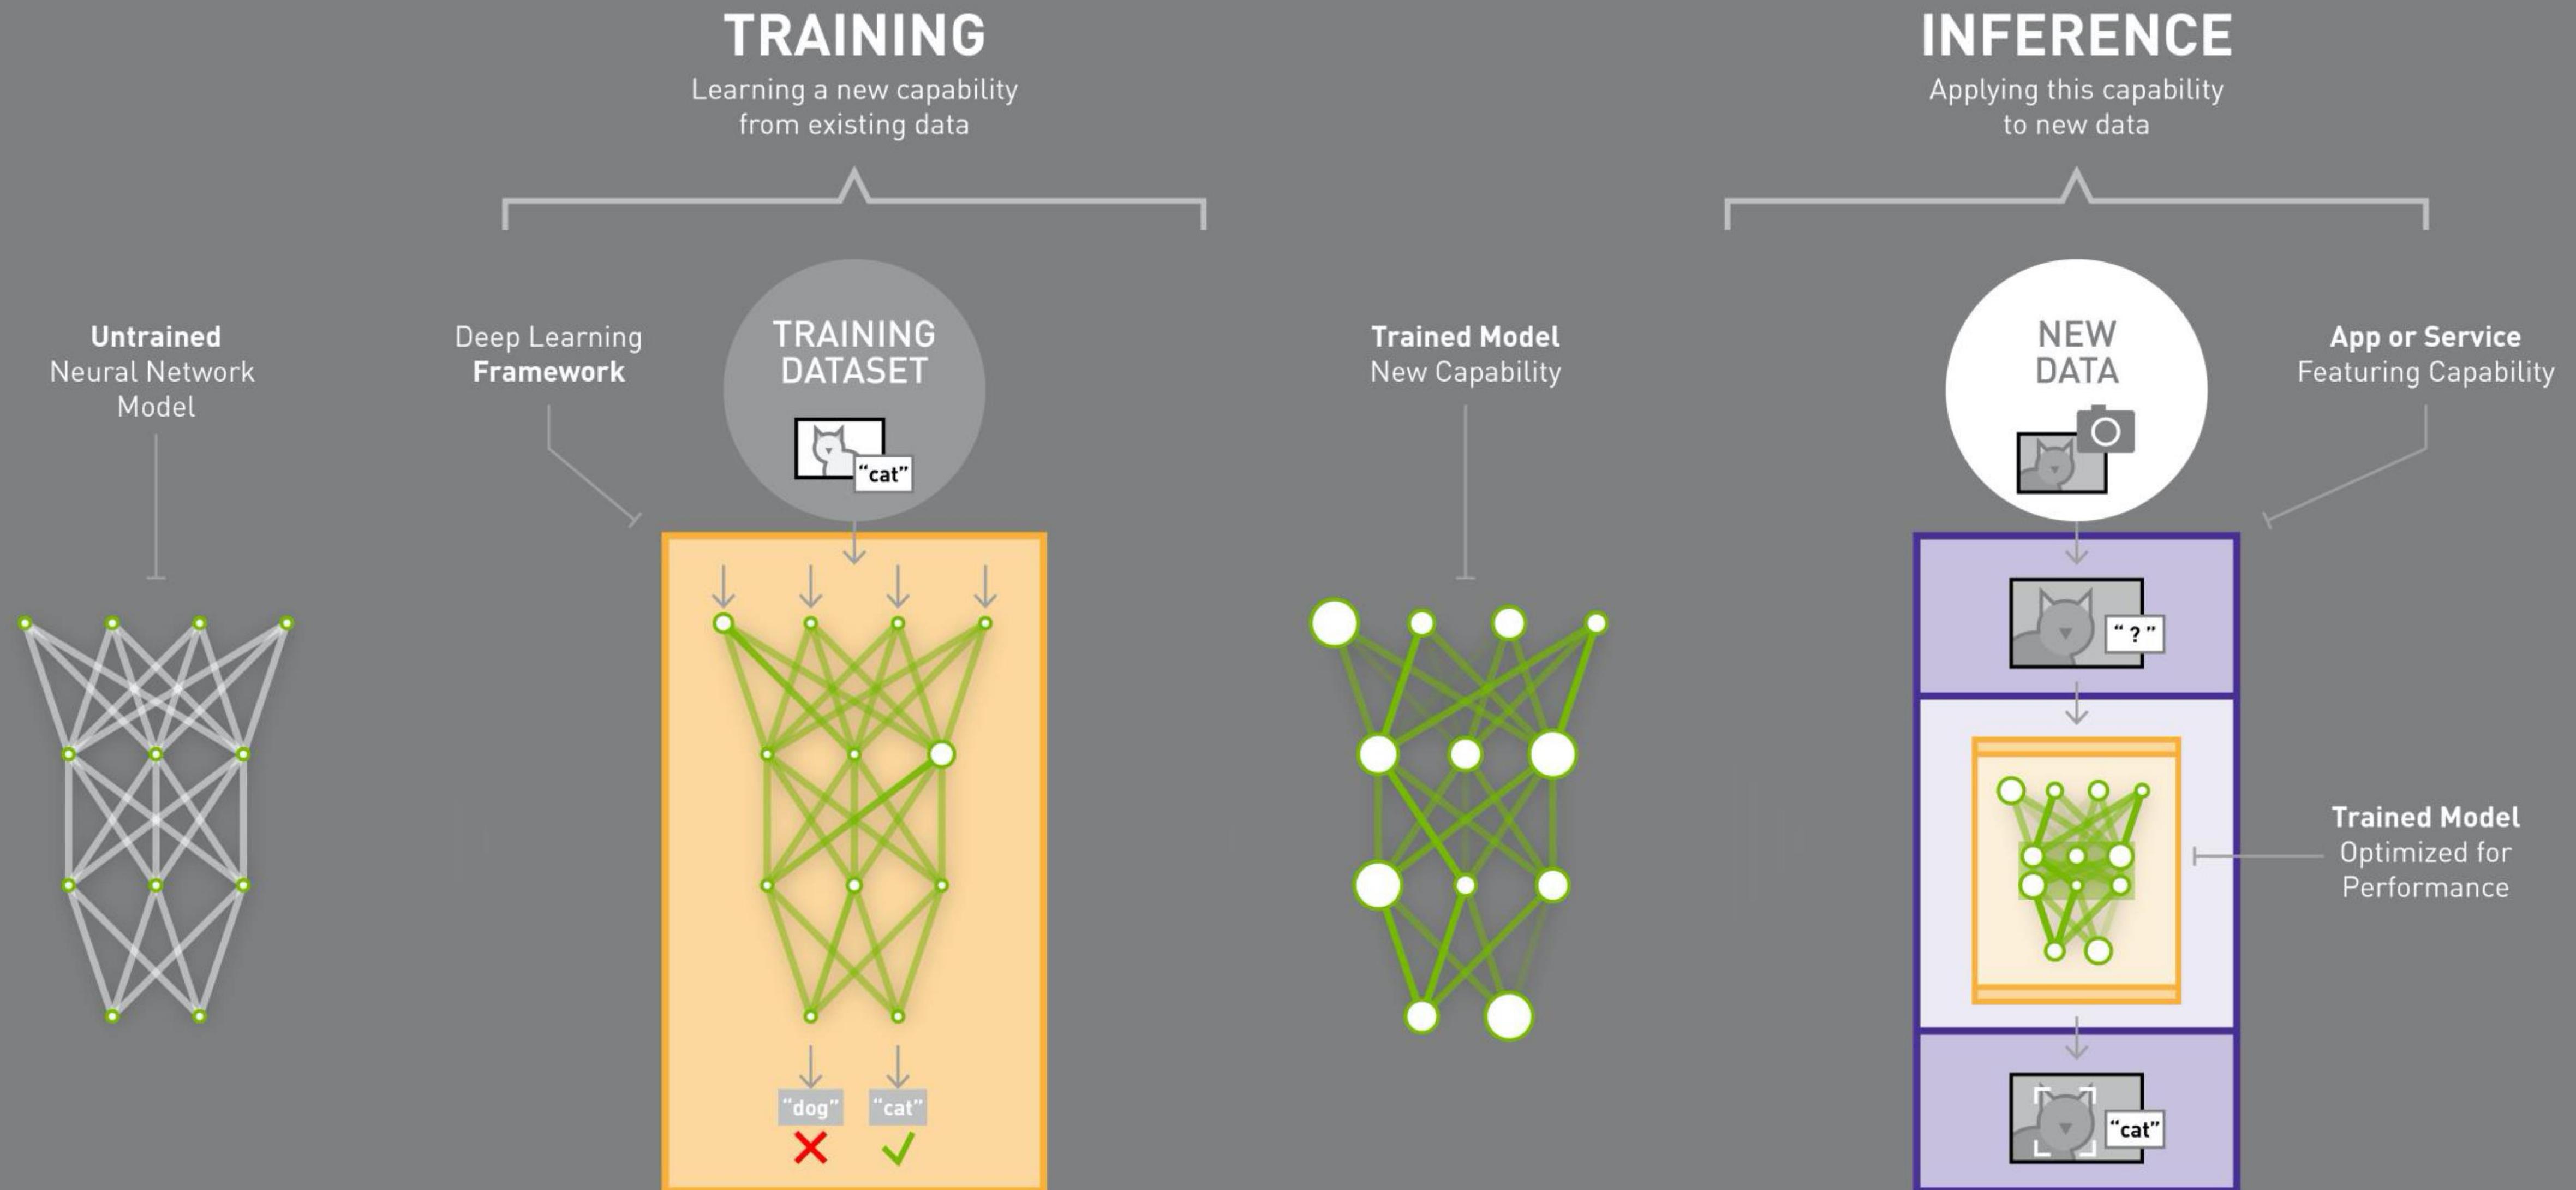


Tip: Take a small sample of your training set and try to overfit it.

# Lifecycle of a computer vision DL project

1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
- 6. Test the model a final time (test set performance)**
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

# DEEP LEARNING



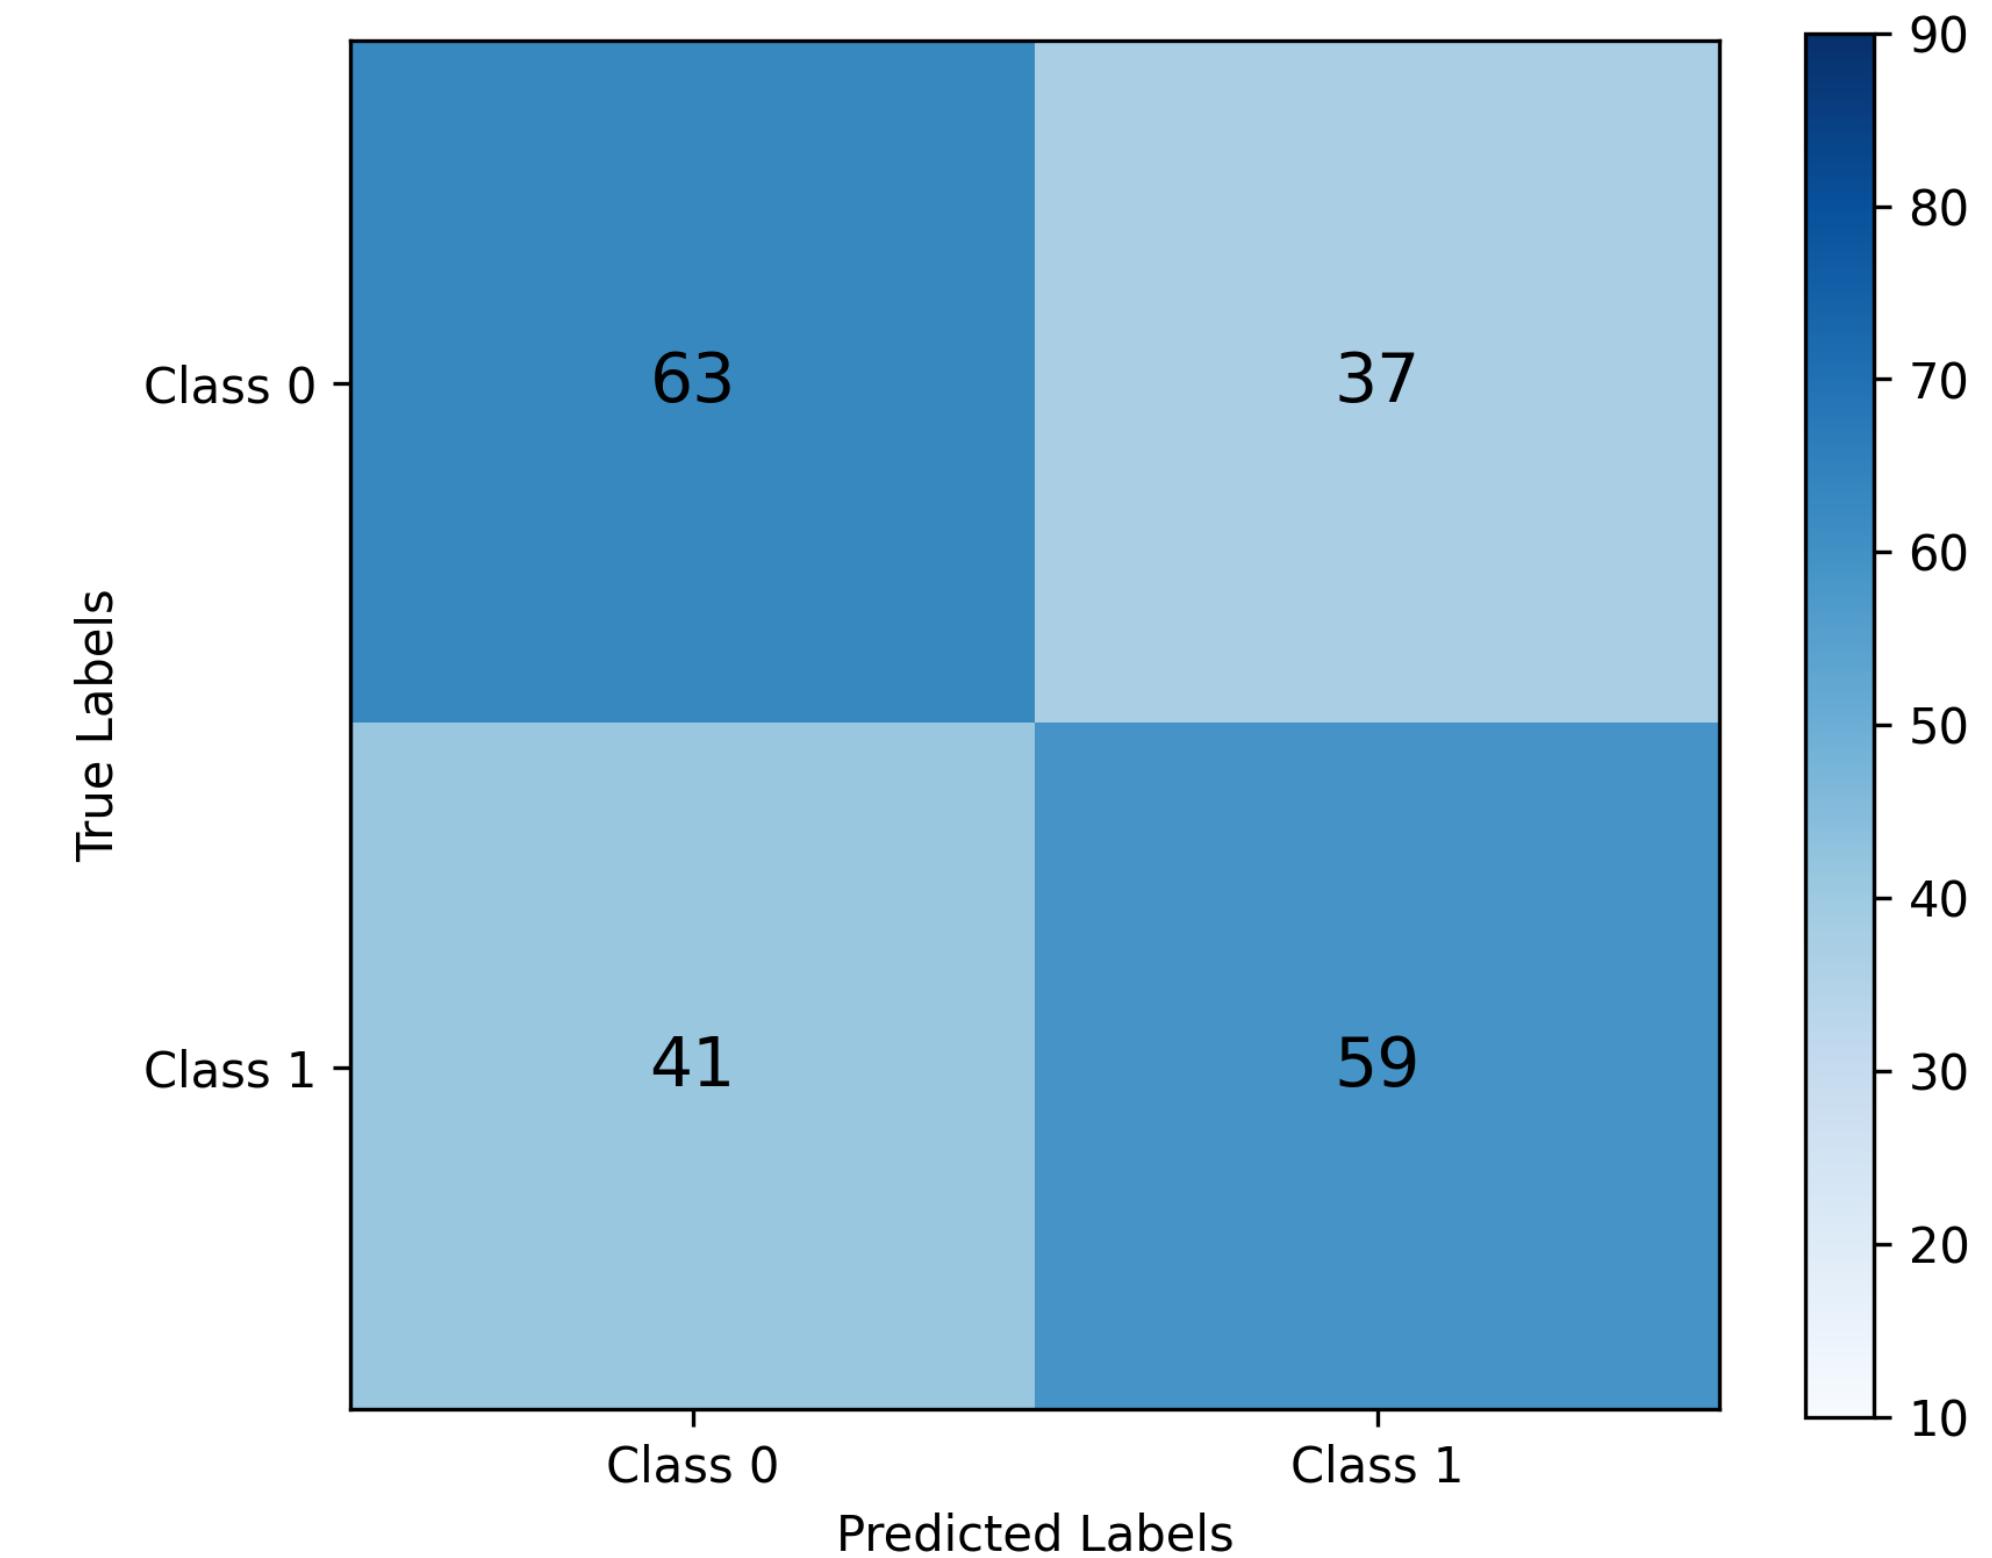
# Confusion Matrix

- (Basically a 2x2 contingency table for binary classification)

Binary Confusion Matrix (Good Model Performance)

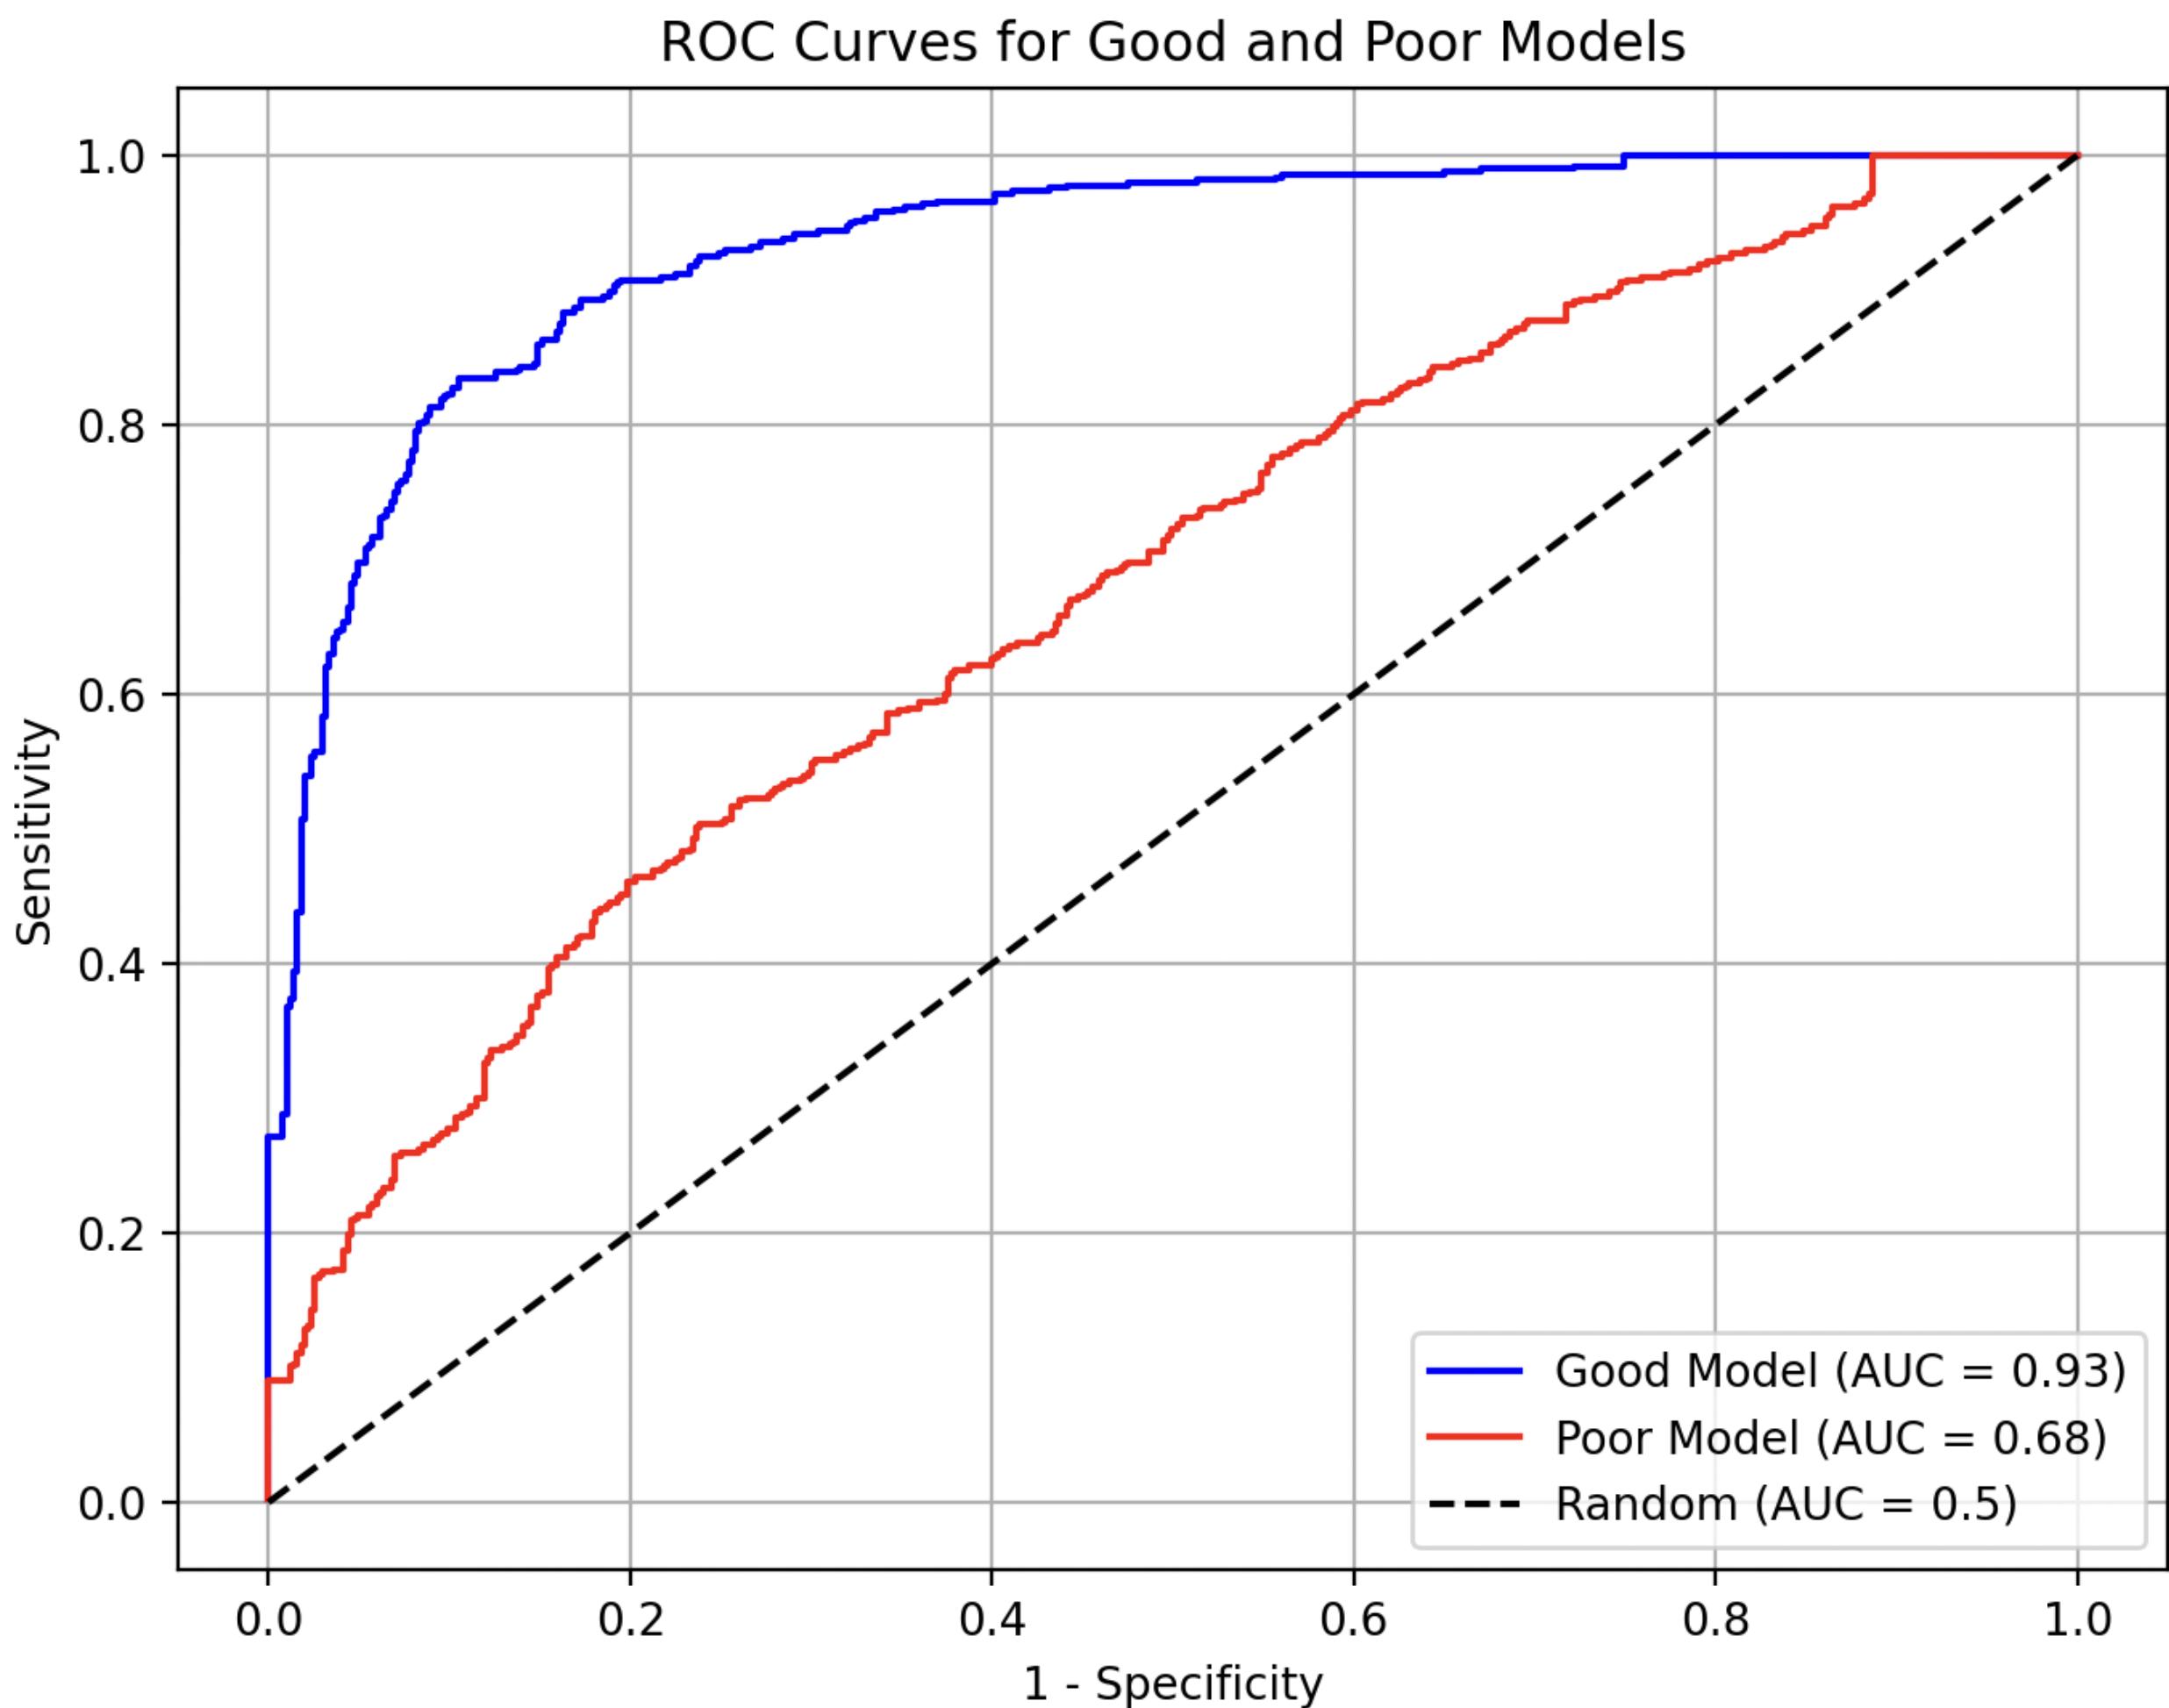


Binary Confusion Matrix (Bad Model Performance)



# Receiver Operator Characteristic curve (for balanced datasets)

- The model will produce an output between 0 and 1.
- We can vary the threshold and recalculate the **Sensitivity** and **Specificity** for each threshold
- Random behavior = diagonal line
- Area under curve = **AUROC**



# Lifecycle of a computer vision DL project

1. Define your problem (model)
2. Collect, clean, collate, partition data (input)
3. Torture humans to label the data (output)
4. Train model (hyperparameters)
5. Repeat 4, sometimes 2-4 (model parameters/weights)
6. Test the model a final time (test set performance)
7. Write a paper, publish a github repo (data + model + weights or at least model + weights)

# Do you need FDA approval before you can use this in clinical care?

**Your software function must meet all four criteria to be Non-Device CDS.**

Summary interpretation  
of CDS criteria

1. Your software function does **NOT** acquire, process, or analyze medical images, signals, or patterns.

2. Your software function displays, analyzes, or prints medical information normally communicated between health care professionals (HCPs).

3. Your software function provides recommendations (information/options) to a HCP rather than provide a specific output or directive.

4. Your software function provides the basis of the recommendations so that the HCP does not rely primarily on any recommendations to make a decision.

**Your software function may be non-device CDS.**

Non-Device Examples

Non-Device examples display, analyze, or print the following examples of medical information, which must also not be images, signals, or patterns:

- Information whose relevance to a clinical decision is well understood
- A single discrete test result that is clinically meaningful
- Report from imaging study

AND

Non-Device examples provide:

- Lists of preventive, diagnostic, or treatment options
- Clinical guidelines matched to patient-specific medical info
- Relevant reference information about a disease or condition

AND

Non-Device examples provide:

- Plain language descriptions of the software purpose, medical input, underlying algorithm
- Relevant patient-specific information and other knowns/unknowns for consideration

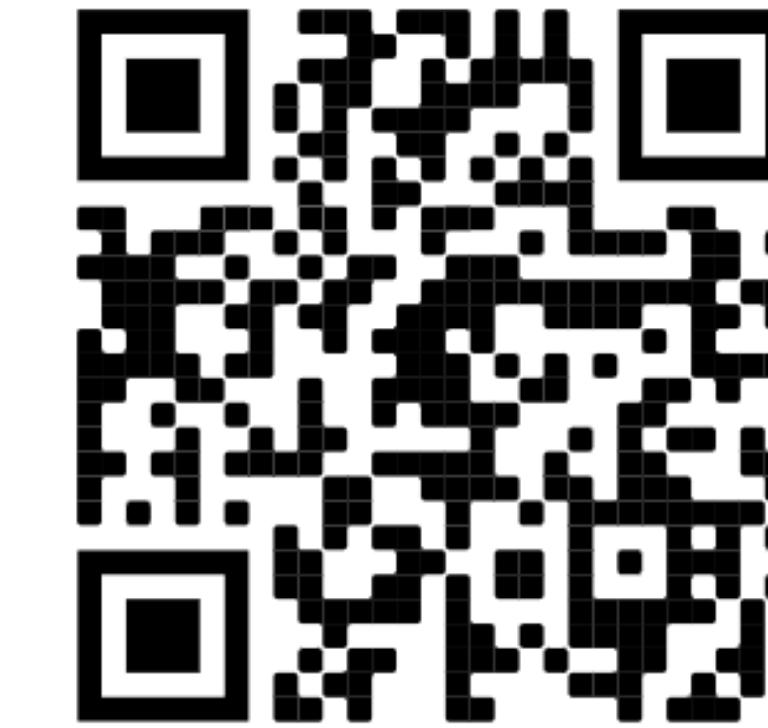
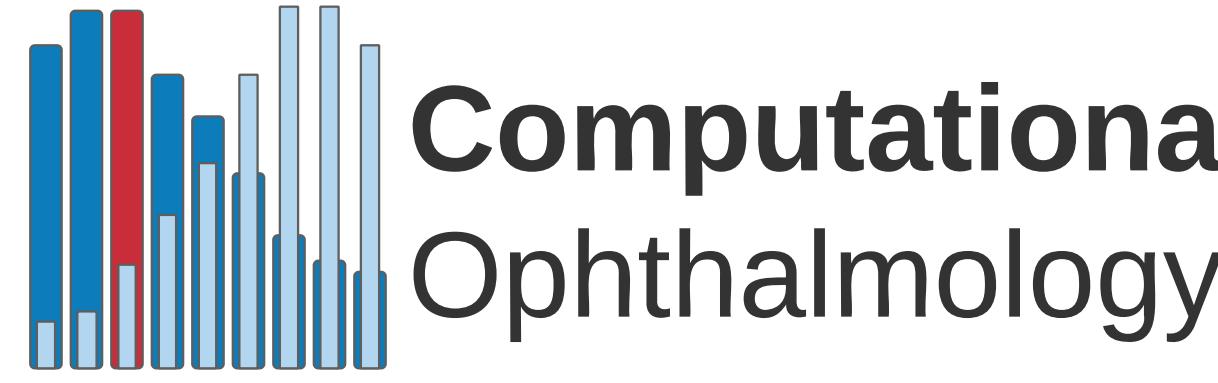
# Acknowledgements

Funding Sources: NIH 1OT2OD032644, NIH/NIA R01 AG060942, NIH/NEI K23EY024921, NIA/NIH U19AG066567, Research to Prevent Blindness, Latham Vision Research Innovation Award, Latham Fund for Vision Research, Donors to Computational Ophthalmology Fund



THE LOWY MEDICAL  
RESEARCH INSTITUTE

Aaron Lee, MD MSCI  
Cecilia Lee, MD MS



<https://comp.ophthalmology.uw.edu>

Megan Lacy MS  
Julia Owen PhD  
Yue Wu PhD  
Scott Song BA

Missy Takahashi BS  
Ashley Batchelor MS  
Matthew Hunt BS  
Theodore Spaide PhD

Emily Heindsmann MA  
Christina Duong BS COA  
Yelena Bagdasarova PhD

Marian Blazes MD  
Jamie Shaffer MS  
Yuka Kihara MS  
Randy Lu BS



Research to  
Prevent Blindness