

Hands-On Exercise 1: Image Classification

Julia Owen, PhD
University of Washington

Setting up today's classification problem

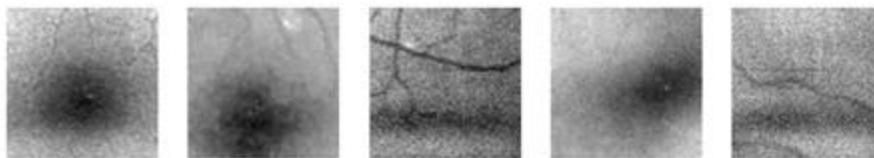
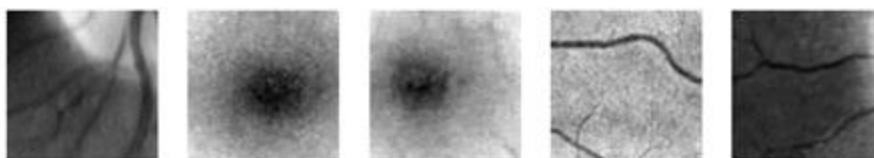
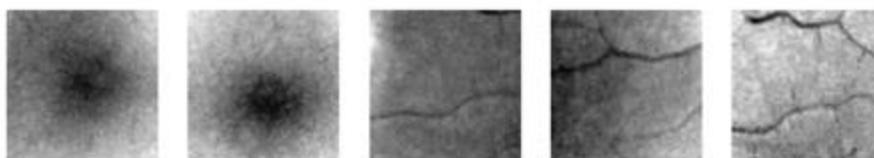
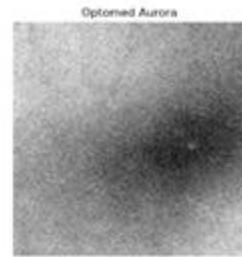
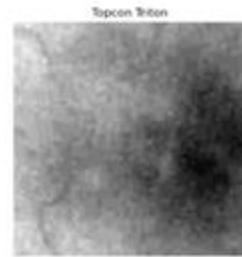
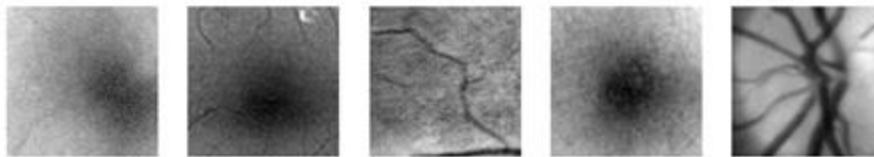
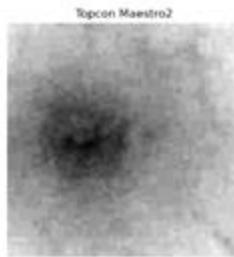
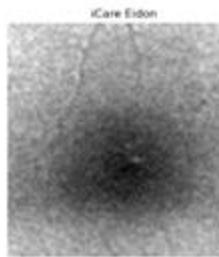
Classifying images using DL is an important area of clinical application in ophthalmology

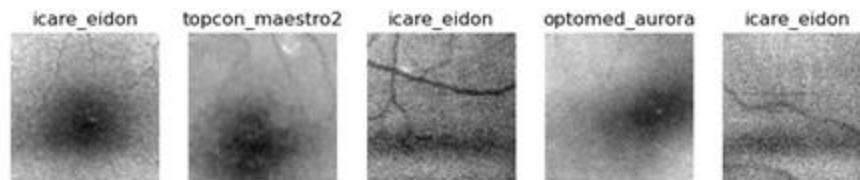
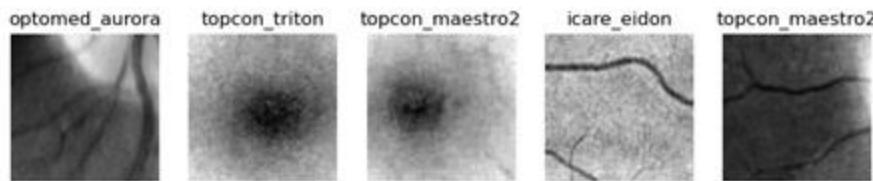
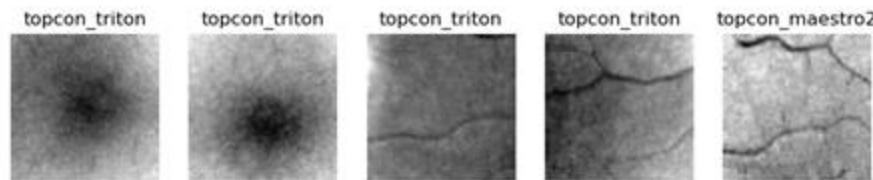
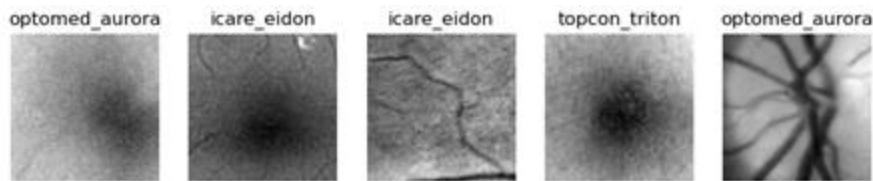
Today we will explore a toy DL problem, but everything we learn can be applied to real-world challenges

In the AI-READI dataset, we have CFPs from 4 imaging devices

These 4 devices create images that are fairly distinct in color and the classification problem is not very difficult. So we design a more difficult problem.

Can you classify these images derived from the images on the previous screen?





So our challenge today is to develop a deep learning model to classify the device from these cropped, grayscale images.

This will demonstrate the power of deep learning.

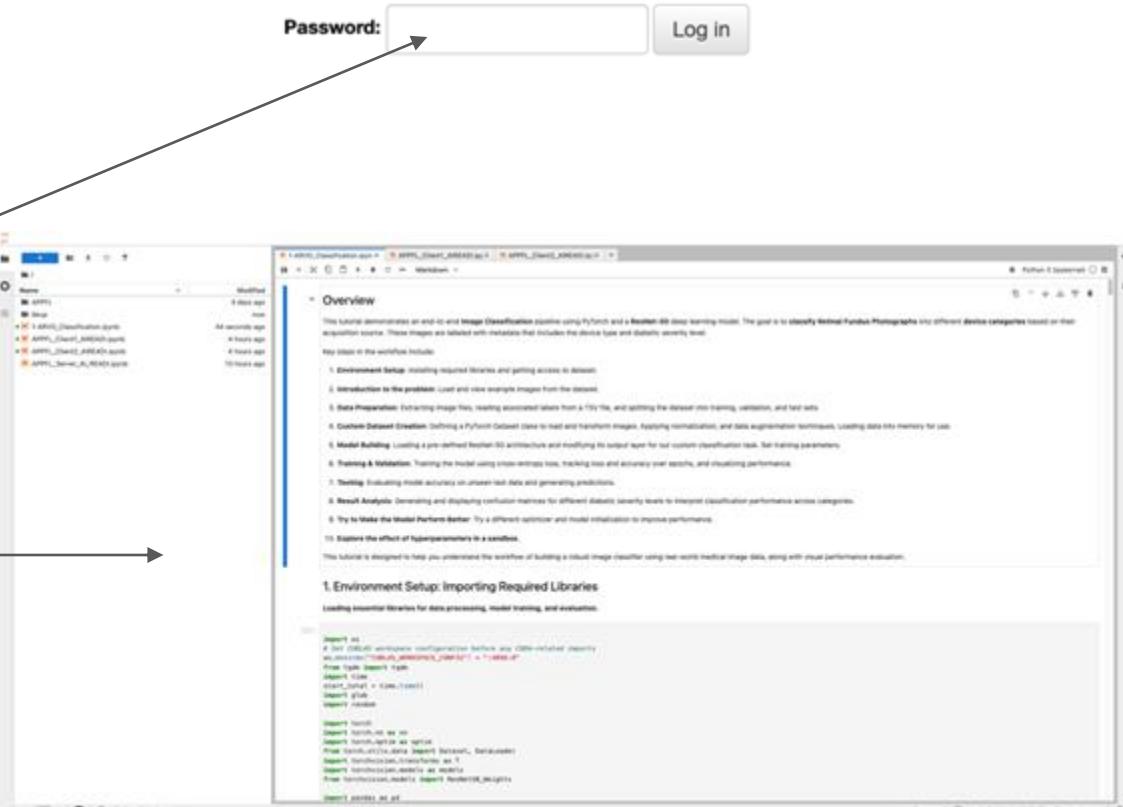
Today's plan

1. Introduce you to Jupyter notebooks
2. Make the plots I just showed you
3. Set up a deep learning model to perform the task we tried to do by eye
4. Train the model
5. Examine the results
6. Explore other settings to improve performance

Let's get everyone logged onto an instance



- Type the url you were given into an internet browser, example
<http://5.936.198:8027>



The screenshot shows a Jupyter Notebook interface with several tabs at the top: "1-ARVO_Classification.ipynb" (active), "APPFL_Client1_AIREADI.ipynb", "APPFL_Client2_AIREADI.ipynb", and "Python 3 (ipykernel)". On the left, a sidebar lists files in the directory: "Name", "APPFL", "bkup", "1-ARVO_Classification.ipynb", "APPFL_Client1_AIREADI.ipynb", "APPFL_Client2_AIREADI.ipynb", and "APPFL_Server_AI_READI.ipynb". A red circle highlights the three top icons in the toolbar: "Run Cell", "Stop Cell", and "Restart Kernel". Below the sidebar, a box labeled "Files in directory" has an arrow pointing to the sidebar. The main area contains an "Overview" section with text about a deep learning pipeline for retinal fundus photographs. Three buttons are shown: "Run the current cell", "Stop the current cell", and "Restart the kernel, resets things". A numbered list from 5 to 10 describes steps in the workflow. The bottom of the overview section states the tutorial's purpose: "This tutorial is designed to help you understand the workflow of building a robust image classifier using real-world medical image data, along with visual performance evaluation." The next section, "1. Environment Setup: Importing Required Libraries", shows code cells. The first cell, which imports various Python libraries, has a red circle around its status indicator "(1)". The output for this cell is: "Blank: never been run". The second cell, which imports torch and torchvision, has a status indicator "(2)". Its output is: "Number: finished running (numbers are sequential)". The third cell, which imports pandas, has a status indicator "*". Its output is: "*: The code is still running." To the right of the notebook, a large text box reads: "Do not run the cells out of order!".

Files in directory

1-ARVO_Classification.ipynb X APPFL_Client1_AIREADI.ipynb X APPFL_Client2_AIREADI.ipynb X Python 3 (ipykernel) 0

Modified
4 days ago
now
44 seconds ago
4 hours ago
4 hours ago
13 hours ago

Overview

This tutorial demonstrates an end-to-end Image Classification pipeline using PyTorch and a ResNet-50 deep learning model. The goal is to classify Retinal Fundus Photographs into different device categories based on their acquisition source. These images are labeled with metadata that includes the device type and diabetic severity level.

Run the current cell

Stop the current cell

Restart the kernel, resets things

and splitting the dataset into training, validation, and test sets.
images. Applying normalization, and data augmentation techniques. Loading data into memory for use.

5. Model Building: Loading a pre-defined ResNet-50 architecture and modifying its output layer for our custom classification task. Set training parameters.
6. Training & Validation: Training the model using cross-entropy loss, tracking loss and accuracy over epochs, and visualizing performance.
7. Testing: Evaluating model accuracy on unseen test data and generating predictions.
8. Result Analysis: Generating and displaying confusion matrices for different diabetic severity levels to interpret classification performance across categories.
9. Try to Make the Model Perform Better: Try a different optimizer and model initialization to improve performance.
10. Explore the effect of hyperparameters in a sandbox.

This tutorial is designed to help you understand the workflow of building a robust image classifier using real-world medical image data, along with visual performance evaluation.

1. Environment Setup: Importing Required Libraries

Loading essential libraries for data processing, model training, and evaluation

(1)

```
import os
# Set CUDA_VISIBLE_DEVICES environment variable
os.environ["CUDA_VISIBLE_DEVICES"] = "0"
from tqdm import tqdm
import time
start_time = time.time()
import glob
import random

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as T
import torchvision.models as models
from torchvision.models import ResNet50_Weights
import pandas as pd
```

Blank: never been run

Number: finished running (numbers are sequential)

*: The code is still running.

Notebooks are an interactive way to write and run code.

Do not run the cells out of order!

Ok now we are ready to start training a model!

1. Environment Setup: Importing Required Libraries

Loading essential libraries for data processing, model training, and evaluation.

```
[1]: import os
# Set CUBLAS workspace configuration before any CUDA-related imports
os.environ["CUBLAS_WORKSPACE_CONFIG"] = ":4096:8"
from tqdm import tqdm
import time
start_total = time.time()
import glob
import random

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as T
import torchvision.models as models
from torchvision.models import ResNet50_Weights

import pandas as pd

from PIL import Image

import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import label_binarize
from sklearn.metrics import roc_curve, auc

from itertools import cycle
from IPython.display import clear_output

torch.use_deterministic_algorithms(True)
```

2. Introduction to the problem: Load and view example images from the dataset.

2a) Problem: Classify device from cropped, grayscale CFP images

We have CFP images from 4 devices: iCare Eidon, Topcon Maestro2, Topcon Triton, and Optomed Aurora.

2b) We can load an image for one participant/eye and view.

```
[2]: img_eidon = np.asarray(Image.open('APPFL/examples/full/icare_eidon/1270_eidon_uwf_centered_cfp_r_1.2.826.0.1.3680043.8.641.1.20240510.232115.43583.jpg'))
img_maestro = np.asarray(Image.open('APPFL/examples/full/topcon_maestro2/1270_maestro2_3d_macula_cfp_r_2.16.840.1.114517.10.5.1.4.907063120240510160006.2.1.jpg'))
img_triton = np.asarray(Image.open('APPFL/examples/full/topcon_triton/1274_triton_macula_12x12_cfp_r_2.16.840.1.114517.10.5.1.4.94005520240514143050.2.1.jpg'))
img_optomed = np.asarray(Image.open('APPFL/examples/full/optomed_aurora/1270_optomed_mac_or_disk_centered_cfp_r_2.25.2183161925995491838121778870991254610818.jpg'))

[3]: plt.figure(figsize=(20,50))
plt.subplot(1,4,1)
plt.imshow(img_eidon)
plt.title('iCare Eidon')
plt.axis('off')

plt.subplot(1,4,2)
plt.imshow(img_maestro)
plt.title('Topcon Maestro2')
plt.axis('off')

plt.subplot(1,4,3)
plt.imshow(img_triton)
plt.title('Topcon Triton')
plt.axis('off')

plt.subplot(1,4,4)
plt.imshow(img_optomed)
plt.title('Optomed Aurora')
plt.axis('off')
```

3. Data Preparation ¶

3a) Loading Dataset Labels: We read the labels from a TSV file, which contains metadata about images.

```
[13]: # importing data label
tsv_path = "APPFL/examples/cfp_images/labels.tsv"
df = pd.read_csv(tsv_path, sep='\t')

# printing top 5
df.head(5)
```

```
[13]:   subject_id      device    protocol laterality partition          dm_severity           file_path
0       7079  optomed_aurora  mac_or_disk_centered      r     train  pre_diabetes_lifestyle_controlled  optomed_aurora/7079_optomed_mac_or_disk_center...
1       7174  optomed_aurora  mac_or_disk_centered      l     train  oral_medication_and_or_non_insulin_injectable...  optomed_aurora/7174_optomed_mac_or_disk_center...
2       7103  optomed_aurora  mac_or_disk_centered      r     train  oral_medication_and_or_non_insulin_injectable...  optomed_aurora/7103_optomed_mac_or_disk_center...
3       7103  optomed_aurora  mac_or_disk_centered      r     train  oral_medication_and_or_non_insulin_injectable...  optomed_aurora/7103_optomed_mac_or_disk_center...
4       7097  optomed_aurora  mac_or_disk_centered      l      val  oral_medication_and_or_non_insulin_injectable...  optomed_aurora/7097_optomed_mac_or_disk_center...
```

```
[14]: #print the first 20 lines
print(df[['subject_id','device','partition']].head(20))
```

```
   subject_id      device partition
0       7079  optomed_aurora     train
1       7174  optomed_aurora     train
2       7183  optomed_aurora     train
3       7183  optomed_aurora     train
4       7097  optomed_aurora      val
5       4028  optomed_aurora     train
6       4152  optomed_aurora     train
7       4152  optomed_aurora     train
8       1035  optomed_aurora     test
9       1132  optomed_aurora     train
10      7198  optomed_aurora      val
11      7073  optomed_aurora     train
12      1231  optomed_aurora     test
13      1231  optomed_aurora     test
14      1231  optomed_aurora     test
15      7370  optomed_aurora     train
16      7293  optomed_aurora     train
17      1336  optomed_aurora     train
18      1336  optomed_aurora     train
19      1193  optomed_aurora     train
```

3b) Splitting Dataset: The dataset is split into training, validation, and test sets.

```
[15]: # dividing data into different dataframes based on train, validation and test
train_df = df[df["partition"] == "train"].copy()
val_df = df[df["partition"] == "val"].copy()
test_df = df[df["partition"] == "test"].copy()

print("Number of data sample \nTraining", len(train_df), "\nValidation", len(val_df), "\nTest", len(test_df))

Number of data sample
Training 1544
Validation 652
Test 666

[16]: # Finding unique classes and giving it a number index
# in this case we have 4 classes
unique_classes = sorted(train_df["device"].unique())

# Mapping classes to a particular index
class_to_idx = {cls_name: idx for idx, cls_name in enumerate(unique_classes)}
print("Class to index mapping:", class_to_idx)

train_df[["label_idx"]] = train_df["device"].map(class_to_idx)
val_df[["label_idx"]] = val_df["device"].map(class_to_idx)
test_df[["label_idx"]] = test_df["device"].map(class_to_idx)

num_classes = len(unique_classes)

Class to index mapping: {'icare_eidon': 0, 'optomed_aurora': 1, 'topcon_maestro2': 2, 'topcon_triton': 3}
```

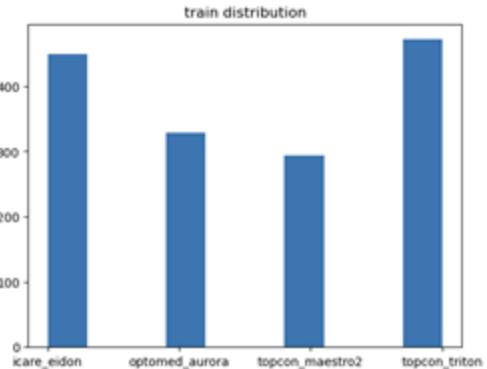
2,862 images total
Roughly 60-20-20 split

54% is from training set -> will be used to train our model
23% is for validation -> will be used to evaluate model during development period
23% is for testing -> will be saved for the end when we evaluate our model

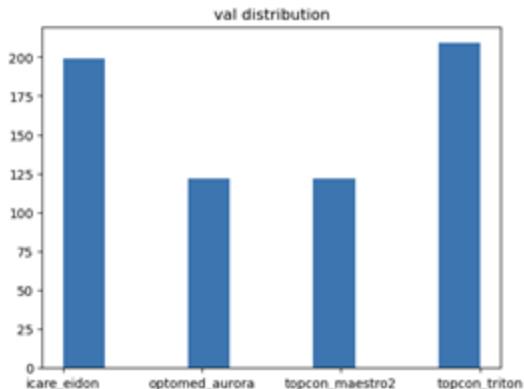
Patient-level split - no leakage between the partitions

3c) Distribution of the Labels: We can plot the distribution of the labels across the three partitions.

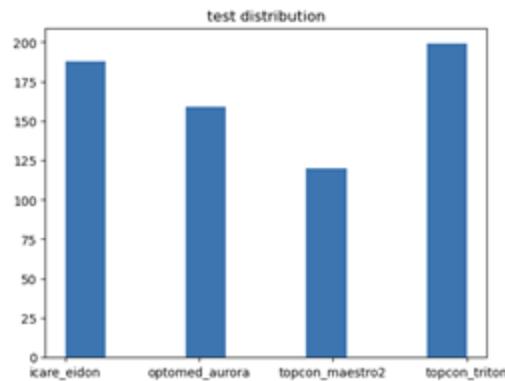
```
[17]: train_labels = sorted(train_df["device"])
_ = plt.hist(train_labels)
plt.title('train distribution')
plt.show()
```



```
[18]: val_labels = sorted(val_df["device"])
_ = plt.hist(val_labels)
plt.title('val distribution')
plt.show()
```



```
[19]: test_labels = sorted(test_df["device"])
_ = plt.hist(test_labels)
plt.title('test distribution')
plt.show()
```



4. Defining the Custom Dataset Class

4a) This class is used for loading images (cropping and converting to grayscale) and applying transformations.

```
[20]: # Defining a dataloader class, which returns image and its label
class AIRAIDDataset_grayscalecrop(Dataset):
    def __init__(self, df, transform=None, preload=False):
        ...
        Args:
            df: a DataFrame with at least ['file_path', 'label_idx'] columns
            transform: torchvision transforms (augmentations) to apply
        ...
        self.df = df.reset_index(drop=True)
        self.transform = transform
        self.preload = []
        if preload:
            partition = self.df["partition"][0]
            npsvfn = f"APPFL/examples/cfp_images/grayscalecrop-{partition}.npy"
            if os.path.isfile(npavfn):
                self.preload = np.load(npavfn)
                print(f'loaded: {partition}')
            else:
                for i in tqd(range(0, len(self.df))):
                    row = self.df.iloc[i]
                    img_path = "APPFL/examples/cfp_images/" + row["file_path"]
                    image = Image.open(img_path).convert("RGB")
                    img_array = np.array(image)
                    #Take only the 224x224 center square of the image
                    h, w, _ = img_array.shape
                    top = (h - 224) // 2
                    left = (w - 224) // 2
                    center_crop = img_array[top:top+224, left:left+224]
                    #Convert to grayscale
                    gray = np.dot(center_crop[:, :, 0], [0.299, 0.587, 0.114])
                    gray = gray / 255
                    self.preload.append(gray)
    def __len__(self):
        return len(self.df)

    def __getitem__(self, index):
        if self.preload:
            img = self.preload[index]
        else:
            row = self.df.iloc[index]
            img_path = "APPFL/examples/cfp_images/" + row["file_path"]
            image = Image.open(img_path).convert("RGB")
            img_array = np.array(image)
            #Take only the 224x224 center square of the image
            h, w, _ = img_array.shape
            top = (h - 224) // 2
            left = (w - 224) // 2
            center_crop = img_array[top:top+224, left:left+224]
            #Convert to grayscale
            gray = np.dot(center_crop[:, :, 0], [0.299, 0.587, 0.114])
            gray = gray / 255
            self.preload.append(gray)
        if self.transform:
            img = self.transform(img)
        return img, row['label_idx']
```

4b) Transformation of all images, normalization to ImageNet statistics Augmenting training images to have horizontal flips and rotations for online augmentation.

```
[21]: # Defining transformation of image
# Augmenting training data to have horizontal flips and rotations
train_transform = T.Compose([
    T.RandomHorizontalFlip(p=0.5), # Randomly flip images horizontally
    T.RandomRotation(degrees=15), # Randomly rotate images
    T.ToTensor(),
    T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]) # Normalize using ImageNet stats
])

val_transform = T.Compose([
    T.ToTensor(),
    T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

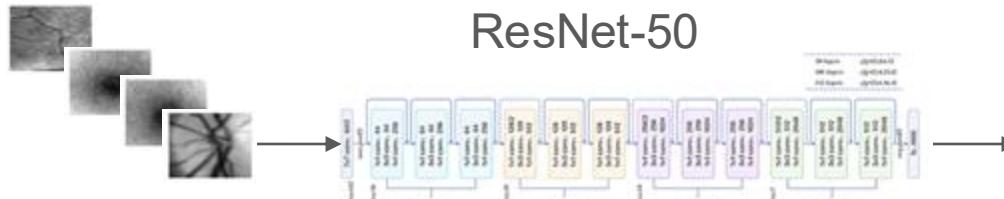
4c) Call the loader class with partition filepaths and transforms

```
[22]: # Loading datasets into memory for faster training
train_dataset = AIRAIDDataset_grayscalecrop(train_df, train_transform, preload=True)
val_dataset = AIRAIDDataset_grayscalecrop(val_df, val_transform, preload=True)
test_dataset = AIRAIDDataset_grayscalecrop(test_df, val_transform, preload=True)
```

```
loaded: train
loaded: val
loaded: test
```

Image augmentation, helps prevent overfitting and helps with generalization (especially with small datasets)

Quick review of deep learning fundamentals



Loss function:
Cross Entropy

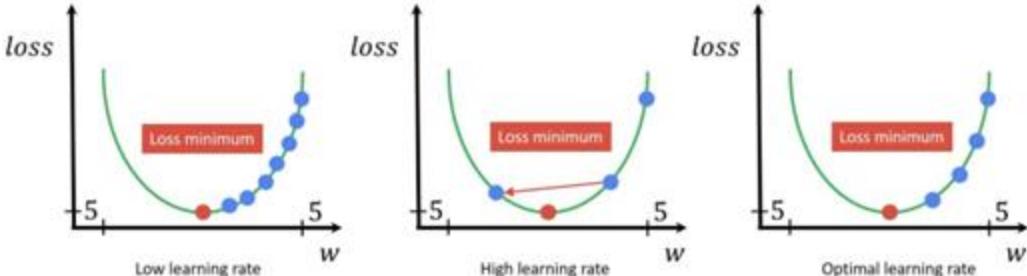
True Labels:
0 = iCare Eidon
1 = Optomed Aurora.
2 = Topcon Maestro2
3 = Topcon Triton

Predicted Labels:
0 = iCare Eidon
1 = Optomed Aurora.
2 = Topcon Maestro2
3 = Topcon Triton

Batch Size:
The number of images to use to compute each step

Number of epochs:
The number time to go through the whole training set

Optimizer and learning rate



5. Model Building

5a) Set batch size and create dataloader instances.

```
[23]: # Define batch size for data loading, batch size is the number of training examples used during one iteration of model training
batch_size = 64
torch.set_num_threads(os.cpu_count())
num_workers = os.cpu_count()

# Create Dataloader instances for efficient batch processing
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
```

5b) Define device to use for training and inference.

```
[24]: # Define the device (GPU if available, else CPU), if this says CPU please raise your hand
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

Using device: cuda
```

5c) Define methods used during model training, defines what to do each train step, validation inference, and plotting the learning curves.

```
[25]: # Methods to help with model training
def train_step(model, dataloader, optimizer, criterion, device, batch_losses):
    """Performs a single training step over the dataset."""
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    progress_bar = tqdm(dataloader, desc="Training", leave=True)

    for batch_idx, (images, labels, _) in enumerate(progress_bar):
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)

        _, preds = torch.max(outputs, 1)
        correct += (preds == labels).sum().item()
        total += labels.size(0)

        batch_losses.append(loss.item())
        progress_bar.set_postfix(batch_loss=loss.item())

    epoch_loss = running_loss / total
    epoch_acc = correct / total

    return epoch_loss, epoch_acc

def validate_step(model, dataloader, criterion, device):
    """Performs a single validation step over the dataset."""
    model.eval()
    total_loss = 0.0
    correct = 0
```

5d) Defining model architecture: a pre-existing RESNET50 model architecture.

```
[26]: #this will make the initialization the same each time it is run for the same setting of parameters, just run once
def set_seed(seed=42):
    random.seed(seed)                      # Python built-in random
    np.random.seed(seed)                    # NumPy
    torch.manual_seed(seed)                # PyTorch CPU
    torch.cuda.manual_seed(seed)           # PyTorch CUDA (if using GPU)
    torch.cuda.manual_seed_all(seed)        # All CUDA devices (if using multi-GPU)

set_seed(12)
#Ensure deterministic behavior
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
```

```
[27]: # Load the ResNet50 model
model = models.resnet50(weights=None)

# Modify the last layer for our classification task
# Changing to 4 classes as we have 4 devices to classify it into
model.fc = nn.Linear(model.fc.in_features, num_classes)

# Transfer model to device (GPU or CPU)
model = model.to(device)
```

5e) Choose a cost function, an optimizer and set learning rate.

```
[28]: # Define loss function, the function that we want to minimize to find the best assignment of labels
criterion = nn.CrossEntropyLoss() #CE quantifies how well a model's predicted probability distribution aligns with the true, or actual, probability distribution of the data

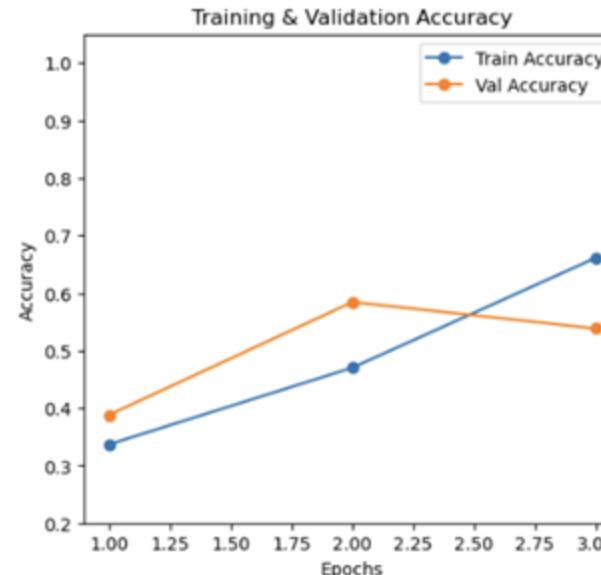
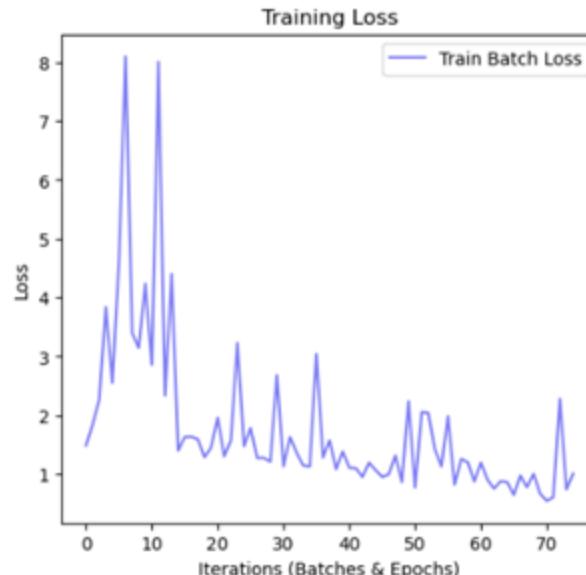
#set the optimizer and learning rate (lr), there are different mathematical methods to minimize loss functions, the learning rate dictates how big a step you take toward the minima
optimizer = optim.SGD(model.parameters(), lr = 1e-2, momentum=0.9)
```

6. Running training: Call training function with desired number of epochs.

```
[29]: #The number of epoch sets how many the model learns over the entire dataset in
      start = time.time()
      train_model(model=model, train_loader=train_loader,
                  val_loader=val_loader,
                  optimizer=optimizer,
                  criterion=criterion,
                  device=device,
                  num_epochs=24)
      print("Training time: ", time.time() - start)
```

Start training

Epoch [3/24] - Train Loss: 1.0995, Train Acc: 0.6613 | Val Loss: 4.3258, Val Acc: 0.5383



7. Testing model: Perform inference with model on validation dataset

7a) Our model hasn't seen this data during training so gives us an indication on how it will perform on new data

```
[30]: # ----- VAL -----
all_preds = []
all_labels = []
all_dm_severity = []
print("Evaluating on val set...")
if len(val_df) > 0:
    model.eval()
    val_loss = 0.0
    val_correct = 0
    val_total = 0

    with torch.no_grad():
        for images, labels, dm_severity in val_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            loss = criterion(outputs, labels)
            val_loss += loss.item() * images.size(0)

            _, preds = torch.max(outputs, 1)
            val_correct += (preds == labels).sum().item()
            val_total += labels.size(0)
            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
            all_dm_severity.extend(dm_severity)

    val_loss /= val_total
    val_acc = val_correct / val_total
    print(f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}")

Evaluating on val set...
Val Loss: 1.3090, Val Acc: 0.7914
```

7b) Saving test set results to a dataframe for use below.

```
[31]: # Create a DataFrame to store results
df_results_val = pd.DataFrame({
    'true_device': all_labels,
    'pred_device': all_preds,
    'dm_severity': all_dm_severity
})
```

Ba) Plot confusion matrix of device type for the val set. This allows for interrogation of model performance on each class (device).

```
[32]: #define a function to plot a confusion matrix (cm)
def plot_confusion_matrix(cm, device_names, title="Confusion Matrix"):
    plt.figure()
    plt.imshow(cm, interpolation='nearest', aspect='auto')
    plt.title(title)
    plt.colorbar()

    tick_marks = np.arange(len(device_names))
    plt.xticks(tick_marks, device_names, rotation=45, ha='right')
    plt.yticks(tick_marks, device_names)

    # Optionally annotate cells with counts
    thresh = cm.max() / 2.
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            plt.text(j, i, str(cm[i, j]),
                     horizontalalignment="center", color="white")
            #color="white" if cm[i, j] > thresh else "black")

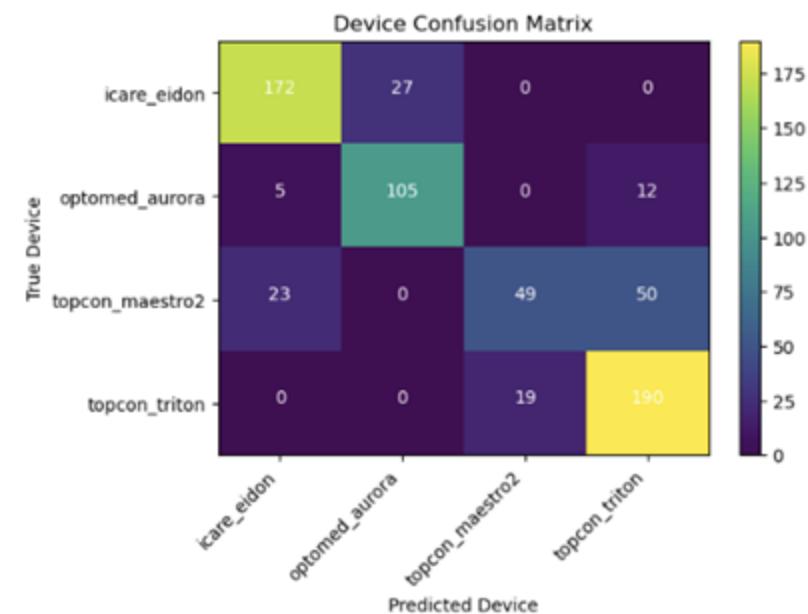
    plt.tight_layout()
    plt.xlabel('Predicted Device')
    plt.ylabel('True Device')
    plt.show()
```

```
[33]: # Generate confusion matrices with confusion_matrix function from scikit-learn
true_devs = df_results_val['true_device'].values
pred_devs = df_results_val['pred_device'].values

cm = confusion_matrix(true_devs, pred_devs)

idx_to_device = {v: k for k, v in class_to_idx.items()}
device_names = [idx_to_device[i] for i in sorted(idx_to_device.keys())]

#calls function written above
plot_confusion_matrix(cm, device_names, title=f"Device Confusion Matrix")
```



8b) Plot Receiver-Operator Curve (ROC)

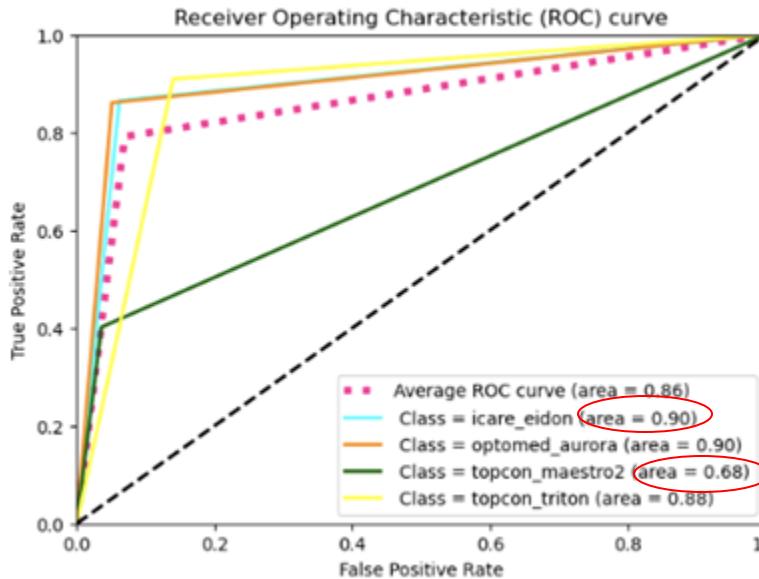
An ROC curve plots the rate of false positives versus the rate of true positives. In our case, we just have one (FPR,TPR) pair and the (0,0) -> everything is predicted 0 and (1,1) points -> everything predicted 1.

[34]: `def plot_roc_curve(y_test, y_pred, class_labels):`

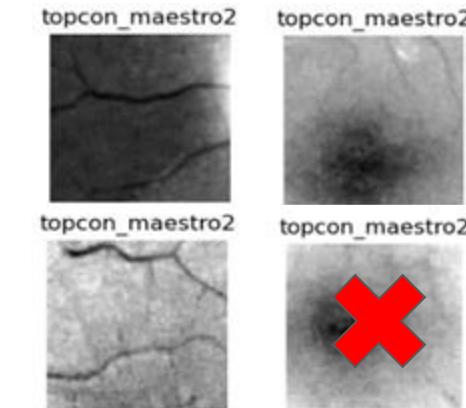
```
n_classes = len(np.unique(y_test))
y_test = label_binarize(y_test, classes=np.arange(n_classes))
y_pred = label_binarize(y_pred, classes=np.arange(n_classes))

# Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], y_pred[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
```

[35]: `plot_roc_curve(df_results_val1['true_device'].values, df_results_val1['pred_device'].values, device_names)`



If we consider the topcon_maestro2 class:
True positive - correctly labeled
topcon_maestro2
False positive - image incorrectly labeled
topcon_maestro2



9. Improve model performance

9a) Let's explore a different optimizer and learning rate.

```
[36]: # Define batch size for data loading
batch_size = 64
torch.set_num_threads(os.cpu_count())
num_workers = os.cpu_count()

# Create DataLoader instances for efficient batch processing
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))

# Define the device (GPU if available, else CPU)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

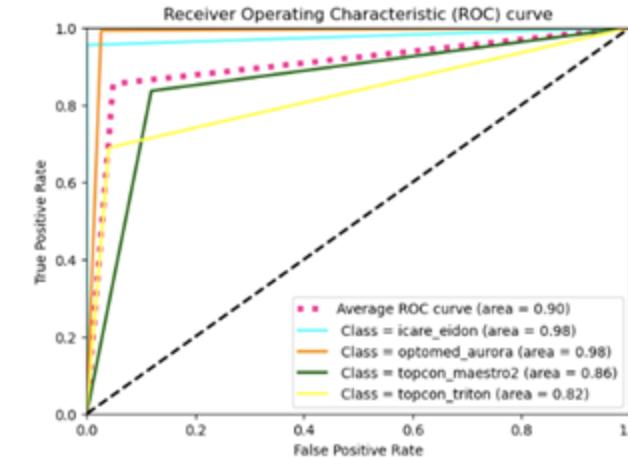
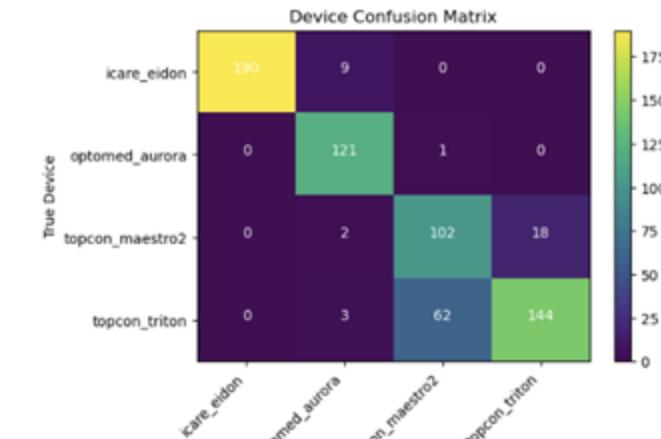
# Load the ResNet50 model
model = models.resnet50(weights=None)
# Modify the last layer for our classification task
# Changing to 4 classes as we have 4 devices to classify it into
model.fc = nn.Linear(model.fc.in_features, num_classes)
# Transfer model to device (GPU or CPU)
model = model.to(device)

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=7e-6, weight_decay=1e-3) #-- this is what we are changing!!

start = time.time()
train_model(model=model, train_loader=train_loader,
            val_loader=val_loader,
            optimizer=optimizer,
            criterion=criterion,
            device=device,
            num_epochs=24)
print("Training time: ", time.time() - start)
```

Start training

Run the next 3 cells



9b) Second the weight initialisation will be changed by using the pretrained weights (ImageNet) for the ResNet50.

```

# Define batch size for data loading
batch_size = 64
torch.set_num_threads(os.cpu_count())
num_workers = os.cpu_count()

# Create DataLoaders instances for efficient batch processing
train_loader = DataLoaders(train_dataset, batch_size=batch_size, shuffle=True, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(torch.getpid()))
val_loader = DataLoaders(val_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(torch.getpid()))
test_loader = DataLoaders(test_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(torch.getpid()))

# Define the device (GPU if available, else CPU)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

# Load the ResNet50 model
model = models.resnet50(weights=ResNet50_Weights.IMAGENET_V2) #--> This is what we are changing!
# Modify the last layer for our classification task
# Changing to 4 classes as we have 4 devices to classify it into
model.fc = nn.Linear(model.fc.in_features, num_classes)
# Transfer model to device (GPU or CPU)
model = model.to(device)

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()

# Set the optimizer and learning rate (lr)
optimizer = optim.Adam(model.parameters(), lr=7e-6, weight_decay=1e-3)

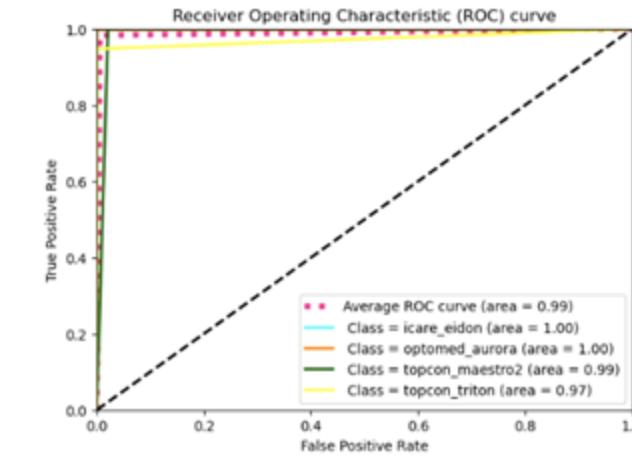
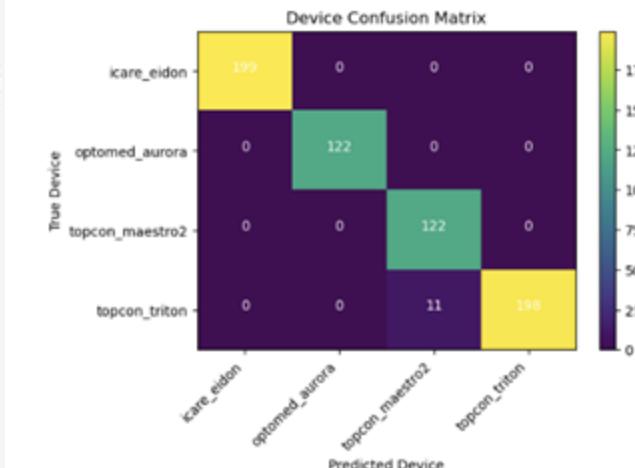
start = time.time()
train_model(model=model, train_loader=train_loader,
            val_loader=val_loader,
            optimizer=optimizer,
            criterion=criterion,
            device=device,
            num_epochs=30)
print("Training took: ", time.time() - start)

```

Start training

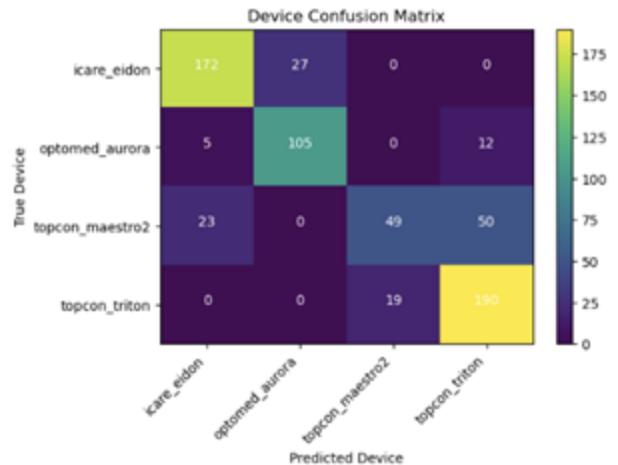


Run the next 3 cells

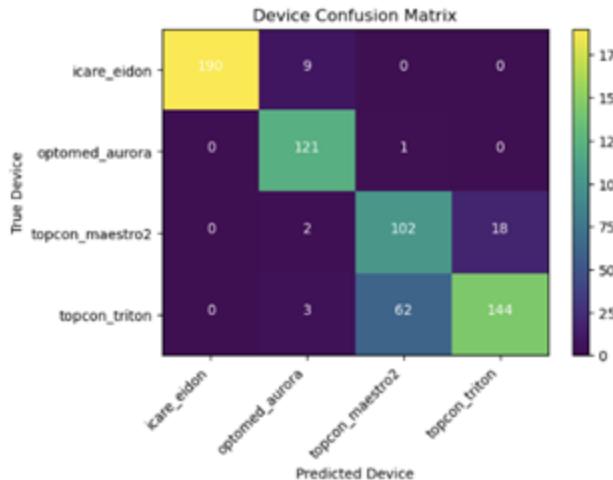


A decision to make!

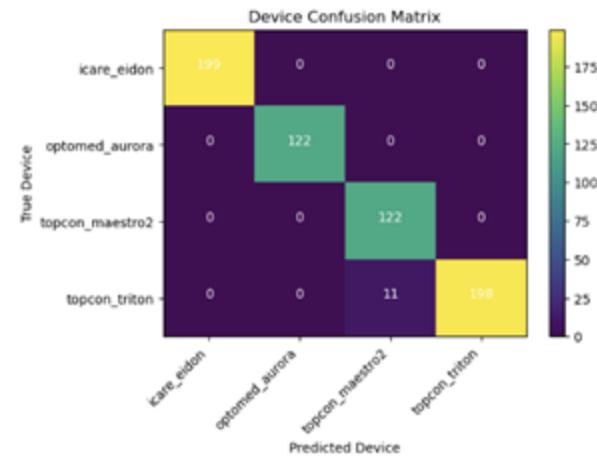
Model 1: Val accuracy: 79.14%



Model 2: Val accuracy: 85.43%



Model 3: Val accuracy: 98.31%



10. Choose a model to evaluate on the test data

The model has not seen these images so this is the least biased way to evaluate a model's performance. But you can only run one model on your test set, so choose carefully.

```
[44]: # Evaluate on the test set

all_preds = []
all_labels = []
all_dm_severity = []
print("Evaluating on test set...")
if len(test_df) > 0:
    model.eval()
    test_loss = 0.0
    test_correct = 0
    test_total = 0

    with torch.no_grad():
        for images, labels, dm_severity in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            loss = criterion(outputs, labels)
            test_loss += loss.item() * images.size(0)

            _, preds = torch.max(outputs, 1)
            test_correct += (preds == labels).sum().item()
            test_total += labels.size(0)
            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
            all_dm_severity.extend(dm_severity)

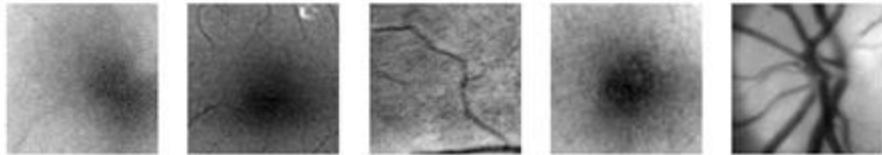
    test_loss /= test_total
    test_acc = test_correct / test_total
    print(f"Test Loss: {test_loss:.4f}, Test Acc: {test_acc:.4f}")

    # Create a DataFrame to store results
df_results_test = pd.DataFrame({
    'true_device': all_labels,
    'pred_device': all_preds,
    'dm_severity': all_dm_severity
})
```

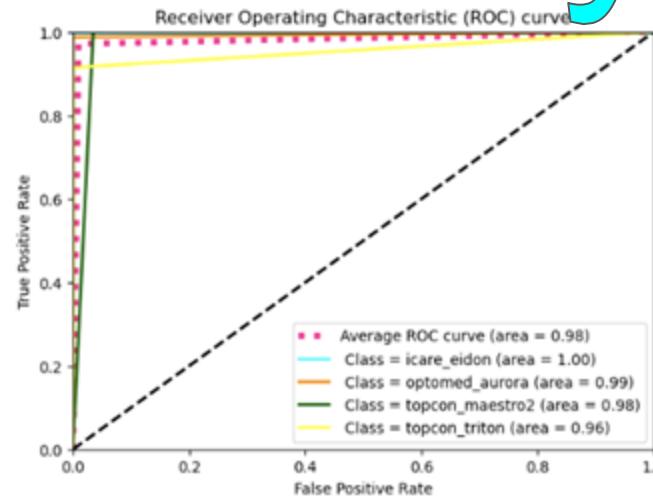
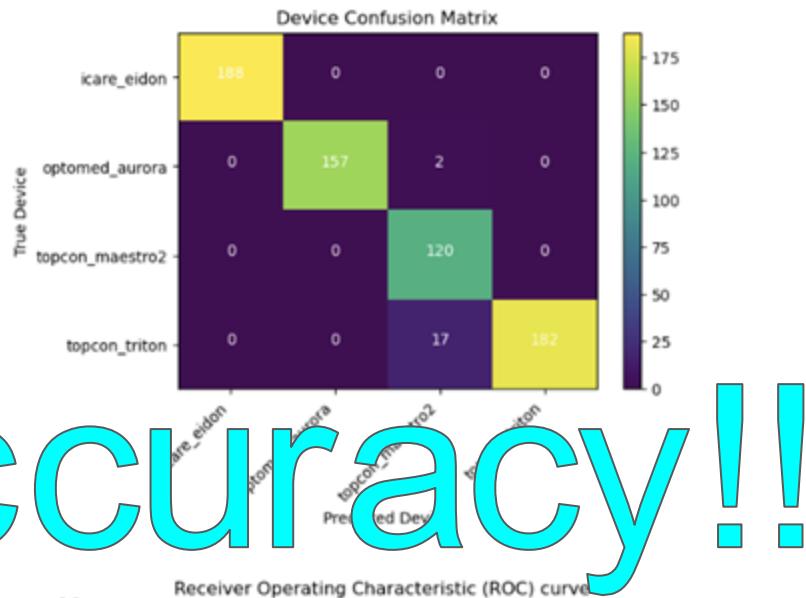
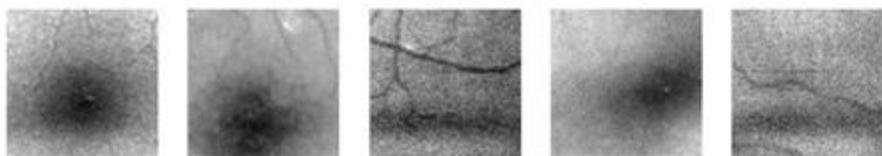
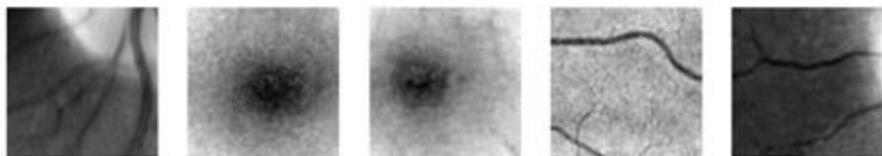
```
[45]: # Generate confusion matrices with confusion_matrix function from scikit-learn
true_devs = df_results_test['true_device'].values
pred_devs = df_results_test['pred_device'].values
cm = confusion_matrix(true_devs, pred_devs)

idx_to_device = {v: k for k, v in class_to_idx.items()}
device_names = [idx_to_device[i] for i in sorted(idx_to_device.keys())]

#calls function written above
plot_confusion_matrix(cm, device_names, title=f"Device Confusion Matrix")
```



97.5% accuracy!!



If you have time.....

11. Sandbox: Explore hyperparameter settings.

This code was set up for you to explore what happens when you change hyperparameters.

Things you could try:

1. Change data augmentation parameters.
2. Change batch size (warning you may run into memory constraints)
3. Change the initialization seed to any number
4. Change the initialization of the weights (model = models.resnet50(weights = ResNet50_Weights.IMAGENET1K_V2))
5. Change the optimizer and/or learning rate (<https://pytorch.org/docs/main/optim.html>)
6. Change the number of epochs (you want to see model convergence, training curve is flat at end)

```
# Defining transformation of image
# Augmenting training data to have horizontal flips and rotations
train_transform = T.Compose([
    T.RandomHorizontalFlip(p=0.5), # Randomly flip images horizontally
    T.RandomRotation(degrees=35), # Randomly rotate images
    T.ToTensor(),
    T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]) # Normalize using ImageNet stats
])
val_transform = T.Compose([
    T.ToTensor(),
    T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Loading datasets into memory for faster training
train_dataset = ADREAD3Dataset_grayscalecrop(train_df, train_transform, preload=True)
val_dataset = ADREAD3Dataset_grayscalecrop(val_df, val_transform, preload=True)
test_dataset = ADREAD3Dataset_grayscalecrop(test_df, val_transform, preload=True)

# Define batch size for data loading
batch_size = 64
torch.set_num_threads(os.cpu_count())
num_workers = os.cpu_count()

# Create Datalader instances for efficient batch processing
train_loader = DataLoaderr(train_dataset, batch_size=batch_size, shuffle=True, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
val_loader = DataLoaderr(val_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))
test_loader = DataLoaderr(test_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers, pin_memory=True, worker_init_fn=lambda _: np.random.seed(42))

# Define the device (GPU if available, else CPU)
device = torch.device("cuda") if torch.cuda.is_available() else "cpu"
print(f"Using device: {device}")

# This will make the initialization the same each time it is run for the same setting of parameters
set_seed(12)

# Load the ResNet50 model
model = models.resnet50(weights=None)
# Modify the last layer for our classification task
# Changing to 4 classes as we have 4 devices to classify it into
model.fc = nn.Linear(model.fc.in_features, num_classes)
# Transfer model to device (GPU or CPU)
model = model.to(device)
```

Acknowledgements

Ravi K. Madduri

Argonne National Laboratory

Aditya Sinha

University of Illinois, Urbana-Champaign

Aaron Lee

University of Washington

All the Lee Lab who helped us test!

Go check out the notebooks later! <https://github.com/AdioosinUIUC/aireadi-course>